*Royal Institute of Philosophy*

A-Level Philosophy Moral Philosophy Resources

Michael Lacewing

*Christ's Hospital*

Table of contents

The guide and its purpose

Introduction

Glossary

THE GUIDE AND ITS PURPOSE

This guide provides an introduction to the topics in moral philosophy studied at A level. It is primarily focused on the AQA (2017) Philosophy syllabus, but almost all of it will be equally useful to moral philosophy topics in Religious Studies syllabi, as well as serving as a general introduction for anyone coming to moral philosophy for the first time.

As a study guide, it is intended to support learning and teaching. While it is written to be clear and accessible, philosophy always requires careful reading and thoughtful questioning, and this guide is no different. Each of the three sections closes with a list of study questions, and some suggestions for further reading and internet resources. At the end there is also a glossary of key terms (when a term appears in the glossary, it will be in **bold** the first time it appears in the text).

INTRODUCTION

'How should one live?' Perhaps uniquely among animals, human beings not only act, we also consider how we *should* act. We think that there are better and worse ways of acting, we reflect on our experience of making mistakes, and try to improve. Much of this, of course, relates to our own self-interest – meeting our needs, successfully achieving our personal goals, and so on. But that is not all. We are social creatures, we live together, and our lives and actions affect the lives and actions of other people. We are concerned not only for ourselves, but for other people as well, and how other people treat us is critical to our own happiness.

How should we relate to one another, how should we treat one another? How should we live so that each of our lives goes 'best'? What is '**good**' in life and how may we go about trying to attain it?

These questions form the basis for moral philosophy. (I shall make no distinction between the terms 'ethics' and '**morality**'.) We can approach them in three different ways. First, we can try to answer them 'in general', developing and defending theories about people care about or what makes their lives go well. This approach is called ***normative ethics***, and aims to give us general guidance on what is morally right or wrong, good or bad. Second, we can try to answer more specific questions about how we should live or act, such as when it is right to go to war or whether abortion is ever okay. This is the domain of *practical* or ***applied ethics***, which considers particular issues in real life and asks what is the right thing to do. Third, we can ask more about the questions and answers themselves. For example, do we think that these questions have a correct answer, which we must then try to discover? Or are we exploring our personal attitudes towards life and other people? Or perhaps each society – each group of people living together – decides what it will accept for itself? This approach leads us into ***metaethics***, the philosophical study of what morality is.

In this guide, we will consider all three approaches to moral philosophy. We begin by looking at three theories in normative ethics. We then discuss four issues in applied ethics, seeing how each of our three normative theories answers questions about stealing, lying, eating animals, and **simulated killing** (pretending to kill in plays, films, video games, etc.). We end by looking at a range of theories in metaethics in the debate over whether moral judgments, such as 'stealing is wrong', can be said to be objectively true or false or not.

I.  NORMATIVE ETHICAL THEORIES

When thinking about how we should live, we may consider what it is to be a good person, what rules we should live by, or what the consequences of our actions are. Different theories emphasise these different considerations.

1. **Virtue ethics** focuses on what it is to be a *good person*. Ethics is about all of life, not just how we choose to act. We focus on Aristotle's virtue theory. Aristotle appeals to human nature as the basis for how we should live.

2. **Deontology** focuses on *duties* and *rules*. We have moral duties to do things which it is right to do and moral duties not to do things which it is wrong to do. This is sometimes explained in terms of ***intentions*** and sometimes in terms of ***rights***. We will look at Immanuel Kant's theory that tries to base our duties in reason.

3. **Utilitarianism** focuses on the *consequences* of our actions. It claims that an action is right if it brings about more *happiness* than any other action you could have done in that situation. We will look at Jeremy Bentham's and John Stuart Mill's versions of utilitarianism.

Some form of virtue ethics or other has been the 'default form' of ethical theory in Western philosophy really until the Enlightenment in the 18th century. Aristotle's theory was very influential, and many other forms developed from it. For example, in the 12th century, St Thomas Aquinas combined Aristotle's theory with Christianity, which had inherited an emphasis on law from Judaism, to develop 'natural law theory'. Then in the Enlightenment, philosophers sought to free morality from religion. The German philosopher Immanuel Kant kept the emphasis on moral laws, but based them on reason, while the English law reformer Jeremy Bentham based morality on pleasure and pain.

A. Aristotle's virtue ethics

*Eudaimonia: The good for human beings*
Aristotle begins the *Nicomachean Ethics* with the question 'What is the good for human beings?' We all want to live a successful, fulfilling, good life – but what is this? Our different activities aim at various 'goods'. For example, education aims at knowledge; medicine aims at health. For any action or activity, there is a purpose (a 'why') for which we undertake it – its **end** (*telos*).

Many activities we do for the sake of something further. For example, in medicine, we research new drugs, make surgical implements, diagnose illnesses, etc. These activities are undertaken for the sake of the overall end of medicine – health. But why be healthy? Is there a further purpose? We might say 'no', or 'to be happy', or 'because a good life is a healthy life'.

Aristotle argues that everything we do, we do in order to live a good life. Successfully living a good life is **eudaimonia**, 'living well and faring well'. Eudaimonia is a **final end**, something we seek for its own sake, not in order to achieve something further. Some philosophers have translated eudaimonia as 'happiness', but it isn't happiness in the usual sense of that word. We tend to think of happiness as a psychological state, but eudaimonia is about how one lives. And happiness is something that comes and goes, while eudaimonia is a feature of a person's life as a whole.

*'Function' and 'virtue'*
We achieve eudaimonia by achieving our purpose or aim in life, our telos. Aristotle believes that everything has a telos, and the telos of anything is related to its 'function' (**ergon**). When a thing functions well, it achieves its telos. Everything, he thinks, has an ideal form of functioning, and all things fit together into a harmonious whole. For example, the ergon of an eye is to see, and bodily parts work together to preserve the life of the organism. We can also apply this to artefacts, things we make in order to perform a certain function, e.g. the ergon of a knife is to cut.

Understanding something's function gives an understanding of its nature, and so also what it is for that thing to be 'good'. A good knife cuts well; a good eye sees well. An **arête** is a quality that aids the fulfilment of a thing's ergon. It can be translated generally as an 'excellence' or a '**virtue**'. So sharpness is a virtue in a knife designed to cut. Good focus is a virtue in an eye.

Aristotle applies this model to species. A good plant – an ideal specimen, we might say – is one that flourishes: it grows well, produces flowers well, etc., according to its species. Because Aristotle thinks that everything fits together, he thinks that we can identify the 'function' of a species in terms of what is characteristic about it. What does this species do that is distinctive, that gives it its particular place in nature?

The idea that human beings have a 'function' sounds unusual, and certainly we don't share Aristotle's views on biology. But if we don't have a 'function', then we can't use our function to identify the good life, and Aristotle's ethical theory fails before it has started. But modern

virtue ethicists have argued that 'function' is better understood in relation to 'functioning' rather than 'purpose'. So the idea of ergon can be understood as 'characteristic form of activity' of something.

What is characteristic of human beings is that we are rational animals – whatever we do, we do for reasons. Our ergon, then, is living as a rational animal, and the virtues of a human being will be what enable us to do this. To live well, we must be guided by the 'right' reasons – good reasons, not 'bad' reasons. Eudaimonia is living a life in accordance with reasons through exercising virtues.

Aristotle allows that being fortunate can also play a small role in leading a good life. Fortunes change, but living virtuously has a much greater permanence. A virtuous person deals with bad fortune in the best possible way, so only very rarely and through terrible circumstances, can someone virtuous fail to lead a good life.

*Virtues*
Virtues are traits that enable us to live rationally. Aristotle identifies virtues of character and virtues of the intellect.

Character traits are dispositions that relate to what, in different circumstances, we feel, how we think, how we react, the sorts of choices we make, and the actions we perform. They relate to our desires and emotions – what Aristotle calls 'passions'. So someone is short-tempered if they are disposed to feel angry quickly and often; quick-witted if they can think on their feet; and so on. Some traits of character, such as being short-tempered or greedy, stop us from leading a good life – these are **vices**. Other traits of character, such as being kind or courageous, help us to lead a good life – and these are virtues. A virtue of character is a disposition to feel, desire and choose 'well', which is necessary if we are to achieve eudaimonia.

**Practical wisdom** (*phronesis*) is an intellectual virtue, a virtue of practical reasoning. (**Theoretical reason** investigates what we can't change and aims at the truth – as in discovering the world through scientific investigation. **Practical reason** investigates what we can change and aims at making good choices.) Reasoning about what we can change is deliberation, so practical reason is expressed in deliberation. The person with practical wisdom deliberates well about how to live a good life. This involves

1.    a general conception of what is good or bad, related to the conditions for human flourishing;
2.    the ability to perceive, in light of that general conception, what is required in terms of feeling, choice and action in a particular situation;
3.    the ability to deliberate well; and
4.    the ability to act on that deliberation.

To make good choices, not only must our reasoning be correct, but we must also have the right desires. That means we will need virtues of character as well.

*The doctrine of the mean*

The question that faces us on any occasion is how to achieve what is good – what the good life involves – in this situation. Aristotle compares living well with other activities, such as eating well or physical training. In these cases, we need to avoid too much or too little food or exercise. We achieve health and physical fitness by following an 'intermediate' course of action, which Aristotle calls the 'mean'. This is differs from person to person. For example, a professional sportsman needs more food and exercise than most people.

Now, in the 'art of living', something similar applies. We can feel our passions either 'too much' or 'too little'. Some people feel angry too often, over too many things (perhaps they take a critical comment as an insult), or maybe whenever they get angry, they get very angry, even at minor things. Other people feel angry not often enough (perhaps they don't understand how people take advantage of them). To be virtuous, says Aristotle, is 'to feel [passions] at the right times, with reference to the right objects, towards the right people, with the right **motive**, and in the right way'.

This is Aristotle's '**doctrine of the mean**'. There better and worse ways to feel passions, and the right way is determined by reason. If we feel our passions 'irrationally' – at the wrong times, towards the wrong objects, etc. – then we don't live well. Likewise, we can choose the right or wrong actions and act for the right or wrong reasons, usually as a result of whether the feelings that influence our choices are themselves rational or irrational. Given the very close connection between what we feel and how we choose to act, virtues are dispositions of choice as well, and there is a 'mean' for actions as well as for feelings.

Aristotle's doctrine of the mean does not claim that when we get angry, we should only ever be 'moderately' angry. We should be as angry as the situation demands, which can be very angry or only slightly irritated.

We give names to character traits that involve 'too much', 'too little' and, in the middle, the virtue. For example, someone who feels fear 'too much' is cowardly. Someone who feels fear 'too little' is rash. Someone who has the virtue relating to fear is courageous.

| Passion/concern | Vice of deficiency | Virtue | Vice of excess |
|---|---|---|---|
| Fear | Cowardice | Courage | Rashness |
| Pleasure/pain | 'Insensitivity' | Temperance/self-control | Self-indulgence |
| Giving money | Stinginess | Generous | Profligacy |
| Shame | Shamelessness | Modesty | Shyness |
| Anger | 'Unirascible' | Good temper | Short temper |
| Small honours | Lack of ambition | Ambition | Overambition |
| Humour | Boorishness | Wittiness | Buffoonery |
| Social relationships | Surliness | Friendliness | Obsequiousness |

What the right action, time, object, person and so on is, for both feeling and action, practical wisdom helps us to know. However, Aristotle says that there are no rules for applying knowledge of the good life to the current situation. What is right can vary from one occasion to another. This does not make ethics subjective, as we still need to discover what

*is* right on each occasion. But it does mean that practical wisdom cannot be taught 'theoretically'. Rather, it is something that has to be acquired through experience.

*Becoming a good person*
Aristotle argues that we acquire virtues of character through 'habit', in particular, the habits we form in childhood. We come to form dispositions to feel and behave in certain ways by what we do. We can understand this by analogy with acquiring practical skills, such as carpentry, cookery or playing a musical instrument. You cannot learn cookery or a musical instrument just by studying the theory, by merely acquiring knowledge about how to cook or play; you have to practice the activity. Likewise, being told how to be good is not enough to become good; you have to actually practice being good. This is known as 'the **skill analogy**'.

In order to become just, Aristotle says, we have to do just acts. But how can we do just acts unless we are already just? We need to distinguish between actions which are 'in accordance with' justice and just acts, properly so called. The actions that we do when learning to become just are acts in accordance with justice. A child may do what is just (such as not taking more than its fair share) because it is told to do so; or because it likes the person it is sharing with; or because it wants to please an adult, and so on. It neither truly understands what justice is nor does it choose the act because the act is just. But a fully just act is an act that is not only in accordance with justice, but also done as the just person does it. This requires the right motivation, but also an understanding of justice and why it is good. We need practical wisdom for this.

But, we said, practical wisdom also depends on virtue, because it is possible to deliberate from the wrong ends and what appears good to someone (what appeals to them) depends on their character traits. So on Aristotle's theory, we become both good and practically wise together.

*Issues for Aristotle's virtue ethics*
1. Does Aristotle's virtue ethics provide guidance about how to act? The doctrine of the mean does not look like much help. First, 'too much' and 'too little' aren't quantities on a single scale. The list of 'right time, right object, right person, right motive, right way' shows that things are much more complicated than that. Second, it doesn't tell us, for example, how often we should get angry, and how angry we should get. Just about anything could be 'in the mean' if the circumstances were right!

But Aristotle didn't intend the doctrine of the mean to be a rule or an algorithm for what to do. We can't 'figure out' what it is right to do by applying rules; we must have practical wisdom, which involves 'seeing' what to do. This requires virtues of character and lots of experience.

This implies that someone who isn't virtuous doesn't know the right thing to do. Aristotle accepts this. However, knowledge of what is good comes in degrees, and we can improve or destroy our ability to know what is good by the kind of character we develop. If someone has a completely depraved character, perhaps they really don't know what is good or bad. But most people will have enough understanding of the good to make moral decisions.

Furthermore, people can improve their knowledge of what is good by becoming more virtuous.

2. Can virtues conflict? For example, could acting with justice conflict with being merciful? Aristotle denies this. With practical wisdom, if we are sensitive to the context, we can discover how to satisfy both justice and mercy. For example, someone may appeal to their difficult circumstances as a reason for lenient punishment. Taking this into consideration isn't only merciful – it is also important for justice, for understanding just how wrong their action was. But can this approach work in every case? For example, could loyalty to a friend ever require you to be dishonest towards someone else?

3. Does Aristotle's account rest on circular reasoning? He seems to say that:

a.   an act is virtuous if it is an act that would be done by a virtuous person in this situation;
b.   a virtuous person is a person who is disposed to do virtuous acts.

This looks like a circular definition – we have defined a virtuous act in terms of a virtuous person and a virtuous person in terms of virtuous acts.

But the objection is mistaken. A virtuous person has the virtues, which are traits that enable them to achieve eudaimonia. Character traits relate to our choices and actions, but they are equally concerned with our passions. And eudaimonia is defined not in terms of virtuous actions, but in terms of feeling, thinking and choosing. A virtuous person also has practical wisdom, so they also understand why certain actions are right. So while (a) is correct, (b) is too simple. With the right definition, the circularity is avoided.

4. Aristotle doesn't draw a distinction between a life that is good for me and a life that is morally good. Is leading a morally good life the same as eudaimonia? Can self-interest and what is morally required come into conflict? Can there be virtues that are not in our self-interest? Suppose a woman works as a medic in a foreign country, ceaselessly saving lives and relieving suffering, often far from civilization and under difficult circumstances. She is often ill and tired, and doesn't experience joy in her work, just the conviction that it is needed. She dies prematurely from a virus. Her life demonstrates many virtues – kindness, compassion, generosity, perhaps justice – but she doesn't 'flourish'. Or again, if someone lives under an unjust dictatorship, courage and justice may lead them to stand up for what is right but be imprisoned or killed as a result. (These examples are from Christine Swanton, *Virtue Ethics.*)

One response is that we can't say what eudaimonia is before identifying virtues. We can't say what it would be to live 'a good life for me' and then investigate whether having virtues would enable this life or not. Instead, living in accordance with the virtues gives us the best account of what *counts* as a flourishing life. Someone who aims to live in a fair, generous, courageous, just way is committed to other people's well-being, but having these commitments is part of what it is to flourish. To live this best life, you must treat people in certain, morally good ways.

Does this answer the objection? Surely the woman medic's life good could have gone better for her, so she didn't flourish. So perhaps virtues can pull in different directions, and flourishing in one sense may lead to not flourishing in another sense.

B. Kant's deontological ethics

Aristotle focused on what it is to be a good person and then explained the right thing to do as what the good person does. But we can reverse this to say that someone is good if they do what is right. Rather than focusing first on the *agent*, we focus on the *action*. For example, morality could be a set of *rules* that we should follow. Kant's deontological ethics appeals to reason to discover these rules which define our duties.

*The good will and duty*
Actions are the result of choices. Choices are made on the basis of an intention – what you intend to achieve (your purpose) and your reason or motive for acting. The intention defines the action. For example, if someone kills someone else, it makes a big difference whether they intended to kill them or if the person was only intending to defend themselves against an attack or if the killing was accidental in some other way.

Kant believed that, whenever we make a decision, we act on a **maxim**. Maxims are Kant's version of intentions. They are our personal principles that guide our decisions; e.g. 'to have as much fun as possible', 'to marry only someone I truly love'. Kant talks of our ability to make decisions as '**the will**'. He assumes that our wills are rational; that is we make choices on the basis of reasons. We do not act without thought on instinct or desires.

But what should guide our choices? Kant argues that any end we aim at can, in some circumstances, be bad. For instance, knowledge and self-control are good – but they can enable someone to do clever or difficult bad things, if that is what they choose. Happiness can be bad as well, e.g. if it comes from hurting others. Kant argues that the only thing that is good 'without qualification' is the '**good will**'.

The good will is motivated by duty. Suppose a shopkeeper sells his goods at a fixed price, giving the correct change, and acting honestly in this way. This is the morally right thing to do, he shouldn't cheat people. But if all he cares about is keeping his customers, then even though he does the right thing, he does it because it will benefit his business, not because it is the morally right thing to do. The shopkeeper is *acting in accordance with duty*. To act in accordance with duty is simply to do what is morally right, whatever one's motive for doing so. But the shopkeeper isn't *motivated by duty*, i.e. he doesn't act from or out of duty. To act out of duty is to do what is morally right because it is morally right.

To have a good will is to do one's duty because it is one's duty. But what is our duty? It can't be to aim at some end, say happiness, because no end, even happiness, is always morally good. So the good will must, in some way, be good 'in itself', just on the basis of what it is like as a will. How?

Maxims are principles of choice. They are subjective – you have yours, I have mine. What makes them different is what they are about, what they aim at and why. But what they have in common is that they are all principles. Now, Kant assumes that morality is a set of principles that are the same for everyone and that apply to everyone. So the concept of duty is the concept of a principle for everyone. So, somehow, the good will is a will that makes choices on the basis of a principle for everyone.

*The Categorical Imperative*
Kant has arrived at the principle, 'Act only on that maxim through which you can at the same time will that it should become a universal law'. To have a good will, I should act only on maxims that I can also will everyone to act on. Suppose I am tempted to make a promise with no intention of keeping it; e.g. I might borrow money (because I want the money) on the promise to pay it back, but I don't intend to pay it back. We can show that this is wrong. Suppose everyone acted on this maxim. Then everyone would know that everyone acts on this maxim. In that situation, making a false promise like this would be impossible. No one would trust my promise, and I can't make a promise unless someone believes it. So I can't will my maxim to be universal – I can't **universalise** it. And this shows that I have a moral duty not to make false promises.

Kant calls his principle the '**Categorical Imperative**'. An '**imperative**' is a command, a statement of what one ought to do. A 'hypothetical imperative' is a statement about what you ought to do assuming some desire or goal. It specifies a **means** to an end, e.g. 'if you want to see the show, you ought to get to the theatre at least 15 minutes early' or 'Eat at least five portions of fruit and vegetables a day (if you want to stay healthy)'. You can avoid the imperative if you give up your goal. Suppose I don't want to see the show – then I don't need to get to the theatre early. In other words, it is possible to 'opt out' of a hypothetical imperative. But we can't 'opt out' of morality. Moral duties are what we ought to do, full stop, regardless of what we want. So they are 'categorical'. All categorical imperatives – our moral duties – are derived from *the* Categorical Imperative.

*Contradiction in conception and contradiction in will*
There are two ways in which we could fail to be able to will our maxim to become a universal law.

'**Contradiction in conception**': Suppose you want something you can't afford, so you steal it. Your maxim is something like: 'To steal something I want if I can't afford it'. This can only be the right thing to do if everyone could do it. However, if we all just help ourselves to whatever we wanted, the idea of 'owning' things disappears. Now, by definition, you can't steal something unless it is owned by someone else. But people can only own things if we don't all go around helping ourselves whenever we want. So it is logically impossible for everyone to steal things. In other words, it is inconceivable – a contradiction in conception – for everyone to steal things. So we can't rationally will 'To steal something I want if I can't afford it' as a universal law. And so stealing is wrong.

'**Contradiction in will**': Consider a refusal to help other people, ever. It is logically possible to universalise the maxim 'not to help others in need'. The world would not be a pleasant place, but this is beside the point. It is not self-contradictory to conceive of this scenario. But Kant argues we cannot *will* it.

P1.   A will, by definition, wills its ends.
P2.   To truly will the ends, one must will the necessary means. (Willing involves not just desiring, but choosing to act.)

C1. Therefore, we cannot rationally will a situation in which it would be impossible for us to achieve our ends. To do so is to cease to will the necessary means to one's ends, which is effectively to cease to will any ends at all. This contradicts the very act of willing.

P3. It is possible that the only available means to our ends, in some situations, involves the help of others.

C2. We cannot therefore will that this possibility is denied to us.

C3. Therefore, we cannot will a situation in which no one ever helps anyone else.

Kant does not claim that an action is wrong because we wouldn't *like* the consequences if everyone did it (many philosophers and students have misinterpreted Kant on this point). His test is whether we can rationally will that our maxim be a universal law. Willing and wanting (or liking) are different. Someone can want something that they don't will – they choose not to act on their desire, e.g. such as cheating on their husband or wife with someone they find very attractive. And someone can will something they don't want, such as going to the dentist for surgery. Kant is concerned with willing not wanting.

As the contradiction in conception and contradiction in will show, disobeying the Categorical Imperative involves a self-contradiction, according to Kant. It is not just morally wrong to disobey the Categorical Imperative, it is also irrational. Morality and rationality are categorical; the demands to be rational and moral apply to everyone and don't stop applying to you even if you don't care about them. Neither morality nor rationality depend on what we want.

*The second formulation of the Categorical Imperative*
Kant gives a second formulation of the Categorical Imperative, known as the **Formula of Humanity**: 'Act in such a way that you always treat humanity, whether in your own person or in the person of any other, never simply as a means, but always at the same time as an end'. By 'humanity', Kant means our ability to rationally choose which ends to adopt and pursue. As rational wills – as the only thing which can be unconditionally good – we have 'intrinsic worth', which Kant identifies as 'dignity'.

To treat someone simply as a means, and not also as an end, is to treat the person in a way that undermines their power of making a rational choice themselves. So to treat someone as an end implies, first, that we should appeal to other people's reason in discussing with them what to do, rather than manipulating them in ways they are unaware of. Coercing someone, lying to them or stealing from them all involve not allowing them to make an informed choice. Second, treating someone as an end implies leaving them free to pursue the ends that they adopt. So we should refrain from harming or hindering them. Third, we should help them pursue their ends, just as we pursue our own ends. In other words, we should help other people.

*Issues for Kant's deontological ethics*
1. Is the Categorical Imperative the right test for whether an action is right? Or could there be maxims that *pass* the test which it is *wrong* to act upon? Suppose I phrase my maxim cleverly. In the example of stealing above, I could claim that my maxim is 'To steal from large shops and when there are seven letters in my name (Michael)'. Universalising this

maxim, only people with seven letters in their name would steal and only from large shops. This would not lead to a general breakdown in the concept of private property. So it is possible for this law to apply to everyone – it passes the test. That doesn't make it right.

Kant's response is that his theory is concerned with my actual maxim, not some made-up one. It is not actually part of my choice that my name has seven letters, as though I wouldn't try to steal if my name was different. For Kant's test to work, we must be honest with ourselves about what our maxims are.

Could there be maxims that *do not pass* the test, but it is *right* to act on it? Suppose someone comes to your house to hide from someone who wants to murder them. After they have hidden, the would-be murderer arrives and asks you where they are. Is lying in this situation right? Kant says not. Lying is always wrong, because we cannot universalise the maxim to deceive people. Doesn't this show his theory is wrong? But perhaps Kant is wrong about our maxim in this case. Perhaps the maxim is 'to save a life by lying', and this maxim can be universalised. In that case, Kant's theory is right, even if his own application of it was too strict.

Another example: I am a hard-working shop assistant, who hates the work. One happy Saturday I win the lottery, and I vow 'never to sell anything to anyone again, but only ever to buy'. This cannot be universalised. If no one ever sold things, how could anyone buy them? It is logically impossible, which makes it wrong according to Kant's test. But while the maxim is perhaps eccentric, it doesn't seem morally wrong.

2. Can our duties can conflict? Could I be faced with a situation where I must either break a promise or tell a lie? If so, then Kant's theory implies that whatever I do is wrong. Aristotle argues that there are no strict rules in ethics, and if we interpret the situation correctly, we can resolve the conflict. But Kant argues that our moral duties do not allow exceptions. Nothing can override a moral duty, because it is categorical. Yet he argues, like Aristotle, that if there appears to be a conflict, we have misunderstood what at least one duty requires of us. Is this convincing?

3. Do consequences matter? Is murder bad because life is valuable? If so, we should try to ensure that there are as few murders as possible. Suppose that unless I kill someone deliberately, many people will die. Shouldn't I prevent more deaths? In this case, I have a duty to kill, because I would be killing in order to save lives.

Kant rejects this approach, because it claims that what makes a will good is that it wills good ends (fewer deaths). But there are no ends that are good without qualification. So we can't analyse what a good person does in terms of the consequences of their action, because no consequences, on their own, are always good or bad. When thinking about morality, we are not thinking about the means to a good end (hypothetical imperatives), but how everyone could act (universalizability).

4. Are any motives, apart from duty, morally good? Consider this example (from Michael Stocker): A woman visits a friend in hospital. The friend thanks her. She replies, 'I was just doing my duty'. If true, her friend can object 'what about loving me?' Kant seems to say that

we should benefit people because it is our duty to so, not because we like them. But surely, if I do something nice for you because I like you, that is a morally good action. Kant's theory doesn't seem to recognise and explain the moral worth of love and friendship.

However, Kant can allow us to be motivated by our feelings. His point is that, when we are choosing what to do, how we feel should not *decide* the matter; our intention to do what is morally right should. To be morally good, you only need to be willing to refuse to help your friend if that involved doing something morally wrong.

5. Are there any real categorical imperatives? Categorical imperatives present the action as some something you should do, independent of what you want. Acting on them is (supposedly) a matter of being rational rather than fulfilling a desire. Philippa Foot rejects this.

It is true that moral judgments are not hypothetical imperatives in Kant's sense. However, think of club rules. In Foot's (now old-fashioned) example, if the club rules say, 'Do not take ladies into the smoking room', there isn't a hidden assumption 'if you want to remain part of the club'. Suppose someone doesn't want to remain part of the club, thinking it fusty and sexist, and he will quit tomorrow for good. Is he now allowed to take ladies into the smoking room? No. In these examples, we don't withdraw the 'should' depending on what someone wants. The club rule is a *non-hypothetical* imperative, but it is not unconditional in the sense that Kant thinks moral judgments are. If you don't like the club rules, don't join.

Foot argues that moral judgments are non-hypothetical imperatives, but Kant does not prove that they are categorical in his sense. To act irrationally, claims Foot, is simply not to take the right means to your ends. An immoral person can still act rationally. If we feel that morality is categorical in some deeper sense, this is because of how morality is taught – we feel that we 'must do' what is morally right, whatever our desires.

C. Utilitarianism

At the same time that Kant developed his deontological theory, Jeremy Bentham developed his theory of utilitarianism as an **empirical** approach to morality. We can know empirically that people want to be happy, and we can work out empirically what will make them happy. To do the right thing, we should consider the consequences of the different actions we could perform and choose that action that brings about, or is likely to bring about, the greatest happiness. This makes complicated decisions easy and avoids appeals to controversial moral intuitions.

*Bentham's quantitative hedonistic utilitarianism*
In its simplest form, utilitarianism is defined by three claims.

1. Actions are morally right or wrong depending on their consequences and nothing else. An act is right if it maximises what is good. This is '**act consequentialism**'.
2. The only thing that is good is happiness, understood as pleasure and the absence of pain. This is '**hedonism**'.
3. No one's happiness counts more than anyone else's. This is a commitment to equality.

This is known as **hedonistic act utilitarianism**. An action is right if it maximises happiness, i.e. if it leads to the greatest happiness of all those it affects. This is the '**principle of utility**'. The theory is called 'utilitarianism' because it maximises '**utility**', which Bentham defines as the property of some object or action to produce 'benefit, advantage, pleasure, good, or happiness'. Happiness, he claimed, was pleasure and the absence of pain.

Bentham argued that we can measure pleasures and pains and add them up on a single scale by a process he called the '**felicific calculus**' ('felicity' means happiness), also known as the 'hedonic calculus'. If a pleasure is more intense, will last longer, is more certain to occur, will happen sooner rather than later, or will produce in turn many other pleasures and few pains, it counts for more. In thinking what to do, we also need to take into account how many people will be affected (the more we affect positively, and the fewer we affect negatively, the better). The total amount of happiness produced is the sum total of everyone's pleasures produced minus the sum total of everyone's pains. As this demonstrates, Bentham took a quantitative approach to happiness.

*Mill's qualitative hedonistic utilitarianism*
John Stuart Mill thought that Bentham's felicific calculus was too simple in how it thought about happiness. Some types of pleasure, he said, are 'higher' than others, more valuable, more important to human happiness, given the types of creatures we are and what we are capable of. If (almost) everyone who has experience of two types of pleasure prefers one type to the other, even if that type of pleasure brings more pain with it, then the type that they prefer must be more valuable – or why else do they prefer it? As long as our physical needs are met, we prefer the pleasures of thought, feeling and imagination to pleasures of the body and the senses, even though our 'higher' capacities also mean we can experience terrible pain, boredom and dissatisfaction. ''Tis better to have loved and lost than never to have loved at all' (Tennyson, 'In Memoriam A.H.H.'). We can say the same about intelligence and artistic creativity – better to have the pleasures that they bring, even though they cause

us pain and distress, than to be unintelligent or lack creativity. We wouldn't choose to be a well-looked-after pig, rather than a dissatisfied human being. This preference, Mill thinks, derives from our sense of dignity, which is an essential part of our happiness.

In introducing this distinction between **higher and lower pleasures**, Mill rejects Bentham's felicific calculus. We can't simply 'add up' how much happiness some action produces, because happiness has an element of quality, not just quantity, of pleasure.

*Mill's 'proof' of utilitarianism*
Utilitarianism claims that happiness is the only good. What is good, says Mill, is what we should aim at in our actions and lives. What we should aim at is what is desirable. People in general desire happiness. This looks like good evidence that happiness is desirable (unless people in general desire what is not worth desiring). Everyone wants happiness, so it is reasonable to infer that happiness is desirable (good).

Of course, we each desire our own happiness, while utilitarianism says that we should aim at the general happiness – everyone's happiness is good. However, this follows from each person's happiness being good plus impartiality. Each person takes their own happiness to be good, and so, adding each person's happiness to that of others, the happiness of everyone is good for people in general. (Mill here assumes that ethics is concerned with what is good in general.)

But is happiness the *only* good? What about truth, beauty, freedom, etc.? People desire all these, so going by the evidence, many different things are good. Mill responds that happiness has many 'ingredients', and each ingredient is desirable in itself. Think about having a good holiday. Suppose you have to get up very early in order to catch the plane. You do this in order to have a good holiday, but it isn't part of having a good holiday. Later on, you are lying on the beach in the sun, listening to your favourite music. Are you doing this 'in order' to have a good holiday? Yes, but not in the same sense. This just is having a good holiday at the moment. In these circumstances, here and now, it is what 'having a good holiday' amounts to. The same applies to happiness, Mill argues. For example, when someone desires to know the truth 'for its own sake', their knowing the truth doesn't cause their happiness as some further and separate thing. Rather, in this situation, their happiness consists in their knowing the truth. Knowing the truth for its own sake is part of happiness for them. So, Mill claims, whatever we desire for its own sake is part of what happiness is for us.

*Pleasure and preferences*
In *Anarchy, State and Utopia*, Robert Nozick asks us to imagine being offered the chance to plug into a virtual reality machine which will produce the experience of a very happy life. If we plug in, we will not know that we are in a virtual reality machine. We will believe that what we experience is reality. However, we must agree to plug in for life or not at all. Nozick argues that most of us would *not* plug in. We value being in contact with reality, even if that makes us less happy. But we can't understand this in terms of the 'pleasure' of being in touch with reality, or a preference for certain types of pleasure over others, because if we were in the machine, we would still experience this pleasure (we would believe we were in

touch with reality). Instead, what we want is not a *psychological state* at all; it is a *relation* to something outside our minds.

**Preference utilitarianism** argues that what we should maximise is not pleasure, but the *satisfaction of people's preferences (desires)*. Why?

1. If Nozick is right, we prefer to be in touch with reality. Having this preference satisfied is valuable. For a preference to be satisfied, it must be satisfied in reality. It is not enough that the person thinks their preference is satisfied.
2. We can also appeal to preferences to explain Mill's claims about higher and lower pleasures. Rather than talk about the quality of types of pleasure, we could argue that whatever people prefer is of more value to them.
3. The satisfaction of many of our preferences will bring us pleasure, but many will not. For instance, Bentham and Mill do not distinguish between producing happiness and decreasing pain. But are these morally equivalent? If people more strongly prefer not to suffer pain than to be brought pleasure, then that would explain the thought that it is more important not to cause harm. Or again, we can also argue that people have preferences about what happens after their death, e.g. to their possessions, and it is important to satisfy these as well, even though this cannot bring them any pleasure.

Preference utilitarians can continue to claim that happiness is the only good but understanding happiness as one's desires being satisfied, rather than pleasure.

*Rule utilitarianism*
We have seen that act utilitarianism judges the rightness of an act by its consequences. This leads to the view that morality isn't about following rules, as Kant thought, but about doing whatever *in this situation* would create the greatest happiness. **Rule utilitarianism** combines the ideas that happiness is the only good and that morality is about rules. It claims that

1. an action is right if it complies with the right rules, and
2. a rule is morally right if people following it creates the greatest happiness (compared to an alternative rule).

In other words, rule utilitarians do not look at the *consequences of individual actions*, but at the *consequences of people following a rule*. So, for example, if people generally following the rule 'Don't tell lies' maximises happiness, it is wrong to tell a lie – even when telling a lie would *in this situation* create more happiness than telling the truth.

This raises an objection. The point of the rules is to bring about the greatest happiness. If there is a situation in which breaking the rule will lead to more happiness than following the rule, why follow the rule? If happiness is the only good and I know that lying in a particular situation will produce more happiness than telling the truth, why tell the truth, causing unhappiness?

Rule utilitarians could respond by saying that we should amend the rule to allow the exception in such cases. For example, 'Don't lie' should become 'Don't lie unless telling the

truth will hurt someone's feelings'. However, life is complicated. Whenever a particular action causes more happiness by breaking the rule than by following it, we should do that action. But if we keep amending the rules like this, we will end up with just one rule, namely to maximise the greatest happiness. So we should be act utilitarians.

Rule utilitarians respond that morality should be understood as a set of rules. Morality needs to provide general guidance, a way of thinking that people can remember and rely on. Since we need rules, the aim of these rules should be to maximise happiness. And so actions are right when they follow a rule that maximises happiness overall – even when the action itself doesn't maximise happiness in this particular situation. The advantages of this approach over act utilitarianism emerge as we consider issues for utilitarianism.

*Issues for utilitarianism*
1. Can we work out the consequences of an action to know how much happiness it will produce? Bentham's felicific calculus shows how challenging this is – for each possible pleasure from each possible action for each person who may be affected, we have to figure out its intensity, duration, etc. It can't be done.

Mill agrees, but responds that we don't need to try. Through history, we have, through trial and error, worked out which actions tend to produce happiness, and this guidance forms our inherited moral rules. 'Be honest', 'don't steal', etc. are embodiments of the wisdom of humanity that lying, theft etc. tend to lead to unhappiness. Only if these 'secondary principles' conflict do we need to apply the primary principle of utility directly to an action. Rule utilitarianism claims that this shows the importance of rules for morality.

It is important to note that utilitarianism says that happiness – not just *human* happiness – is good. As Peter Singer argues in *Animal Liberation*, there is nothing in the theory that gives us a reason to privilege human happiness over the happiness of non-human animals. Just as treating men better than women is sexism, treating humans better than animals (as regards happiness) is '**speciesism**'. We need to improve on our inherited morality here, since it allows animals to be treated as tools for making human lives happier, e.g. as food.

2. Can utilitarianism respect justice – issues of individual liberty, rights and fairness? Act utilitarianism does not rule out any type of action as immoral. Suppose a group of child abusers find and torture abandoned children. Only the child suffers pain (no one else knows about their activities). But the abusers derive a great deal of happiness. So more happiness is produced by torturing the child than not. This doesn't make it morally right!

Many rights involve restrictions on how people can treat each other. For instance, I have a right that other people don't kill me (the right to life). One purpose of rights is to protect individual freedom and interests, even when violating that freedom would produce greater happiness.

Some utilitarians argue that we have no rights. As long as we consider situations *realistically*, then whatever brings about the greatest happiness is the right thing to do. In the case of the tortured child, in real life, other people *would* find out and become very upset. So these actions wouldn't lead to the greatest happiness.

Mill argues that we do have rights. Rights are rules that prohibit harm and protect our freedom, and these interests are 'extraordinarily important' to us. Having these rights contributes most to happiness in the long term. This leads to rule utilitarianism again: rules requiring fairness and justice will produce greater happiness in the long run than rules that do not.

3. Does utilitarianism demand complete impartiality? Many of the things that we do to make people happy are aimed at our family and friends. But according to act utilitarianism, we should consider the happiness of each person equally when deciding what to do. In other words, we should be impartial. It seems we should spend much less time and money on people we love and more on helping people in need.

We can argue that this makes utilitarianism too idealistic. But in addition, it misses something important about friendship. Friendship requires that the friend is valued as the individual person that they are, and that we act out of love for them. Doing something for a friend is morally good, not wrong because it fails to be impartial. Furthermore, attachments of love and friendship are central to our happiness, but they are at odds with utilitarianism, because they are not impartial between everyone's happiness. Utilitarianism fails to recognise the moral importance of partiality.

Rule utilitarians respond a rule that allows partiality to ourselves, our family and friends will create more happiness than a rule that requires us to be impartial all the time. This explains the moral importance of partial relationships – they are necessary to happiness. Of course, we shouldn't be completely partial. We still need to consider the general happiness, but we only need to act in such a way that, if everyone acted like that, would promote the greatest happiness. For example, in the case of charity, I only need to give as much to charity as would be a 'fair share' of the amount needed to really help other people. This combination of partiality and impartiality respects both our natural inclinations and the demands of morality.

4. Does utilitarianism recognise integrity? Having integrity involves acting according to your own values, sticking to them especially in the face of temptation or other situations that would make it easier to do something you consider wrong. Imagine George, a chemist who is looking for a job (the example is from Bernard Williams, *Utilitarianism*). He isn't having much luck, and with small children to look after, the situation is causing a lot of stress. He is offered a job in a laboratory that develops chemical weapons. George is strongly opposed to chemical warfare. But refusing the job will leave the vacancy for someone else who may be more enthusiastic and so develop more chemical weapons faster and more effectively than George will – leading to more harm in the world. Utilitarianism says that George should take the job. That George would be made very unhappy by taking the job is outweighed by the prevention of significant unhappiness to others. So utilitarianism undermines our integrity.

The rule utilitarian might reply by saying that we should have a rule that allows people not to follow other rules if they conflict with one's integrity.

5. Can act utilitarianism recognise the moral value of our intentions? Whether someone intends to harm us or not makes a big difference to how we respond to their action, whether or not they succeed. Trying to harm someone and failing – so they are unharmed – is (usually) blameworthy; trying not to harm someone and failing – so they are accidentally harmed – is not. How can this be if all that matters are consequences, not intentions?

Mill responds that intentions are irrelevant to whether an action is right; however, they are relevant to whether a person is good. An intention that tends to produce morally wrong actions, such as intending to harm someone, is itself a bad intention, while intentions to produce happiness are good intentions. And someone with bad intentions is a bad person. But we need to separate the judgment of whether an action was right or wrong from the judgment of whether the intention was good or bad.

Rule utilitarianism explains the importance of intentions in relation to moral rules. Rules prohibit or require certain types of action, and these differ depending on the intention, e.g. whether a killing is murder or manslaughter or accidental death. A rule prohibiting murder is not a rule prohibiting self-defence, even if self-defence may lead to death. So whether an action complies with a rule or not depends, at least in part, on the individual's intention in performing the act. Two actions may have the same consequences, and yet one is right and one is wrong, because they are different types of action and fall under different moral rules, because the person's intention in each was different. And so rule utilitarianism can recognise that the intention does matter.

1. What does Aristotle mean by saying that eudaimonia is living a life in accordance with reason? What is his argument for this claim?
2. What is Aristotle's doctrine of the mean? Does it help in real life?
3. Do all virtues contribute to the virtuous person's eudaimonia?
4. Can moral virtues or rules ever conflict? If so, what should we do when they do?
5. Why does Kant argue that we should only act on maxims that we can will as universal laws?
6. What is involved in treating someone as an end in themselves? Are there times when it is permissible not to treat someone this way?
7. Are moral duties categorical?
8. Explain the similarities and differences between Bentham and Mill on happiness.
9. Is happiness always good?
10. Can act utilitarianism be successfully defended? Does rule utilitarianism do any better?

*Further reading*

Bennett, C. *What is This Thing Called Ethics?*, 2nd ed., Routledge 2015 [see Ch. 4, 5, 6 for accessible introductions to the three normative theories]

Rachels, J. *The Elements of Moral Philosophy*, McGraw-Hill 2014 [a favourite of introductory courses, now in its eighth edition; see Ch. 7, 8, 9, 10, 12]

Blackburn, S. *Being Good*, OUP 2001 [a lively, engaging, and quick introduction to the main issues of ethics, written in a very clear style; see Ch. 12, 17, 18]

Norman, R. *The Moral Philosophers*, Clarendon Press 1983 [see Ch. 2, 5, 7 and 11 for one of the best introductions to normative theories of ethics via the great philosophers who held them]

Cave, P. *Ethics*, Oneworld 2015 [see Ch. 2, 3, 4 for stimulating and entertaining accounts of normative theories]

Hughes, G. *Aristotle's Nicomachean Ethics*, Routledge 2013 [a clear, detailed, authoritative discussion of Aristotle's virtue theory]

Annas, Julia (2006), 'Virtue Ethics', in David Copp (ed.), *The Oxford Handbook of Ethical Theory*, Oxford: Oxford University Press, pp. 515–536 [an excellent overview of virtue theory, focusing on Aristotle but commenting on developments]

Scruton, R. *Kant: a Very Short Introduction*, OUP 2001 [Ch. 5 is a good introduction to the Categorical Imperative and the challenges it faces]

Crisp. R. *Mill on Utilitarianism*, Routledge 1997 [the best discussion of Bentham and Mill's utilitarianism]

Smart, J. & Williams, B. *Utilitarianism*: *For and Against*, CUP 1973 [one of the most informative discussions of utilitarianism - its structure, its motives, and its problems]

*Internet resources*
Aristotle
Oxford University philosophers discuss Aristotle's ethics:
   http://podcasts.ox.ac.uk/what-virtue-ethics
   http://philosophybites.com/2009/04/terence-irwin-on-aristotles-ethics.html
   http://philosophybites.com/2008/10/roger-crisp-on.html
   http://philosophybites.com/2007/11/miles-burnyeat.html

Kant

Johnson, Robert and Cureton, Adam (2018), 'Kant's Moral Philosophy', *Stanford Encyclopedia of Philosophy* <https://plato.stanford.edu/archives/spr2018/entries/kant-moral/> [difficult and detailed but gets to Kant's underlying motivations]

A two-part accessible introduction to Kant's ethics from teachphilosophy (Prof Paul Stearns):
<https://www.youtube.com/watch?v=W_Q8cNzjTv0> and
<https://www.youtube.com/watch?v=KQqcD3_3_Y8>

Utilitarianism

Crimmins, James E. (2018), 'Jeremy Bentham', *Stanford Encyclopedia of Philosophy* <https://plato.stanford.edu/archives/sum2018/entries/bentham/> [an informative overview of Bentham's utilitarianism, including his later amendments]

Brink, David (2016), 'Mill's Moral and Political Philosophy', *Stanford Encyclopedia of Philosophy* <https://plato.stanford.edu/archives/win2016/entries/mill-moral-political/> [§2 has a thorough discussion of the main themes of Mill's utilitarianism]

Nathanson, S. (2018), 'Act and Rule Utilitarianism', *Internet Encyclopedia of Philosophy* < https://www.iep.utm.edu/util-a-r/> [a very helpful summary of the debate]

Philosophers in conversation on aspects of utilitarianism:
http://philosophybites.com/2012/02/philip-schofield-on-jeremy-benthams-utilitarianism.html
http://philosophybites.com/2009/04/thomas-hurka-on-pleasure.html
http://philosophybites.com/2007/07/roger-crisp-on-.html

An thorough introduction to utilitarianism from teachphilosophy (Prof Paul Stearns):
https://www.youtube.com/watch?v=kp4a5u_IIL4&index=3&list=PLFGHE1xQFhhxZ0jYd9K0woLbeNhJsv0Ni

II. Applied ethics

Normative ethical theories are intended to guide us in knowing and doing what is morally right. So we can understand the theories better if we consider them in relation to practical issues. We will look at four practical issues – telling lies, stealing, eating animals and simulated killing (killing in a fictional context – videogames, plays, TV dramas etc).

Before we start, a brief reminder of the main principles of our three theories.

1.  Act utilitarianism, in its simplest form, says that an action is right if it maximises happiness, and wrong if it does not. Rule utilitarianism says that an action is right if it complies with those rules which, if everybody followed them, would lead to the greatest happiness (compared to any other set of rules).
2.  Kant's Categorical Imperative says that we should 'Act only on that maxim through which you can at the same time will that it should become a universal law'. The Second Formulation is 'Act in such a way that you always treat humanity, whether in your own person or in the person of any other, never simply as a means, but always at the same time as an end'.
3.  Aristotle's virtue ethics says that eudaimonia is living in accordance with reason and the virtues, which means feeling and acting 'at the right times, with reference to the right objects, towards the right people, with the right motive, and in the right way', and it takes practical wisdom to know what is right.

For any theory and any issue, if we think the theory gives the wrong answer on the issue – e.g. it allows lying too easily or prohibits stealing too strictly – we can use this result as an argument that the theory is mistaken and should be rejected.

A. Telling lies

A simple act-utilitarian approach to lying would consider whether telling a lie creates greater happiness than telling the truth (or keeping silent). If it does, then it is morally right. If it doesn't, then it is morally wrong.

Mill notes that a person's being truthful is of great benefit to people's happiness generally. Being able to trust what others say is the basis of living together. Weakening either our tendency to be truthful or other people's trust is, therefore, severely damaging to happiness. To tell a lie just for the sake of convenience is therefore morally wrong. But lying is sometimes permissible, e.g. when it is the only way we can withhold information from someone who intends to do great harm. When else? Should we lie when the truth would hurt (e.g. in cases of terminal illness or sexual infidelity); or because the truth would be damaging to some long-term good (e.g. in politics)? It is hard to know what the right 'rule' for lying should be.

We discussed Kant's view of lying in 'Issues for Kantian deontological ethics'. If everyone told lies when they wanted to, people would stop believing each other. But you can deceive someone with a lie only if they believe you. So the maxim cannot be universalised, and lying is wrong. Or again, to lie to someone is to treat them as a means to our own ends. They are not able to make an informed choice about what to do, but are manipulated in a way that they are unaware of.

Kant says we should not lie even when the other person's ends are immoral, and we are trying to prevent those ends being realised. Even in the example of the would-be murderer who asks you where his victim is, Kant says we should not lie. We responded that perhaps 'to save a life by lying' could be universalised. Yet it would still treat the would-be murderer as a means to an end. Is this wrong?

Aristotle doesn't discuss lying in general, only being truthful about oneself (as opposed to boasting and false modesty). But he does says that truth is, in itself, praiseworthy and falsehood blameworthy. This could mean that lying is an act, like adultery and murder, that has no mean, a deficiency of truthfulness that cannot be virtuous. Alternatively, Aristotle may mean that truth is a final end, something that we should seek for its own sake. This doesn't entail that lying is always wrong, since there are few strict rules in ethics. Instead, we will need practical wisdom to judge when lying is justified. If we deceive someone, to do so virtuously, we would need to do so at the right time, with the right motive, about the right truths, and in the right way. For example, perhaps the virtuous person will exhaust all the alternatives (e.g. refusing to answer) before resorting to a lie.

B. Stealing

To steal is to take someone else's property with no intention of returning it and without their permission (or without the legal right to do so). To own property involves a system of rights. I own my books, for instance. This means that I am free to use them, but no one else is (unless I say so). Because property involves rights, whether or not we may steal is a matter of justice. To deprive me of my property through stealing looks like a violation of justice. So we start from the presumption that stealing is wrong, but arguments might overturn that presumption.

Act utilitarianism says that if stealing, on some occasion, creates greater happiness than not stealing, then it is morally right on that occasion. Otherwise, it is morally wrong. While it usually decreases happiness, it may be okay to steal in order to save a life or when we are in dire need, and especially if we steal from someone wealthy who won't be significantly harmed by the loss of their property. Or again, it may be okay to steal from the wealthy where property is unjustly distributed, especially if we intend to benefit the poor – a 'Robin Hood' kind of stealing. These are occasions on which stealing would lead to greater happiness than not doing so.

For rule utilitarians, a clear, simple rule on stealing will probably lead to better consequences overall than trying to build in lots of exceptions. Although 'don't steal' won't lead to the best consequences in every case, allowing exceptions may tempt people to think that their situation is 'exceptional'. We don't want people to be afraid that others will steal from them whenever the thief convinced himself that it would bring about greater happiness. In such a society, people would cease to trust each other. Mill recognises the importance of feeling secure. Respecting property rights contributes most to happiness in the long term, and Mill comments that justice is 'most sacred and binding part of all morality'. This suggests that we should never steal for any reason other than justice itself.

For Kant, acting on the maxim 'To steal something I want if I can't afford it' leads to a contradiction when universalised. Stealing can only be the right thing to do if everyone could do it. However, if we could all just help ourselves to whatever we wanted, the idea of 'owning' things would disappear. Now, by definition, you can't steal something unless it is owned by someone else. But people can only own things if they don't all go around helping themselves whenever they want. So it is logically impossible for everyone to steal things.

Does this allow exceptions on the basis of a different maxim? 'To steal in order to save a life' would not, if universalised, lead to the end of property, because it is rare that anyone would need to steal for this purpose. However, it is unlikely that Kant would agree with this amendment. Stealing to save a life is still stealing, it is the same *kind of action* as stealing because one wants something. To steal from someone involves not allowing them to make an informed choice, so it treats them as a means to an end. Why not just ask them to give you what you want? The whole idea of justice and rights is that there are constraints on *how* we pursue good ends. The duties of justice are more stringent than the duties of virtue (doing good). So we should not steal.

Aristotle's account of justice is largely deontological. Although Aristotle does not understand justice in terms of rights, he says that stealing is never in the mean, but always wrong. Stealing involves depriving someone of their 'due' or 'fair share'.

But is this right? We should not assume that simply because someone owns something that they own it 'fairly'. What about stealing in an unjust situation, where some people have more than is fair while others are poor? For Aristotle, it may depend on how the situation came about. If an unjust situation is not the result of anyone's action, this is unfortunate, but deliberately doing something unjust, such as stealing, is worse. But if the unjust situation is the result of people's actions, e.g. the rich using power to keep the poor poor, then this situation should be put right. Stealing is not normally an act of justice in this sense, but if it were, then perhaps it could be justified. If someone steals your phone, you aren't stealing if you take it back without their permission. If, in the tales of Robin Hood, the rich did not own their wealth because they had literally stolen it from the poor, then we could say that Robin Hood didn't steal from the rich at all. He simply returned to the poor what was stolen from them.

C. Eating animals

Is eating meat morally wrong? In 'Issues for utilitarianism', we saw that there is nothing in utilitarianism that gives us a reason to privilege human happiness over the happiness of non-human animals. To do so, argues Peter Singer, is 'speciesism'. The obvious differences between human beings and other animals – reason, the use of language, the depth of our emotional experience, our self-awareness, our ability to distinguish right and wrong – are not relevant when it comes to thinking about pleasure and pain. When it comes to suffering, animals should be treated as equal to people.

This doesn't mean that eating animals is always wrong. Utilitarianism only objects to suffering, not to killing. If you painlessly kill an animal and bring another animal into being, you haven't reduced the total amount of happiness in the world. According to Singer, we need only ensure that animals are happy when they are alive, and slaughtered painlessly. This would make eating meat much more expensive, because animals would have to be kept in much better conditions. The way in which we rear animals for food at present means that eating meat is wrong.

Kant argues that animals do not have a rational will. They have desires, but to have a will is to be able to stand back from one's desires and reflect on whether or not one's desires are good and whether or not to act on them. The maxim that human beings eat animals can be universalised – we can all will this without contradiction. Similarly, because animals are not rational wills, they are not ends in themselves, and can therefore be treated as means to our ends. We may therefore eat them.

However, Kant argued that if we lack kindness towards animals, we may become unkind towards other people – and this would be morally wrong. Therefore, we need to treat animals in such a way that we don't become vicious.

One objection to Kant's theory here is that it implies the same result for human beings who aren't rational wills, including babies and some people with severe mental disabilities. Can we treat them as means to an end? While Kant could respond that babies have the potential to become rational wills, people with severe mental disabilities never will. Can we eat them? Whatever is the basis of our duty of care towards people with severe mental disabilities, we may argue, should equally give us a duty of care towards animals.

Aristotle argues that animals have no share in eudaimonia because they are incapable of either practical or theoretical reason. Our primary concern with eudaimonia has little place, therefore, for the consideration of animals. Recent virtue theorists have argued that a deeper understanding of the relationship between human beings and animals provides a different answer.

Cora Diamond argues that there are many different practices in eating animals, and some may be wrong while others are not. For example, rearing your own pigs, looking after them well and killing them humanely for yourself is very different from picking up sliced pork in the supermarket. To state the obvious, the meat industry is an industry, with animals reared, slaughtered and shipped around the world on an industrial scale, and many of the

things that are done to animals in the process are done not because they benefit the animals, but to make the process more efficient. By contrast, we recognise that animals are 'fellow creatures'. Like human beings, they can die, they lead their lives without our knowledge of what they do, they can provide company for us. These thoughts respect the independent lives that animals have. To treat them just as means to our ends, as in the meat industry, is callous and selfish. Killing animals for food may be morally permissible, if it is done for the right reason and in the right way, i.e. with the appropriate feelings. Because that is so often not the case, eating animals is often morally wrong.

D. Simulated killing

Simulated killing is the dramatization of killing within a fictional context, e.g. in video games, films and plays. It is not merely the description of a killing, as in a novel, but a fictional enactment of killing. Of course, no one is actually killed; so, if simulated killing is wrong, it is not wrong for the same reasons that killing is usually wrong. But is it morally permissible to create or participate in any fictional representation? While it has become widely socially acceptable to play violent video games, video games involving rape and paedophilia are banned in the UK. And we may ask about the effects of simulated killing both on the people involved and on how they then treat other people in real life. We will focus just on killing people within video games.

Utilitarianism will note that in playing a video game, no one is actually harmed in simulated killings, so as long as the gamer is enjoying themselves, there is a gain of happiness. But if the gamer is, as a result of simulated killing, more likely to kill, be aggressive or antisocial in real life, then this could decrease overall happiness. The evidence for this is not clear (see Garry Young, *Ethics in the Virtual World*). Some studies indicate that there is an increased risk of aggressive thoughts and behaviour for up to 75 minutes after playing. However, others found that this increased risk occurred only in some people. There is no firm indication of long term effects.

For Kant, it is important that simulated killing involves no violation of moral duty. But if simulated killing damages one's rational will or leads to neglecting or violating one's duty to other people, then we can object. Cultivating cruelty and an indifference to virtual suffering through simulated killing could undermine our willingness and ability to treat others as ends in themselves in real life. But, again, there is currently not enough evidence to say that this is true.

Aristotle argues that the kinds of actions we do repeatedly lead to similar habits. We become just by doing just acts. Likewise, we become unjust by doing unjust acts. Killing is often an unjust act. So the cumulative effect of playing games which involve simulated killing may lead to the development of injustice and unkindness, or at least inhibit the development of justice and kindness. But even if *doing* unjust acts may make someone unjust, why think that *simulating* unjust acts will develop injustice? Once again, we can argue that the evidence doesn't support this claim.

We can ask whether a virtuous person would engage in playing video games that involve simulated killing. If so, then they will do so in the right way, with the right motive, and at the right times, as this is what defines virtuous action. What might that involve? For example, why would someone want to simulate killing someone else? Is taking pleasure in this activity virtuous? Virtue ethics also recognises that the right thing to do is not the same for everyone. Someone who draws a clear conceptual and emotional distinction between simulated killing and real life may be at no risk of being more aggressive after playing or developing bad character traits. Someone who cannot draw such a distinction may be at risk, and so should not play.

*Study questions*
1. Are there any circumstances in which either telling a lie or stealing is permissible? Why?
2. Is vegetarianism a moral duty?
3. What, if anything, is wrong with simulated killing?
4. Do issues in applied ethics support or challenge utilitarianism?
5. Should we expect clear moral guidance from ethical theories if life is messy?

*Further reading*
Singer, P. *Practical Ethics* (3rd ed.), CUP 2011 [Ch. 2 and 4 defend a utilitarian standpoint on the treatment of animals]
DeGrazia, D. *Animal Rights: a Very Short Introduction*, OUP 2002 [accessible with lots of real-world cases and information]
Young, G. *Ethics in the Virtual World*, Routledge 2013 [*the* book for discussion of simulated killing]

*Internet resources*
Wilson, S. (2018) 'Animals and ethics', *Internet Encyclopedia of Philosophy*, <https://www.iep.utm.edu/anim-eth/> [a very helpful overview of the debate]
Doggett, T. (2014) 'Is it morally permissible to kill animals for food?', <https://www.youtube.com/watch?v=3HAMk_ZYO7g&index=8&list=PLtKNX4SfKpzWO2Yjvkp-hMS0gTI948pIS> [a fun and easy introduction to the ethics of vegetarianism]

III. Metaethics

Normative ethical theories are intended to provide guidance on how to live. Metaethics, by contrast, does not do this. It asks about what morality is, philosophically speaking. It asks questions like 'What do statements like "Murder is wrong" mean?, Can they be true or false? Or do they express feelings rather than state facts?'; 'What is it to hold a particular moral view, like "murder is wrong"? Is this a belief or a feeling or something else?'; 'Are there objective moral truths and objective moral properties?'; 'How do we know what is right and wrong?'

*Cognitivism, non-cognitivism, realism and anti-realism*
Theories in metaethics fall into two broad families – **cognitivism** and **non-cognitivism**. Cognitivism claims that ethical language expresses beliefs. Beliefs can be true or false, so ethical claims that can be true or false. To believe that murder is wrong is to believe that the sentence 'Murder is wrong' is true. Because (usually) a claim is true because it correctly describes how the world is, cognitivists (usually) also claim that ethical language aims to describe the world.

Non-cognitivism claims that ethical language does not express beliefs, but some other mental state. Different non-cognitivist theories disagree on the kind of mental state, but it is usually an attitude or feeling. Ethical claims do not try to describe the world and cannot be true or false. So 'Murder is wrong' is neither true nor false, but an expression of, say, the speaker's disapproval of murder.

Here are three quick arguments in favour of cognitivism:

1. We think we can make mistakes about morality. Children have to be taught what is right and wrong. If there were no facts about moral right and wrong, it wouldn't be possible to make mistakes.
2. Morality feels like a demand from 'outside' us. We feel answerable to a standard of behaviour which is independent of what we want. So morality is objective, not determined by what we think about it.
3. Many people believe in moral progress. But how is moral progress possible, unless some views about morality are better than others? And how is *that* possible unless there are facts about morality?

But non-cognitivists have their own arguments in response:

1. Moral judgements are, like desires, motivating. Holding the view that murder is wrong involves being motivated not to murder. But factual beliefs are not motivating. The sun is 93 million miles from the Earth – so what? Believing that fact inclines me to do nothing in particular at all. Because moral views are motivating, they are not beliefs, but non-cognitive attitudes.
2. Cognitivism needs to explain *how* moral claims can be objectively true or false. Are there moral properties 'in the world'? What kind of property could they be, and how can we

find out about them? Non-cognitivism argues that it is a simpler theory. It has a simpler metaphysics and a simpler epistemology.

This last point leads to a distinction between **moral realism** and **moral anti-realism**. Moral realism claims that good and bad are properties of situations and people, right and wrong are properties of actions. Just as people can be 5 feet tall or run fast, they can be morally good or bad. Just as actions can be done in 10 minutes or done from greed, they can be right or wrong. These moral properties are a genuine part of the world.

Moral anti-realism claims that there are no such moral properties. Non-cognitivist theories are anti-realist. Since moral judgments do not describe the world and are neither true nor false, then there are no **mind-independent moral properties** that would make moral judgments true or false. For example, to say 'racism is wrong' is not to claim that racism has any kind of property. It is, say, to express disapproval of racism.

We will look at two cognitivist, realist theories – **naturalism** and **non-naturalism** – and two non-cognitivist, anti-realist theories – **emotivism** and **prescriptivism**. We will also look at the interesting case of a cognitivist theory that is anti-realist – **error theory**.

A. Two types of moral realism: naturalism and non-naturalism

*Moral naturalism*
Moral naturalism claims that moral properties are **natural properties**, i.e. properties that we can identify through sense experience and science. It claims that moral properties, e.g. goodness and rightness, are *the same properties* as certain properties that we can identify empirically. The most plausible natural properties that might be moral properties are certain psychological properties, e.g. happiness.

Utilitarianism can be interpreted as a form of moral naturalism, if we say that happiness is goodness, i.e. they are the same property (like water is H$_2$O). So we can find out what is morally right and wrong through experience because we can discover what creates happiness empirically. Morality is objective because what contributes to people's happiness is objective.

This is how G. E. Moore interprets Mill's 'proof' of utilitarianism. Mill argues that what is good is what is desirable. We can find out what is desirable by finding out what people desire. He then argues that happiness is desired. From this, he infers that happiness is good. This only works, says Moore, if Mill thinks that what is good is the same thing as what is desirable, and that what is desirable is the same thing as what is desired. So Mill must be thinking that goodness is the natural property of what is desired.

However, we can understand Mill's argument differently. Mill takes what people desire (which is a natural property) as *evidence* for what is desirable (good). He does not say that goodness is the *same property* as being desired. And when he claims that what is good is what is desirable, nothing he says shows that 'being desirable' (as opposed to being desired) is a natural psychological property or a distinct moral property. So we can't say whether Mill is a naturalist or not.

*Moral non-naturalism*
Moral non-naturalism claims that moral properties are not natural properties, i.e. they are not properties that can be discovered through science and experience. Instead, they are distinct, **non-natural properties**. We will look at just one kind of non-naturalism, Moore's **intuitionism**.

In *Principia Ethica*, Moore called the attempt to identify goodness with any natural property the '**naturalistic fallacy**'. Goodness, he argued, is a simple and unanalysable property. It cannot be defined in terms of anything else. We can only say how people use the term 'good' and what has the property of goodness. Colours are similar. We can say what things are yellow, e.g. the sun, ripe lemons, etc. but these things aren't the same thing as the colour yellow, and they don't define it. Yellow is a simple property, and no one can explain what yellow is to someone who doesn't know. Similarly, happiness isn't the same thing as goodness, just as lemons aren't the same thing as yellow. But unlike colours, goodness cannot be investigated empirically. It is real, but it is not part of the world of science.

Moore supports his view by the '**open question argument**'. An open question is a question which has more than one logically possible answer. If goodness is happiness, say, it wouldn't

make sense to ask 'Is happiness good?' This would be like asking 'Is happiness happiness?' This second question isn't an open question, because the answer has to be 'yes'. It cannot, logically, be 'no'. It's like asking 'Is a bachelor an unmarried man?' But, says Moore, asking 'Is happiness good?' is an open question – the answer can logically be 'yes' or 'no'. And so goodness cannot be happiness. In fact, 'Is x good?' is always an open question. And so goodness cannot be defined as any other property.

But Moore's open question argument doesn't work. Here is a similar argument. 'The property of being water cannot be any other property in the world, such as the property of being $H_2O$. If it was then the question "Is water $H_2O$?" would not make sense – it would be like asking "Is $H_2O$ $H_2O$?" So being water is a simple, unanalysable property.' This is not right, as water just is $H_2O$.

Moore's argument confuses concepts and properties. Two different concepts – 'water' and '$H_2O$' – can pick out the same property in the world. Before the discovery of hydrogen and oxygen, people knew about water. They had the concept of water, but not the concept of $H_2O$. So they didn't know that water is $H_2O$. 'Water is $H_2O$' is not a conceptual truth. However, water and $H_2O$ are one and the same thing – the two concepts refer to just one thing in the world. Water is identical to $H_2O$.

Likewise, the concept 'goodness' is a different concept from 'happiness'. 'Happiness is good' is not a conceptual truth. But perhaps the two concepts refer to exactly the same property in the world, so that goodness is happiness.

If moral properties are not natural properties, how do we discover them? How do we know what is good? Moore says that we can know claims about what is good to be true (or false) by considering the claims themselves. He calls these claims 'intuitions'. Intuitions are 'self-evident' judgements. A self-evident judgement rests on the 'evidence' of its own plausibility, which is grasped directly. ('Self-evident' is not the same as 'obvious'. Our ability to make a self-evident judgement needs to develop first, and we need to consider the issue very carefully and clearly.) Intuitionists compare moral intuitions to mathematical ones (how do you know 2 + 2 = 4? – just think about it). They don't tend to claim that we have a faculty of intuition that 'detects' whether something is good or not, a bit like a supernatural sense. Intuitionism is simply a form of non-naturalism that claims that some of our moral judgements are self-evident.

*An objection from empiricism*
Empiricists argue that our knowledge is limited to what we learn through experience. For example, **Hume's fork** claims that we can have *knowledge* of just two sorts of claim: relations of ideas and matters of fact. Relations of ideas can be discovered just by thinking, by recognising the truth of an **analytic proposition** or by deductive reasoning. A statement is analytic if it is true or false in virtue of the meanings of the words. For example, 'Bachelors are unmarried' is analytic and true; 'Squares have three sides' is analytic and false. To deny a relation of ideas is to contradict oneself. Matters of fact are claims about what exists, and they are established by sense experience and causal inference. In other words, they are empirically verifiable, i.e. empirical evidence would go towards establishing whether a matter of fact statement is true or false.

A J Ayer's **verification principle** adapts Hume's fork to say that these are the only two types of claim that are *meaningful*. A statement only has meaning if it is either analytic or empirically verifiable. Unless there is some way of showing, at least in principle, that a statement is true or false, then it doesn't really say anything, it doesn't make a meaningful claim.

This view of what can be known (or what is meaningful) challenges moral cognitivism, and so moral realism, and so naturalism and non-naturalism. Hume claims that moral judgments are neither relations of ideas nor matters of fact. Similarly, Ayer argues that they are neither analytic nor empirically verifiable. So we cannot have moral knowledge. Instead, we should reject cognitivism in favour of non-cognitivism.

Moral judgments are not analytic. To deny an analytic claim, e.g. 'black is the same colour as white' or '3 x 5 ≠ 30 ÷ 2', is a contradiction. Denying a moral claim isn't a self-contradiction. On the other hand, if moral judgments are matters of fact, *which* fact? Only natural facts are facts, say empiricists, and Ayer agrees with Moore that you can't identify good with any natural property. We can show that murder causes grief and pain, or that it is often done out of anger. But we cannot show empirically that it is wrong.

Naturalists respond that moral judgments *are* (natural) matters of fact. But it takes philosophical reasoning to show which matters of fact they are, e.g. to show that goodness is happiness, we need to defend Mill's 'proof' against objections. Non-naturalists reject empiricist constraints on knowledge. Moral judgments are neither analytic nor empirically verifiable. They are self-evident (or inferred from self-evident claims). Ayer rejects this: we can't establish the truth of a moral claim by appealing to 'intuition', because if two intuitions conflict, there is no test that will establish which intuition is correct.

*Hume's argument from motivation*
Hume attacks cognitivism (and so realism) in another way: moral judgements can motivate actions, but relations of ideas and matters of fact cannot. Simply understanding that some relation holds between two ideas, e.g. that 3 x 5 = 30 ÷ 2, doesn't get us to act one way or another. Knowing facts about the world might well tell us how to *achieve* what we want, but it can't make us want anything in the first place. What could knowing that there is food in the kitchen lead me to do anything without some desire (to eat, to cook…) to act upon? As these are the only types of judgment about what is true or false, moral judgements can't be true or false – they must express non-cognitive attitudes.

Some realists respond by denying that moral judgments motivate actions. To do good actions, we have to have the *desire* to be good as well. To say 'Murder is wrong' is simply to describe murder. Someone could hold this belief, but not care about what is wrong, and so be quite willing to murder if it suited them (a sociopath perhaps?). If moral judgements *don't* motivate us on their own, then Hume's argument gives us no reason to believe that moral judgements aren't about what is true.

*Hume's is-ought gap*

Hume's third argument against cognitivism draws a distinction between what *is* the case and what *ought* to be the case. How, for instance, do we infer from the fact that some action will cause pain to the claim that we ought not to do it? What's the connection? Suppose I say 'Eating meat causes animal suffering. Therefore, you shouldn't eat meat.' According to cognitivism, the conclusion states a truth, and this truth is inferred from the premise. But how is this a rational inference, Hume asks? The premise tells me how the world is; the conclusion tells me how the world ought to be. There is a logical 'gap' between what is and what ought to be, so that we can't reason from one to the other.

If moral judgments were true or false, we *would* be able to infer them from other truth claims, such as matters of fact. But we can't. So moral judgments don't make truth claims; cognitivism, and so realism, and so naturalism and non-naturalism, are false.

B.  Error theory: cognitivist anti-realism

*Mackie's error theory*
In *Ethics: Inventing Right and Wrong*, John Mackie argues that moral realists are right to say moral judgements are cognitive. We use ethical language to make objective claims about right and wrong. However, he argues, there is no objective right and wrong. And so he argues that all moral judgements are false. 'Murder is wrong' is false, because the property of being wrong does not exist. 'Murder is right' is false, because the property of being right does not exist. There are no moral properties. Ethical language rests on a mistake. Mackie is an 'error theorist'.

Error theory is a form of moral anti-realism. It accepts cognitivism, but rejects moral realism: there are no mind-independent moral properties to make moral judgments true.

Here is an analogy. Some people really believe in fairies. They don't think that when we are talking about fairies, we are using language 'fictionally' (like when we talk about Sherlock Holmes). Imagine that everyone believed in fairies in this way. An 'error theory' of fairies would say that while talk of fairies is cognitivist, there are no fairies. It is not true that fairies have wings, because there are no fairies. It is not true that fairies don't have wings, because there are no fairies. All claims about fairies are false, because there are no fairies.

Mackie offers three arguments for thinking that there are no objective moral properties.

*Mackie's argument from relativity*
Moral codes differ from one society to the next. For example, according to one society, slavery is permissible under certain conditions; according to another, it is never permissible. This is 'descriptive relativism'. This doesn't yet show that moral realism is wrong. Societies have also disagreed on empirical facts, e.g. some have thought that the Earth is flat, others that it is round. But in this case, there is an objective truth. So disagreement over a claim doesn't show that there is no objective truth. However, with science, different societies are trying to *discover* the truth about the world. With morality, Mackie says, there are simply different ways that human beings live, different conventions, and these conventions are *reflected* in their moral judgments. This is '**meta-ethical relativism**'.

A moral realist can deny that moral codes differ very much, e.g. all societies prohibit certain types of killing. Or a realist can argue that 'relative' ≠ 'not objective'. For example, 'Chilli plants will grow well' is a relative truth – it is true in hot countries but not cold ones. But whether a chilli plant will grow well in the country you are in is still objective. Moral realists can say the same about moral judgments. Some ethical practices will be permissible in some circumstances but not in others. Whether a moral judgment is true will depend on whether the practice is *objectively* morally right or wrong *in those circumstances*.

*Mackie's argument from epistemological* **queerness**
If there are objective moral properties, how can we know about them? Mackie argues that none of our usual ways of knowing about the world – sense perception, introspection, hypothetical reasoning, even conceptual analysis – can explain knowledge of morality.

Intuitionism, he claims, says no more than that we have some special faculty – but this is a terrible answer that doesn't explain how we have this knowledge at all.

Like Hume and Ayer, Mackie is appealing to empiricism. The non-naturalist will respond by rejecting empiricist limitations on knowledge. But Mackie presses Hume's is-ought gap as well. Suppose we condemn some action for causing pain just for fun. What is the relation between the natural fact that it causes pain for fun  and the 'fact' that it is wrong? Intuitionism fails to tell us how natural facts contribute to moral thinking. This makes it even more puzzling how we could come to know about moral properties.

*Mackie's argument from metaphysical queerness*
If there were moral properties, Mackie argues, they would have to be very different from anything else in the universe. Moral judgements motivate us – we avoid actions we believe are wrong and try to do actions that are right. But that means, if there were moral properties, simply *knowing* what is right or wrong would be enough to motivate us to act in certain ways. How could an objective property motivate us? How could there be some direct, immediate relation between some fact of the world and our desires? Natural properties don't motivate us this way. So if there are moral properties, they must be non-natural properties. What Mackie's argument is supposed to show is how 'metaphysically queer' (strange) such non-natural properties would have to be.

Moral naturalists challenge Mackie's assumptions here. If there are moral properties, they will be – or be closely related to – psychological properties (happiness, pain, etc.). Psychological properties, such as whether one is happy, are natural facts, but because they are part of our minds, they are also motivating. So, if goodness is happiness, for example, then it is not at all 'queer' that facts about goodness motivate us. We desire happiness and are motivated to pursue it.

Mackie's arguments from queerness originally targeted non-naturalism. It is harder to see how non-naturalist theories such as Moore's intuitionism can respond.

C. Two types of non-cognitivist anti-realism: emotivism and prescriptivism

Non-cognitivist theories of ethics claim that moral judgments do not express beliefs, but some other psychological state. Moral judgments are not true or false and do not describe the world. Non-cognitivist theories are anti-realist – there are no moral properties that make moral judgments true or false.

*Emotivism*
Emotivism claims that moral judgements *express* approval or disapproval. To say that 'Murder is wrong' is to express one's disapproval of murder. Ethical language is 'emotive'. In *Language, Truth and Logic*, A J Ayer defends this view by appealing to his verification principle. Moral judgments do not make claims that are true or false as they are neither analytic nor can they be verified empirically. Instead, he argues, moral judgments express feelings. Through expressing our feelings, they also aim to arouse feelings in others, and so influence what they do.

However, the verification principle faces a famous objection. According to the verification principle, the principle itself is meaningless. The claim that 'a statement only has meaning if it is analytic or can be verified empirically' is not analytic and cannot be verified empirically. So if the principle of verification is true, it is meaningless. But if it is meaningless, it is not true. So it cannot be true. Therefore, it does not give us any reason to believe that ethical language is meaningless.

However, emotivism does not depend on the principle of verification. An emotivist only needs to defend the positive claim that the purpose of moral judgements is to express approval and disapproval. If true, this explains how and why it is that moral judgements motivate us, since these feelings obviously influence our actions.

*Prescriptivism*
In *The Language of Morals*, R. M. Hare argued that moral language is prescriptive. When I express a moral judgement, Hare says, I am saying what you ought to do. I am not trying to influence you, nor am I expressing my feelings. Whether you act as I prescribe is a separate matter. Simply saying you should do x isn't an attempt to persuade you.

Hare identifies two types of prescriptive meaning – 'commanding' (right/wrong) and 'commending' (good/bad). For example, 'Eating meat is wrong' entails the command 'Do not eat meat'. To accept that eating meat is wrong is to not eat meat. And we say something is 'good' when we want to commend something. This provides guidance, though it doesn't command.

To say something, e.g. chocolate, a teacher, a person, is good is to praise it. This is different from describing it. Suppose I say 'That's a good strawberry, because it is sweet and juicy'. If we think 'good' as applied to strawberries just means 'sweet and juicy', then all I have said is 'That's a sweet and juicy strawberry because it is sweet and juicy'. But this isn't what I said. I commended the strawberry, I didn't merely describe it. But in praising something as good, we draw on certain standards. For example, good teachers can explain new ideas clearly and create enthusiasm in their students. When we use 'good' to mean 'morally good', we are

appealing to a set of standards for how to live as a person. There are no facts that establish one set of moral standards as objectively correct. We have to *adopt* the standards.

Whenever we apply a standard in making a prescription, we are committed to making the same judgement of two things that match the standard in the same way. If I say this chocolate is good but that chocolate is not, I must think that there is some relevant difference between the two. Similarly, if one action is right and another wrong, there must be a relevant difference. If there is no relevant difference, they are either both right or both wrong. In other words, we must be willing to 'universalise' our moral judgements.

*Moral language*
Emotivism and prescriptivism each identify a key function of moral language. But, we can object, moral language is not always used in either of these ways. We not only prescribe, but can also implore someone, confess, complain, and so on. Or again, we sometimes express our moral attitudes to others who already agree with them or that we know to be indifferent to our views – so influencing their attitudes is not the purpose. Or again, moral language isn't always emotive, but can be dispassionate.

However, both emotivism and prescriptivism can argue that their analysis identifies the *central purpose* of moral language. If we didn't try to influence or guide others, we would have no moral language or judgements at all. However, that doesn't mean that it always has to be used for this purpose. But then, we can reply, if the purpose of moral language gives us its *meaning*, what does it mean when not being used for this purpose?

*Moral reasoning*
One use of moral language is in moral reasoning – using premises to draw conclusions about what to do. Can emotivism and prescriptivism explain moral reasoning, e.g. in disagreements?

If I say 'abortion is wrong' and you say 'abortion is right', according to emotivism, I am expressing my disapproval of it and you are expressing your approval. I am also trying to influence your attitudes, and you are trying to influence mine. But we are not doing so *rationally*, or by appealing to facts about what is good or bad. Trying to influence people without reasoning is just a form of manipulation.

Ayer responds that moral arguments are not arguments over moral judgements, but over natural facts. When arguing over animal rights, say, we are constantly drawing facts to each other's attention. I point out how much animals suffer in factory farms. You point out how much more sophisticated human beings are than animals. And so on. If we both agree on the facts, but still disagree morally, there is nothing left to discuss, says Ayer. Moral judgements always presuppose a system of values; but no arguments for these values can be given.

Hare argues that prescriptivism can do better. First, we can ask about someone's reasons for prescribing what they do. Second, morality involves consistency – moral judgements must be universalised. For example, Peter Singer claims there is no relevant difference between the suffering of people and the suffering of animals. If we are going to say that

causing the suffering of people is wrong, we are committed to saying the suffering of animals is wrong – unless we can find a relevant difference. Moral disagreements can be about the consistency in applying certain standards, and reason can help resolve this. Third, we can infer prescriptions from other prescriptions. A famous argument against abortion says 'Taking an innocent human life is wrong. Abortion is the taking of an innocent human life. Therefore abortion is wrong.' We can rephrase this as commands: 'Do not take innocent human life. Abortion is the taking of an innocent human life. Therefore, do not commit abortion.' To reject or refuse the conclusion, we must reject or refuse at least one premise. And so our prescriptions are logically related to one another. So we can use reason to discuss these relations. Moral arguments are not only about the facts, but about moral judgments as well.

But while prescriptivism can appeal to rational consistency, we can object that consistency isn't enough on its own. In requiring us to universalise moral judgements, Hare's theory is similar to Kantian ethics. However, Kant argues that the standards for a good person (the good will) are themselves set by reason, and are therefore objective. Hare rejects this. Neither the empirical facts nor reason entails that we must have certain standards rather than others. If I argue that racism is morally right, as long as I am prepared to universalise this claim, there is no objective ground on which to disagree with me. Suppose you say 'But what if you were of a different race. Would you say you should be treated as inferior?' I can reply 'Yes.' Now what? Prescriptivism does not rule out very objectionable values.

*Does moral anti-realism become moral nihilism?*
Moral **nihilism** is the rejection of all moral values and principles. It is the view that nothing is of moral value, that we have no moral duties. Moral anti-realism – including emotivism, prescriptivism and error theory – claims that there are no objective moral truths. We may object that if morality is simply an expression of our attitudes, then really there are no moral values – we invent them. Really, we have no obligation to be moral, because we have no obligation to have certain emotions or adopt certain standards of value. If moral judgments are not objectively true, then why be moral at all?

Moral anti-realists can argue that this is an unfair simplification of their theories. Living without moral values is itself a choice or expression of feeling, and one that morally good people will disapprove of. The theory that moral values are a reflection of our feelings does not imply that we should stop having moral feelings or making moral judgments.

*Moral progress*
If there is no moral reality, we can argue, then our moral views cannot become better or worse. People used to believe that slavery was morally acceptable and now they do not. But how can we say that this is progress if there is no objective moral truth?

Moral anti-realists can claim that people's moral views can improve by becoming more rational. We come to know facts we didn't know before, e.g. that slaves are not stupid. We can become more consistent, e.g. in how we treat animals. And we can make our judgments more coherent, resolving conflicts and dilemmas. Moral progress just means becoming more rational in our moral thinking, not becoming more 'correct' in our moral judgements.

Alternatively, moral anti-realists say that if we disapprove of past moral codes and approve of our own moral code, then we think that society has moved from moral principles that were bad (i.e. principles we disapprove of) to moral principles that are good (i.e. principles we approve of). That is what moral progress is – a non-cognitive moral judgment.

*Concluding challenges*
While moral anti-realists have answers to the four objections raised above, we must decide whether their answers are persuasive, not only theoretically (in terms of philosophical argument) but also practically (in terms of how we live our lives). Realists continue to press their case that non-cognitivists fail to do justice to how people think about ethics in real life – can we really suppose that our criminal law, say, or our feelings of anger toward racism or slavery, rest on nothing more than subjective commitments that other human beings may reject without failing to understand what is truly, objectively right? On the other hand, anti-realists appeal to our scientific conception of reality to question how it is really possible for us to think that there are objective values waiting to be discovered. As they seek to meet these challenges, both realists and anti-realists develop their theories. It is easy to misunderstand both sides and challenging to think about the complex place morality has in human life and thought.

*Study questions*
1. Explain and illustrate the differences between cognitivism and non-cognitivism.
2. Are moral properties a kind of natural property?
3. Does Moore successfully defend intuitionism?
4. Does Hume show that moral cognitivism is false?
5. Is Mackie right that moral language rests on a mistake?
6. Explain the similarities and differences between emotivism and prescriptivism.
7. Can emotivism explain moral reasoning?
8. If moral judgments are not objectively true, do we have any reason to be moral?

*Further reading*
Fisher, A. *Metaethics*, Routledge 2014 [a genuinely introductory text – as Fisher says, 'an introduction to an introduction to metaethics'!]
Chrisman, M. *What is This Thing Called Metaethics?*, Routledge 2017 [an excellent, clear introduction to the theories, with lots of student-friendly features]
Dancy, J. 'Intuitionism' in Singer, P. (ed.) *A Companion to Ethics*, Blackwell 1991 [a good introduction to intuitionist theories, including Moore]
Pidgin, C. 'Naturalism' in Singer, P. (ed.) *A Companion to Ethics*, Blackwell 1991 [a detailed discussion of the open question argument and responses to it]
Hare, R. 'Universal prescriptivism' in Singer, P. (ed.) *A Companion to Ethics*, Blackwell 1991 [Hare gives an accessible summary of his own theory]

*Internet resources*
Sayre-McCord, G. (2014), 'Metaethics', *Stanford Encyclopedia of Philosophy*, <https://plato.stanford.edu/archives/sum2014/entries/metaethics/> [a thorough and clear overview of the whole area, putting the debates here in a wider context]
DeLapp, K. (2018), 'Metaethics', *Internet Encyclopedia of Philosophy*, <http://www.iep.utm.edu/metaethi/> [a very accessible and helpful review of the main debates]
A series of good introductory talks from 'Kane B':
Introduction to metaethics: <https://www.youtube.com/watch?v=OBE50_tfAIA>
Emotivism: <https://www.youtube.com/watch?v=BIK4n3oYfrs>
Error Theory: <https://www.youtube.com/watch?v=MbTcXDMyFrA>
Non Naturalism 1: <https://www.youtube.com/watch?v=P60yrm6AWHM>
Non Naturalism 2: <https://www.youtube.com/watch?v=xr3i8oYlG_g>

# Glossary

| | |
|---|---|
| anti-realism, moral | The theory that claims that there are no moral properties. Error theory and moral non-cognitivism are both anti-realist. |
| applied ethics | The branch of ethics concerned with the application of normative ethical theories to particular issues, such as lying or stealing. |
| *arête* | An 'excellence', or more specifically, a 'virtue' – a quality that aids the fulfilment of a thing's *ergon* (Aristotle). |
| Categorical Imperative | 'Act only on that maxim through which you can at the same time will that it should become a universal law.' (Kant) |
| cognitivism | A cognitivist account of ethical language argues that moral judgements express beliefs, can be true or false and aim to describe the world. So 'lying is wrong' expresses the belief that lying is wrong, and is either true or false. |
| consequentialism, act | The theory that actions are morally right or wrong depending on their consequences and nothing else. An act is right if it maximises what is good. |
| contradiction in conception | In Kantian ethics, the test for whether we can will a maxim to become universal law can be failed if it would somehow be self-contradictory for everyone to act on that maxim. |
| contradiction in will | In Kantian ethics, the test for whether we can will a maxim to become universal law can be failed if, although the maxim is not self-contradictory, we cannot rationally will it. |
| deontology | The study of what one must do (*deon* (Greek) means 'one must'). Deontology claims that actions are right or wrong in themselves, not depending on their consequences. We have moral duties to do things which it is right to do and moral duties not to do things which it is wrong to do. |
| doctrine of the mean | Aristotle's claim that virtue requires us to feel, choose and act in an 'intermediate' way, neither 'too much' nor 'too little', but 'to feel [passions] at the right times, with reference to the right objects, towards the right people, with the right motive, and in the right way'. |
| emotivism | The theory that claims that moral judgements express a feeling or non-cognitive attitude, typically approval or disapproval, and aim to influence the feelings and actions of others. |
| end | What an action seeks to achieve or secure, its aim or purpose. |
| end, final | An end that we desire for its own sake, we can't give some further purpose for why we seek it. |
| *ergon* | 'Function' or 'characteristic activity' of something, e.g. the *ergon* of a knife is to cut, the *ergon* of an eye is to see. |
| error theory | The theory that moral judgements make claims about objective moral properties, but that no such properties exist. Thus moral judgements are cognitive, but are all false. Moral language, as we mean to use it, rests on a mistake. |

| | |
|---|---|
| *eudaimonia* | Often translated as 'happiness', but better understood as 'living well and faring well'. According to Aristotle, eudaimonia is not subjective and is not a psychological state, but an objective quality of someone's life as a whole. It is the final end for human beings. |
| felicific calculus | In Bentham's ethics, the means of calculating pleasures and pains caused by an action and adding them up on a single scale. The total amount of happiness produced is the sum total of everyone's pleasures minus the sum total of everyone's pains. |
| Formula of Humanity | A version of the Categorical Imperative: 'Act in such a way that you always treat humanity, whether in your own person or in the person of any other, never simply as a means, but always at the same time as an end' (Kant). |
| good | In ethics, what is good provides a standard of evaluation and what we should aim at in our actions and lives. |
| hedonism | The claim that pleasure is happiness and the only good. |
| Hume's 'fork' | We can have knowledge of just two sorts of claim: the relations between ideas and matters of fact. |
| imperative | A command or order. A hypothetical imperative is a statement about what you ought to do, on the assumption of some desire or goal, e.g. if you want to pass your exam, you ought to study hard. A categorical imperative is a statement about what you ought to do, without regard to what you want. |
| intention | A mental state that expresses a person's choice. It specifies the action they choose and often their reason or end in acting. |
| intuitionism | The theory that some moral judgements are self-evident, i.e. their truth can be known just by rational reflection upon the judgement itself. Moral intuitions are a type of synthetic a priori knowledge. |
| is-ought gap | Hume's claim that judgments about what ought to be the case are very different from judgments about what is the case, and cannot be deduced from them. The claim is made as an objection to moral cognitivism. |
| maxim | A personal principle that guides our decisions, e.g. 'to get a good education' (Kant). |
| means | What is done to achieve an end. |
| metaethics | The philosophical study of what morality is, enquiring into the meaning of moral language, the metaphysics of moral values, the epistemology of moral judgements, and the nature of moral attitudes. |
| mind-independent | Not depending on a mind for existence or definition. According to realism in perception, physical objects are mind-independent; according to moral realism, moral properties are mind-independent (in some important way). |
| morality | The rules, ideals and expectations governing fundamental aspects of human conduct. It concerns right and wrong, good and bad, in human action and character. |
| motive | A mental state or consideration that inclines someone to act in a certain way. Someone's motive could be a reason for acting, an end, or a desire. |

| | |
|---|---|
| naturalism, moral | A form of moral realism that claims that moral properties are natural properties. |
| naturalistic fallacy | According to Moore, the mistake of identifying moral good with any natural property. |
| nihilism | The view that there are no moral values. |
| non-cognitivism | The theory that claims that moral judgements express non-cognitive attitudes. Moral judgements do not make claims about reality and are not true or false (they are not fact-stating). |
| non-naturalism, moral | A form of moral realism that claims that moral properties are not natural properties. |
| normative ethics | The branch of ethics concerned with developing theories concerning what (e.g. which actions, which character traits, which intentions) is right or wrong, good or bad. |
| open question argument | Moore's argument that identifying the property 'good' with any other property is never correct because whether that property is, in fact, good is an open question (logically, it can receive a yes or no answer), whereas whether some property is itself is not an open question. |
| pleasures, higher and lower | According to Mill, one pleasure is higher than another if almost everyone who is 'competently acquainted' with both prefers one over the other. Higher pleasures include thought, feeling and imagination, while lower pleasures involve the body and senses. |
| practical reason | Reasons and reasoning concerned with what we can change and making good choices. |
| practical wisdom | (*phronesis*) An intellectual virtue of practical reason. It involves knowledge of what is good or bad in general and what is good in a particular situation, and the abilities to deliberate well and act on that deliberation. |
| prescriptivism | The non-cognitive theory that moral judgements are prescriptive, that is, moral judgements provide commands and recommendations about how to act. |
| property, moral | An attribute or characteristic of an object that is ethically normative, e.g. goodness or being a virtue. |
| property, natural/non-natural | Natural properties are those that we can identify through sense experience and science. Non-natural properties cannot be analysed in terms of or reduced to natural properties. |
| queerness, argument from | Mackie's argument that moral properties, understood as non-natural properties, are (metaphysically and epistemologically) puzzling and improbable, which is a reason to believe they do not exist. |
| realism, moral | The theory that claims that moral judgements are made true or false by objective moral properties that exist and are mind-independent (in some sense). |
| relativism, meta-ethical | The theory that some morality is 'relative to' a society or person, i.e. an action may be morally right in one society but not another. |

| | |
|---|---|
| rights | Justified moral demands regarding how other people may treat us, especially the freedoms (e.g. from harm) or benefits (e.g. education) they ought to provide. We are entitled to our rights in the sense that others have a moral obligation to respect them. |
| simulated killing | The dramatization of killing within a fictional context, e.g. in video games, films and plays. It is not merely the description of a killing, as in a novel, but a fictional enactment of killing that the audience or gamer can see and hear. |
| skill analogy | The analogy between virtues and practical skills. In Aristotle, virtues are said to be like practical skills because we learn both through practice (what we do), not (just) theory; and we gain an expertise that is sensitive to individual situations. |
| speciesism | Unfair discrimination on the basis of what species something belongs to. |
| theoretical reason | Reasons and reasoning concerned with what we can't change and what is true. |
| universalize | To apply to everything/everyone. |
| utilitarianism, act | The theory that only happiness is good, and the right act is that act that maximises happiness. Hedonist act utilitarianism understands happiness in terms of the balance of pleasure over pain. |
| utilitarianism, preference | The theory that we should maximise happiness, which is understood not in terms of pleasure and pain, but in terms of the satisfaction of people's preferences. |
| utilitarianism, rule | The theory that only happiness is good, and the right act is that act that complies with those rules which, if everybody followed them, would lead to the greatest happiness (compared to any other set of rules). |
| utility | The property of an object or action in virtue of which it tends to produce happiness. |
| verification principle | The principle that all meaningful claims are either analytic or empirically verifiable (Ayer). A statement is analytic if it is true or false just in virtue of the meanings of the words. A statement is empirically verifiable if empirical evidence would go towards establishing that the statement is true or false. |
| vice | A trait that is morally bad. Aristotle argues that vices are dispositions to feel or choose not in the mean, but either too much or too little. |
| virtue | Traits or states of a person that enable them to achieve some good purpose, esp. living a morally good life. Aristotle argues that virtues are traits in accordance with reason, and distinguishes virtues of intellect and virtues of character. |
| virtue ethics | The normative theory that starts from the question of what it is to be a good person, then derives an account morally right action as what a good person would do. Aristotle argues that a good person has the virtues, which enable them to achieve eudaimonia. |
| will, good | In Kant, the good will is the will that is motivated by duty, which Kant argues means that it chooses in accordance with reason. It is the only thing that is morally good without qualification. |

will, the           Our ability to make choices and decisions. Our wills are rational, that is we can make choices on the basis of reasons.