Bioethics ISSN 0269-9702
Volume 12 Number 3 1998

CONSEQUENTIALISM, REASONS, VALUE AND JUSTICE[1]

JULIAN SAVULESCU

ABSTRACT

Over the past 10 years, John Harris has made important contributions to thinking about distributive justice in health care. In his latest work,[2] Harris controversially argues that clinicians should stop prioritising patients according to prognosis. He argues that the good or benefit of health care is providing each individual with an opportunity to live the best and longest life possible for him or her. I call this thesis, opportunism. For the purpose of distribution of resources in health care, Harris rejects welfarism (the thesis that the good of health care is well-being) and argues that utilitarianism in general may lead to de facto discrimination against groups of people needing health care. I argue that well-being is a superior theory of the good of health care to Harris' opportunism. Harris' concerns about utilitarianism can be better addressed by: (i) relating justice more closely to reasons for action; (ii) by conceptualising the relationship between reasons for action and the value of the consequences of those actions as a plateau rather than scalar relationship. Justice can be understood as satisfying as many equally rational claims on resources as possible. The rationality of a person's claim on health resources turns on the strength of that person's reasons to promote certain health-related states of affairs. I argue that the strength of that reason does not track the expected value of that state of affairs in a fully scalar fashion. Rather a person can have most reason to promote some state of affairs, even though he or she could promote other more valuable states of affairs. Thus there can be equal reason for a distributor of public resources to save either of two people, even though one will have a better and more valuable life. This approach, while addressing many of Harris' concerns about utilitarianism, does not imply that doctors should give up prioritising patients according to prognosis altogether, but it does allow that patients with lower but reasonable prognosis should have a share of public resources.

[1] I would like to thank David McCarthy, Derek Parfit, Tony Hope and especially Klemens Kappel for helpful comments on earlier drafts.

[2] J. Harris, 'What Is the Good of Health Care?' *Bioethics* 10 (1996), pp. 269–291. All page references in parentheses refer to this work.

In Sheffield, England, publicly funded in vitro fertilisation (IVF) is a scarce resource and is available only to women most likely to successfully bear a live child. Roughly, a 30 year old infertile woman has a 15% chance of bearing a child with IVF, but the chance drops by 2/3 by the time she gets to 40, that is, to roughly 5%. On the basis of this observation, older women are effectively not offered IVF. While the primary motive of this distributive procedure is to achieve as good an outcome as possible (maximisation of outcome), the result is that older women are not treated. The prioritisation of individuals according to outcome goes far beyond the allocation of IVF treatment in England. It is probably the most widely used method of prioritising patients.

Such a system is unfair in one sense. Dora, a 40 year old infertile woman, could claim,

> 'I admit that my chance of becoming pregnant is only 5%, but IVF might be successful. That is my only chance. I have paid taxes for 20 years and now when I need the health service, I am denied a chance of effective treatment. Why has a 30 year old woman like Jean who happens to have 3 times my chance of becoming pregnant a greater claim on the services that I contributed to, perhaps even more than she has?'

Indeed, Dora could argue that her need for a child is the same as Jean's. Since their needs are the same, why are they being treated differently?

The theory of distributive justice which motivates much medical decision-making such as that of gynaecologists in Sheffield is utilitarianism. Utilitarianism has three components.[3] The first is an account of what is good. According to welfarism, the predominant account of the good used in health care, the value of a state of affairs is given entirely by the welfare (well-being) of individuals in that state.[4] The second component of utilitarianism is the thesis that actions are to be chosen on the basis of their consequences. This is called consequentialism. Utilitarianism is a version of welfarist consequentialism which requires simply adding up individual's well-being to evaluate consequences. This is the third component and has been called 'sum-ranking'.[5] Both because this terminology is not descriptive for non-specialists and because Harris does not himself use it, I will refer to sum-ranking versions of consequentialism as maximising versions, and all other versions as submaximising versions. Submaximising versions

[3] A. Sen, 'Utilitarianism and Welfarism' *Journal of Philosophy* 76 (1979), pp. 463–89.

[4] A. Sen, *Choice, Welfare, Measurement*, (Basil Blackwell, Oxford, 1982), pp. 28, 227.

[5] A. Sen and B. Williams, 'Introduction' in A. Sen and B. Williams, *Utilitarianism and Beyond*. (Cambridge University Press, Cambridge, 1982), p. 4.

of consequentialism include satisficing theories (according to which some states of affairs are good enough to justify promoting them, even though more valuable states could be promoted[6] and versions which give weight to the fairness of distribution of welfare. Utilitarianism is, in my terminology, welfarist maximising consequentialism. When consequences are not certain, utilitarianism requires that we should choose that course of action which maximises expected value, and it is on this basis that utilitarians would treat Jean in preference to Dora.

In what follows, I will argue that Harris advances a new conception of the good. I aim to show that this is implausible, and that the problems with utilitarianism are better understood as being based on an inaccurate relationship between reasons for action and the value of the expected consequences of action.

HARRIS' WORRIES ABOUT UTILITARIANISM

All three elements of utilitarianism – welfarism, consequentialism and maximisation – have attracted scholarly criticism.[7] One widely discussed problem with utilitarianism is that, in the pursuit of maximising welfare, it gives insufficient weight to desert, rights or fairness. It does not matter how welfare is distributed across individuals, as long distribution produces maximum welfare.[8] In this vein, Harris objects that utilitarianism fails to treat people as equals. It results in the systematic neglect of certain groups of individuals in need.[9] His arguments have been directed against the use of Quality Adjusted Life Years (QALYs) as a measure of the good of health care, but they apply equally to all forms of welfarist maximising consequentialism. Consider two of Harris' examples.

In the **life-boat example**, a naval disaster has occurred. Two ships are sinking, one with 13 crew (the big boat) and the other with 6 (the little boat). Utilitarians would save the 13. However, Harris also argues we should save the 13 (281–2).

In the **medical example**, a surgeon has 26 patients on her waiting list who will all die in a couple of weeks if not operated on. She only has time to operate on 13. If she operates on the first 13, A–M, they will all survive; if she operates on the second 13, N–Z, only about half will

[6] See below and J. Broome, *Weighing Goods*, (Basil Blackwell, Cambridge, Ma, 1991), p. 7.

[7] Sen and Williams, 'Introduction', pp. 4–5.

[8] J. Rawls, *A Theory of Justice*, (Oxford University Press, Oxford, 1972), p. 26.

[9] See for example, his: 'QALYfying the Value of Life', *Journal of Medical Ethics*, Sept 1987, pp. 117–123; 'More and Better Justice' in ed. J.M. Bell and S.M. Mendus *Philosophy and Medical Welfare*, (Cambridge University Press, Cambridge, 1988); The Value of Life, (Routledge, London, 1985).

survive. Let's say that A–M have appendicitis and N–Z have brain tumours. The surgeon is faced with the choice of saving 13 or 6 lives. A utilitarian would save A–M, those with appendicitis, since this would save the most lives and so produce the most welfare.[10]

According to utilitarianism, we should provide aid so as to maximise human welfare. One consequence of this utilitarian decision procedure is what Harris calls *de facto discrimination*.[11] The utilitarian decision procedure rules those with brain tumours out of eligibility, even though they have similar needs. Indeed, arguably, they are already worst off and suffer a double injustice. In other cases, utilitarianism can be, Harris has claimed, 'economist' (it discriminates against those whose disease is more expensive to treat), racist (it discriminates against racial groups with a poorer prognosis[12]), ageist (discriminating against the old) and sexist.[13] In the case of IVF treatment, it has been observed that lower socio-economic groups have a worse prognosis.[14] Utilitarians would seek out patients from the upper socio-economic classes. Such a system would be classist. [15]

One terminological point. Harris is right to draw attention to the implications of a utilitarian decision procedure. *De facto* discrimination is morally objectionable in one sense. Recently, the age limit for post-doctoral scholarships was increased in the England because it was recognised that the previous limit discriminated against women who take time out to have children. But there was no intention in the original policy to disadvantage women, nor was there any presumption behind this policy that women are in any way less

[10] Consequentialism can be stated in a negative form: to act so as to prevent the most harm. In the life-boat case, the choice is between preventing 13 people dying and preventing 6 people dying. Negative consequentialism requires that we prevent the 13 dying. In the medical example, the choice is between 20 people dying and 13 people dying. Negative utilitarians would prevent the 20 people dying.

[11] There are other problems with consequentialism as a theory of distributive justice. See for example, Rawls, *A Theory of Justice*., Sen and Williams, 'Introduction'. and W. Kymlicka, *Contemporary Political Philosophy*, (Oxford University Press, Oxford, 1990), Chapter 2.

[12] M. Lowe, I.H. Kerridge, K.R. Mitchell, ' "These sorts of people don't do very well": race and allocation of health care resources' *Journal of Medical Ethics* 2 (1995), pp. 356–60.

[13] 'QALYfying the Value of Life'

[14] Ian Cooke, personal communication. Michael Lockwood envisaged this general possibility in 'Quality of Life and Resource Allocation' in Bell and Mendus, *Philosophy and Medical Welfare*, p. 44.

[15] Full-blooded egalitarians could argue that Harris' own intuition that we should save the 13 in the life-boat example is another instance of *de facto* discrimination: 'little boatism' in which people are denied rescue because they happen to be in a little boat rather than a big boat, through no fault of their own.

worthy of scholarships. Racism is discrimination *based on* race. Terms like ageism, sexism and racism imply a high degree of blame. They are not justified in evaluating utilitarian attempts to secure a just distribution since the intention here is to promote all people's well-being. We should stick to the term '*de facto* discrimination.'[16]

Harris concludes that utilitarian providers and funders of health care will choose those who are easiest to treat, those with the least expensive diseases to treat and those with the best prognosis, skimming the cream of those who are ill. In the limit, those who need health care the most are most likely to be denied it.

OPPORTUNISM

Where does utilitarianism go wrong? Harris' solution is to reject welfarism. He plausibly claims that each rational person wants at least 3 things from health care: (i) the maximum possible life-expectancy *for him or her*; (ii) the best quality of life *for him or her*; (iii) the best opportunity or chance *for him or her* of getting both (i) and (ii) (270). According to his rival theory of the good which I call opportunism,

> the good of health care is a state of affairs which provides people with the best chance or opportunity to achieve what is for them the best life (iii).[17]

As Harris puts it,

> 'The fact that each person *counts* – matters morally, is recognised when their moral claims are respected, and this happens when their chances of continued life are given equal weight with the, necessarily different, chances of anyone else. (282)'

Harris retains a commitment to maximisation (281–2). He writes, 'So the maximising requirement of consequentialism is met when the claims to chances of continued life, of equal numbers of people, are given equal respect (282).'[18] While this is true, it does not capture the scope of maximisation. According to what can be called **opportunist maximising consequentialism**,

> the right course of action is that course of action which **maximises** the number of persons who will receive the opportunity to realise

[16] Thanks to Derek Parfit for this example and observation.

[17] While strictly speaking, it is these states which provide opportunity which are good, for brevity I will write as if opportunity itself were the good of health care.

[18] This passage illustrates that Harris assumes that all versions of consequentialism are maximising. On my taxonomy, this is not the case.

what is **the best life for them** (given their situation and all possible courses of action).

In the medical example, an opportunist maximising consequentialist is faced with a dilemma: giving 13 people (A–M) their chance of the best life or giving a different 13 (N–Z) their chance. Harris argues that the surgeon should not prefer those with appendicitis, simply based on prognosis. He raises the possibility of treatment on a 'first come, first served basis' (274,282). In the life-boat example, an opportunist maximiser should save the 13 (281–2), since this gives the greater number of people their opportunity to live. This example establishes Harris as a maximiser.

In my experience, Harris' intuitions about these two examples are shared by many people. Opportunist consequentialism is consistent with these intuitions. However, I will presently examine these intuitions in detail and offer a more plausible alternative.

Opportunist maximising consequentialism, like utilitarianism, has some counterintuitive implications. For example, it requires giving those with brain tumours the same chance of treatment even if they only have a one in a million chance of cure or if a cure would only result in them living for a few more days. That is absurd.

It is also hard to square Harris' opportunism with his rejection of economism. Imagine that, in the medical example, performing an appendicectomy costs $1000 whereas removing a brain tumour costs $2000. If there were 13 people with each disease but only $13,000 in public resources, saving those who happen to have appendicitis is like saving those who happen to be on the big life boat in the life boat example: given the constraints of our situation, we can give more people their opportunity to live if we go for those with appendicitis. Harris' rejection of economism prevents him from giving preference to those with appendicitis if that disease is cheaper to treat, but I can see no morally significant difference in these examples between the cost of one's disease (which allows distributors of public resources to put that person in one of two groups in which either more or less people are saved) and whether one happens to find oneself be in a big or little boat, with more or less people.

Both welfarist and opportunist maximising consequentialism have counterintuitive implications. However, there are more principled reasons to reject Harris' opportunism.

IS THE GOOD OF HEALTH CARE OPPORTUNITY?

How should we interpret Harris' claim that the good of health care is opportunity? One obvious interpretation is that giving people an

opportunity is a constraint on distributive schemes aimed at maximising welfare. That constraint is necessary to promote a just and fair society. It is good *that* people have opportunity. Harris in the past has certainly said this kind of thing. He wrote,

> 'Equality requires both that we treat as many people as we can [the maximising element] and that we ensure so far as possible that certain sorts of people be not systematically ignored.'[19]
>
> '[W]e have two equally plausible moral principles here [maximising welfare and equality] and that pull in opposite directions. When this is true some means of doing justice to each must be found and it is not enough merely to opt for one.' [20]

This, however, is not very new and many critics of utilitarianism have given variants of this argument. Moreover, Harris himself seems to be saying something new, and not merely restating his old arguments. If he is saying something new and interesting, I think we should understand him literally as saying that the good of health care *is* opportunity, not welfare, and that opportunities are to be distributed among people, as utilitarians distributed welfare. This squares with his talk of the 'benefit' of health care, 'conceptions of the good', and his apparent sympathy for some maximising version of consequentialism, albeit a different one to that employed by clinicians and health economists. Thus Harris concludes, 'While it is true that funders and providers might legitimately wish to take into account the amount of benefit that their money and/or efforts will provide, discounted by the probability of that benefit being achieved. I have argued that there is another perspective to consider and another interpretation of what 'benefit' legitimately means.' (290; see also 274, 281–2)

There are three ways in which Harris' claim that opportunity is literally a good can be interpreted. He might mean that opportunity is good in itself (that is, an intrinsic good) or is a means to other states of affairs which are good in themselves (that is, opportunity is an instrumental good), or both. An example of an intrinsically good state is being happy; an example of an instrumentally good state is having money. Some good states like having knowledge may be both intrinsically good and a means to other intrinsically good states like being happy. In this section I will argue that on any of these

[19] 'More and Better Justice', p. 95.

[20] *Ibid.*, pp. 94–5. This is in response to a similar challenge from Michael Lockwood ('Quality of Life and Resource Allocation', p. 54). This is like John Broome's suggestion that we should trade some good for fairness ('Good, Fairness and QALYs' in Bell and Mendus, *Philosophical Medical Welfare*.)

interpretations, Harris' intuitions about these cases cannot be justified.

The claim that opportunity has instrumental value is most easily dealt with and clearly will not justify Harris' intuitions. On this reading, we should interpret the claim that opportunity is good to mean that it has value as a means to prolonging a good life. But if it has only this instrumental value, then what is ultimately good is a good life. Let's give every person's good life the value of 1, with death being 0. Operating on those with brain tumours realises 6 units of value while operating on those with appendicitis realises 13 units. We should operate on those with appendicitis.

However, Harris is best understood as suggesting that opportunity is of value in itself, of intrinsic value.[21] If opportunity is an intrinsic good, then this would distinguish the life-boat from the medical example. When we attempt rescue or medical treatment, we provide opportunity, regardless of outcome. If providing a person with an opportunity counts as value 1, and not providing opportunity is 0, then if we rescue the 13 people in the naval disaster, we provide 13 units of value; if we rescue 6 people, we provide 6 units. If we are to choose that option with the best consequences, we should rescue the 13. If we operate on the 13 patients with brain tumours, we provide 13 units of value; if we operate on the 13 patients with appendicitis, we provide 13 units of value. There is no reason to prefer those with appendicitis over those with brain tumours.

But can opportunity be an intrinsic good? Entering a lottery is a means to our ultimate ends, such experiencing pleasure, achieving worthwhile things with the money and so on. Losers of lotteries often console themselves by pointing out that money did not make the winner any happier. Watching the lottery wheel spin may be fun, but the value of this experience lies in the pleasure it provides, and not in the mere participation in a lottery itself. Chance or opportunity is not of objective intrinsic value.

Some would deny this and claim that there is value in just having a chance. However, Harris himself rejects the concept that chance or opportunity has an objective intrinsic value, as his commitment to the Argument against Potentiality implies. This argument is deployed in the abortion and other debates, and goes something like this:

Premise 1. Potential persons have the same rights as persons.
Premise 2. The fetus is a potential person.

[21] Utilitarians can attribute intrinsic value to both the outcomes of an act and the act itself (S. Scheffler, *The Rejection of Consequentialism* (Oxford University Press, Oxford, 1982), pp. 1–2; Broome, *Weighing Goods*, p. 4).

So, the fetus has the same rights as a person (including a right to life).[22]

Harris has rejected the first premise. '[T]he bare fact that something will become X . . . is not a good reason for treating it now as if it were in fact X. We will all inevitably die, but that is . . . an inadequate reason for treating us now as if we were dead.'[23] Not only does Harris believe that potential persons do not have the same rights as persons, he also believes that they do not have the same value as persons and on this grounds it is not wrong to kill fetuses.[24] In this context, potential is just another word for chance or opportunity. Harris could not make the claims he does regarding the value of the fetus' life if potential or chance or opportunity to realise something valuable had substantial value in itself.

There is one further way in which opportunity may have intrinsic value. Rather than having an objective intrinsic value, it may have subjective intrinsic value.

VALUING OPPORTUNITY

Harris asserts that a person's life matters 'not because it is a life, but because it is *someone's* life, because her life is an enterprise in which she has, and takes, an interest (282).'[25] The suggestion here is that life *per se* is not of value, but of value to the extent that a person values it.[26] As we have seen, this strategy will not justify Harris' intuitions in the medical example because, even if every patient's life has value of 1, we must discount the value of operating on patients with a brain tumour by the probability that the operation will not achieve that valued outcome.

However, the implication of these claims is that each person *values having the operation to the same degree*. On this view, opportunity is good to the extent that it is valued. Thus, if Tom, who has a brain tumour,

[22] *The Value of Life*, p. 11. Strictly, this example is not relevant. It demonstrates that the fact that something will lose value does not justify treating it as if it has already lost that value. However, we are interested in the claim that because something will have a certain value it should be treated now as if it has that value.

[23] *Ibid.*

[24] *Ibid.*, p. 159.

[25] This is similar to Kamm's suggestion that 'we count equally each individual's preference, understood not as the *object* of his preference but as the fact *that he prefers it*' and that the fact that he prefers some state of affairs should make a difference in the process of deciding whether to bring that state of affairs about (F.M. Kamm, 'Equal Treatment and Equal Chances', *Philosophy and Public Affairs*, 14 (1985), pp. 177-194, esp. p. 181).

[26] *The Value of Life*, Ch. 1 and 2.

wants the opportunity to live just as much as Alex, who has appendicitis, this grounds an equal moral claim, regardless of how great their chances are. The relevant difference between the life-boat example and the medical example is that in the former, the choice is between respecting the values of 13 people or 6 people, while in the latter the choice is between respecting the values of 13 people or a different 13 people.

There are several problems with this account of the goodness of opportunity.

(i) It is not the way we value opportunity.

To value opportunity or treatment itself as an end and not as a means would be quite unusual. According to this view, we might value holding a ticket in a lottery not because of the money we might win, but simply for the chance to participate in the lottery. An extreme example of this pattern of concern would be a person who valued entering lotteries, but did not care about the result at all. Seeing that he has won, he walks off, not interested in collecting the prize money. Such a person, perhaps suffering from a psychological disorder, might be said to have missed the point of entering a lottery.

I myself have difficulty understanding how opportunity could have intrinsic value or how one could intelligibly value it for its own sake. Nonetheless, people do value some pretty bizarre things and part of the value of participating in risky sports may be taking risk.[27] And doctors sometimes justify the disastrous result of some medical adventure by saying, 'At least he was given a chance.'

(ii) Present Preferences

On Harris' view, the goodness of a state of affairs is a function of people's desires for that state of affairs. But which desires? Desires we now have or desires we will have? According to a principle of temporal neutrality,[28] if the value of a state of affairs is determined by our desires for that state, we should appeal not only to what people now desire, but also what they will desire. And if we consider future preference satisfaction, we must discount the value of that satisfaction

[27] Though more plausibly the contrast with death enhances the value of life. The psychological heuristic of contrast is described in D. Kahneman and C. Varey, 'Notes on the psychology of utility,' in ed. J. Elster and J.E. Roemer, *Interpersonal Comparisons of Well-Being*, (Cambridge University Press, Cambridge, 1991), pp. 127–63.

[28] H. Sidgwick, *The Methods of Ethics*. (Macmillan, London, 1963), p. 111; T. Nagel, *The Possibility of Altruism*. (Clarendon Press, Oxford, 1970), pp. 60, 72.

by the probability of it not occurring. Thus, the value of operating on patients with brain tumours must be reduced by 50% in Harris' medical example.

Harris must reject this kind of temporal neutrality. On his view, what is good is not preference satisfaction *per se*, but present preference satisfaction. Thus Harris' consequentialism is different from preference utilitarianism which counts all preferences, across all times. I note in passing that this theory shares many formal features with the Instrumental and Deliberative Versions of Derek Parfit's Present-aim Theory of reasons for action. According to the Instrumental version, what each of us has most reason to do is whatever would best fulfil our present desires.. Parfit himself goes on to reject this view, as I will suggest we should, in favour of an objective theory, the Critical Present-aim Theory, which I will describe presently.

(iii) Other Present Preferences

But let us assume for argument's sake that the only relevant preferences are present preferences. Even if opportunity has intrinsic value to the extent that it is valued, life itself, if of a certain quality or kind, is also surely valued. So, on a subjective account of value, life is also of value. If life is also presently valued, the moral imperatives derived from the need to save life would direct us, as I have argued, to save those most likely to live.

To illustrate, imagine that we accord equal weight to valued opportunity and valued life. If we treat 10 people with brain tumours, we expect to realise 10 units of opportunity value and 5 units of life value. If we treat 10 people with appendicitis, we realise 10 units of opportunity value and 10 units of life value. We should give priority to those with appendicitis.

(iv) Subjectivism

One significant problem with Harris' view that an outcome is valuable to the extent that individuals value it is that it is a subjectivist account of value. On this view, if an individual does not want to live, her life has no value, and it might not be wrong to kill her. While this may be true if there is some objective reason for her not wanting to live, such as she is racked with incurable pain from imminently terminal cancer, in the absence of such an objective justification, it is absurd to suggest that because someone happens to want to die her life is not of value. Subjectivists have responses to this argument, and I will only signal that I do not believe that a subjectivist account of

value is plausible and that objectivist alternatives are practicable.[30] Moreover, there are possible responses to my preceding arguments against the claim that opportunity has value because it is valued. Some of these turn on what it means to treat an individual as an end in himself or herself.[31] For present purposes, the preceding objections are sufficient to call into question the concept of the good of health care as opportunity.

BEYOND OPPORTUNISM

If we do reject the concept of the good as opportunity, there are at least two moves open. One is to give up commitment to maximisation and go for some submaximising form of consequentialism, such as Harris has described in the past, or perhaps even nonconsequentialism. Another strategy would be to further revise the conception of the good. Harris could claim that both welfare and fairness are good. On this approach, a plausible form of consequentialism involves weighing both these goods.[32]

There is much more to be said for these alternatives than I can say here, and there are many ways understanding what it is to treat a person as an equal.[33] But let me signal that I have some doubt that giving weight to fairness as Harris has conceived of it will square with Harris' intuitions about the life-boat and medical examples. Harris elsewhere claims that 'the equality principle demands that each person be given an equal chance of benefiting from health care.'[34] If Harris is to be taken literally, giving significant weight to fairness and equality requires that we give everyone an equal chance of being saved. Fairness, on this view of equality, requires that in choosing whether to save those in the big boat or the small boat, rescuers should toss a coin.[35]

[29] D. Parfit, *Reasons and Persons*, (Clarendon Press, Oxford, 1984), p. 117.

[30] J. Savulescu, 'Rational Non-Interventional Paternalism: Why Doctors Ought to Make Judgements of What Is Best for Their Patients', *Journal of Medical Ethics* 21, 6 (Dec 1995), pp. 327-31 J. Savulescu, 'Liberal Rationalism and Medical Decision-Making', *Bioethics*, 11 (1997), pp. 115–129.

[31] F.M. Kamm, *Morality, Mortality*, Part I (Oxford University Press, New York, 1993).

[32] Broome, J. *Weighing Goods*, Chapter 1.

[33] See for example the works by Kamm previously cited.

[34] 'More and Better Justice,' p. 86. This view is shared by Dan Brock ('Ethical Issues in Recipient Selection for Organ Transplantation' in ed. D. Mathieu, *Organ Substitution Technology: Ethical, Legal, and Public Policy Issues*, (Westview Press, Boulder, Colorado, 1988).

[35] J.M. Taurek, 'Should the Numbers Count?' *Philosophy and Public Affairs*, 6 (1977), pp. 293–316. Kamm's response to this paper in both the works cited outlines

I will leave both these possibilities open. I will not address the non-consequentialist literature around equality as respect for persons.[36] Rather, I will offer a version of maximising consequentialism which I believe is closer to what Harris has in mind, and which squares with Harris' intuitions about these cases. But first we need to discuss the relationship between reasons for action and value, and between reasons for action and distributive justice.

REASONS AND JUSTICE

A reason for acting is a fact or circumstance forming a sufficient motive to lead a person to act. Knowing a person's reasons allows us to understand why a person acted as he did. These are explanatory or motivating reasons. For example, a person's reason for buying a lottery ticket might be the fact that he believes that this is his lucky week and he wants to win a large sum of money. Reasons for acting can be good or bad. For example, 'His reason for removing the pollution control device on his car was to reduce petrol consumption, but that wasn't a good reason to do that.' A reason for action is good if it meets a standard, that is, if it conforms to a set of norms governing that behaviour. Good reasons for action are called normative or justifying reasons for action. In what follows, I will only consider normative and not motivating reasons for action.

What is the relationship between a person's normative reasons and her entitlement to health resources? To have a *prima facie* entitlement to health resources, the state which there is reason to promote must be some relevant health-related state, like regaining sight or being free of pain. But it requires more than this.

Justice is concerned with providing what there is good reason to provide for people. Let's say that a person has a rational claim to have some state of affairs, p, promoted if there is good reason to promote p.

several other non-consequentialist procedures such as various gambles and majority rule, which claim to treat people as equals without necessarily according everyone an equal chance of being saved.

[36] In particular, her principle of majority rule (*Morality*, *Mortality* pp. 116ff and 'Equal Treatment and Equal Chances') would justify Harris' intuitions about the life-boat and medical case, if the interests of all those involved in receiving treatment or being saved were the same. However, I have questioned whether the interest in treatment of a patient with a brain tumour (who has a 50% chance of surviving with treatment) is the same as the interest of a person with appendicitis (who has a 100% chance of surviving). I will argue that their reasons for action are the same strength, though the expected value of those actions is different. Kamm herself gives some weight to outcome (at least length of life – 257–260) in the distribution of scarce medical resources. She does not directly address prognosis or chance of good outcome.

According to one version of consequentialism,

C1. the good of health care is satisfying a rational claim for some health-related state.

C2. the right distribution is that distribution which maximises the number of people whose equally rational health-related claims are fully satisfied.

Call this view reasons-based maximising consequentialism, or reasons-based consequentialism for short.[37] According to reasons-based consequentialism, the following claims are true:

C3. If a person (including a distributor of public resources) has equal reason to promote p, q or r, and that person can promote either p and q, or r, then he or she should promote p and q.[38]

C4. If a person, A, has the same strength reason to promote p as another person, B, has to promote q, there is as much reason to promote p as q. Thus, if A has the same strength reason to promote p as B has to promote q, then a distributor of public resources, X, has as much reason to promote p as to promote q.

C3 applied to the life-boat example implies that we should save the 13 rather than the 6, assuming there is equal reason to save each person's life.[39] C3 seems obviously true.

C4 is less obviously true, though I believe it is true. An example of the principle properly specified is that if A (who is a professional footballer) has as much reason to have a knee reconstruction as B (who is a professional tennis player) has to have an elbow reconstruction, then there is as much reason for the distributor of public resources to provide A with resources for a knee reconstruction as there is to provide B with those needed for an elbow reconstruction. The idea is that reasons have a force which is determined by the particular set of circumstances and apply to anyone in a relevantly similar situation. The return of normal function of a joint which is crucial for performing properly in a professional capacity provides a

[37] If duties provide reasons, this becomes a very broad reading of consequentialism which encompasses much of deontology. For a similarly broad interpretation, see D. Sosa, 'Consequences of Consequentialism', *Mind*, 102 (1993), pp. 101–22.

[38] Provided, of course, that there are no negative interactive effects between p and q. Both C3 and C4 concern only the agent-neutral component of reasons, as we shall see.

[39] Our intuitions about this example would change if the 13 were patients with terminal cancer expected to die in the next week and the 6 were healthy.

reason, and that same reason is what should determine allocation of resources to A and B.

C4 is easily misinterpreted and its scope over-extended. Taurek, for example, in an often-cited work, argues that if it is morally permissible to save oneself rather than 5 strangers, then it is morally permissible for another person to save one stranger rather than 5 strangers. Taurek might be interpreted (erroneously) as appealing to something like C4: if (1) there is at least as much reason for A to save A as there is for A to save B-F, then (2) there is at least as much reason for a distributor of public resources, X, to save A as there is to save B-F. Claim (2), however, does not follow claim (1)

Taurek's version of consequentialism involves evaluations of outcomes relative to individuals' own interests (it is permissible for A to save herself rather than 5 strangers), combined with an agent-neutral conception of reasons. An agent-neutral reason applies to any agent in relevantly similar circumstances. Agent-neutral reasons can be specified without making essential reference to the agent.[40] An agent-relative reason applies to some agents in virtue of their relationship with the state to be promoted and make essential reference to the agent. For example, I may have a reason to save X rather than Y and Z because X is my child. However, you may have reason to save Y and Z because they are all strangers to you. Your reason is agent-neutral whereas mine is agent-relative.[41]

Parfit objects that Taurek is really discussing agent-relative reasons when he claims to describing about agent-neutral reasons.[42] The only conclusion one can draw from the claim that A can have as much reason to save himself as 5 strangers is that anyone can have as much reason to save himself as 5 strangers. Thus, at most, all that follows from Taurek's claim (1) that there is as much reason for A to save A as for A to save B-F is (2*) that there is as much for

[40] T. Nagel, *The View from Nowhere*, (Oxford University Press, New York, 1986), p. 153.

[41] M. Smith, *The Moral Problem*, (Blackwell, Oxford, 1994), p. 169.

[42] D. Parfit, 'Innumerate Ethics', *Philosophy and Public Affairs*, 7 (1978), pp. 285–301 at p. 287. Parfit uses the term agent-neutral in a later work (*Reasons and Persons*).

This distinction is in some ways unhelpful. All reasons are relative, in that they are relative to the relevant features of the circumstances including relevant features of the agent and his or her relationships. However, all reasons are agent-neutral in that they apply to any agents (irrespective of identity) in those circumstances. Thus if I have reason to save my child rather than two strangers, any father (in relevantly similar circumstances) has reason to save his child rather than two strangers. If you have a reason to save two strangers rather than one stranger, any person has a reason to save two strangers rather than one stranger. These reasons are both agent-relative and neutral.

a distributor of public resources to save himself as there is for him to save B-F.[43]

My claim, C4, refers to the agent-neutral component of a person's reasons. In the example of joint the reconstructions, what I am claiming is that if there is a reason of strength R to provide A with a knee reconstruction and there is a reason of strength R to provide B with an elbow replacement, then there is equal reason for a distributor of public resources to provide A with a knee reconstruction as there is to provide B with an elbow reconstruction. These reasons apply in virtue of the suffering or disability each experiences.[44]

C4 applied to the medical example would support Harris' intuition that we should not prefer those with appendicitis over those with brain tumours if there is as much (agent neutral) reason to treat a person with appendicitis as there is to treat a person with a brain tumour. Is there as much reason to treat patients with brain tumours as there is to treat patients with appendicitis? How are reasons for action related to the expected value of the consequences of that action?

REASON AND VALUE

Harris' move to opportunism was motivated, in part, by (widely-held) intuitions about cases like the medical example. However, I have argued that this thesis has counterintuitive implications and is an implausible account of value. Harris was close, but not spot-on in his diagnosis. A better solution is based on a new understanding of the relationship between reasons and value.

The relationship between reasons and value which is assumed by most discussions of reasons for action is a scalar one: increases in the expected value of action result in roughly linear increases in the strength of our reasons to perform those actions. That is, the more good an action would achieve, the more reason (the stronger) there is to perform the action.[45] On this account, those with brain tumours have less reason to seek out operation than those with appendicitis because they have half the chance of achieving the good outcome (prolonging a good life).

[43] Other things being equal, which they would not typically be because distributors of public resources have special duties to B-F which other individuals would not have.

[44] There may of course be other agent-relative reasons but I am not referring to these.

[45] I am not distinguishing in this paper between the strength of reason, the rationality of a reason and the amount of reason.

However, reasons and value may be related in a different, non-scalar way. The relationship may be of a plateau kind such that the strength of reason to act increases as the value promoted by that action increases, until some plateau is reached where strength of reason no longer increases despite increments in value. Thus a person may have most reason to perform some act, even though other actions would promote more value, if the consequences of the chosen act are good enough.

I have argued elsewhere[46] that the relationship between reasons and value is a plateau type on Derek Parfit's Critical Present-aim Theory (CP). The central features of this theory are:

1. each of us has most reason to satisfy his set of rational present desires.
2. a set of rational present desires includes those desires we would have if we knew the relevant facts and were thinking clearly.
3. all intrinsically irrational desires are excluded from this set. An intrinsically irrational desire is a desire which is in no sense worth achieving.[47]
4. all rationally required desires are included in this set. A rationally required desire is a desire which each of us has reason to cause to be fulfilled, whether or not we actually have this desire.[48]
5. the set of desires is itself not irrational (e.g., no inconsistent or intransitive preferences).[49]

Elsewhere,[50] I have argued that according to the Critical Present-aim Theory:

- for a choice or act to be rational, the state of affairs promoted by that choice or act must be worth promoting. That is, it must promote some objectively valuable state such as well-being, achievement, knowledge, justice, and so on.
- the state of affairs promoted must have an expected value which is good enough relative to other available alternatives.
- we are not rationally required to give up a concern for one objectively valuable state which is good enough for a relevantly

[46] J. Savulescu, 'The Present-aim Theory: A Submaximizing Theory of Reasons?', *Australasian Journal of Philosophy*, forthcoming.

[47] Parfit, *Reasons and Persons*, p. 122.

[48] *Ibid.*, p. 131.

[49] *Ibid.*, p. 119. Framing CP in terms of desires and aims is potentially misleading. CP is an objective theory of reasons for action. What generates a reason is the objective value of the object of that aim.

[50] Savulescu, 'The Present-aim Theory: A Submaximizing Theory of Reasons?'

different state which is more valuable. Some present rational concerns are good enough.

Thus, I said that CP is a 'submaximising' theory of reasons for action, that is, a theory which allows that a person can have most reason to act in some way even though other actions would realise more value. Submaximisation has been proposed by some philosophers as being rational and morally acceptable,[51] but, as Harris notes, it is generally argued to be irrational.[52] I have here described the relationship between reasons and value as plateau rather than submaximising to avoid confusion with a submaximising theory of justice. These are different, and one can be (as I am) a submaximiser about individual reason for action but a maximiser about justice.

The Present-aim Theory, when interpreted this way, gives some weight *to what agents now actually care about*.[53] Can an individual have most reason to promote a states of affairs which has less value than other states of affairs which she could promote?

I have given arguments elsewhere[54] that she can. Here is a one example. Imagine that Peter's wife, Andrea, becomes an alcoholic. If Peter stayed with her, he could help her, and their relationship would be good in some ways. However, it would be a difficult life, and they would probably not be able to bring up children together. If he left her, he could love an old friend, Mary. They would have a rich and happy relationship, and be able to have and care properly for children. Loving Mary would likely produce more value for him and overall. If Peter chose to stay with Andrea, his justification would be, in part, that Andrea happens to be the woman he now actually cares about.

Let's assume that, if Peter left Andrea for Mary, Peter's pattern of concern would change and he would care most for Mary. If Peter is rationally justified in staying with Andrea, there is at least as much reason for Peter to love Andrea as there is for him to love Mary. My claim is that Peter is rationally justified in staying with Andrea, if she is the person he most cares about.

[51] M. Slote, 'Satisficing Consequentialism', *Proceedings of the Aristotelian Society*, Suppl. 58 (1984) pp. 139-63. M. Slote, *Common-Sense Morality and Consequentialism*, (Routledge and Kegan Paul, London, 1985). M. Stocker, *Plural and Conflicting Values*, (Clarendon Press, Oxford, 1990), Part IV.

[52] P. Pettit, 'Satisficing Consequentialism', *Proceedings of the Aristotelian Society*, Suppl. 58 (1984), pp. 164–76.

[53] For a number of examples, see Savulescu, 'The Present-aim Theory: A Submaximizing Theory of Reasons?'

[54] *Ibid.*

FROM REASON AND VALUE TO JUSTICE

If the relationship between reasons and value is a plateau one, this has important implications for distributive justice. Imagine that only the following conditions obtain:

1. there is a distributor of public resources who can promote either p or q, but not both.
2. A has at least as much reason to promote p (because that is what she now cares about) as she would have to promote q (if that were what she cared about).
3. A could change her pattern of concern to care about q (A*).
4. q is more valuable than p.

If A can have most reason to promote p if that is what she most cares about, then according to C4, the distributor of public resources can have most reason to provide A with the resources so that p rather than q is promoted, if that is what A wants, and this does not conflict with other people's rational claims.

Here is a medical example. Imagine that A must have an operation on his spine. If the operation is performed one way, there will be no damage to the nerves to his legs but A will certainly be impotent. If another operation is performed, A will not be impotent but there will be a small chance that he will be left paralysed. A is 48 years old and values his potency. He chooses to have the operation which will preserve his potency. Now it may be that it is better to be impotent than paralysed by a long way (let's assume that it is), and that this choice does not maximise expected value. Nonetheless, it may be the choice A has most reason to make. If that is so, distributors of public resources should not require him to have the operation which will avoid paralysis with certainty.

Distributors of public resources should not require that agents change what matters most to them, provided that the object of that pattern of concern is worth achieving and good enough relative to other alternatives. But, *a fortiori*, agents should not be required to give up altogether what matters most to them (unless it is not worth achieving) for the sake of what matters to *others* in order to maximise value.

Imagine now that the following conditions obtain:

1. there is a distributor of public resources who can promote either p or q, but not both.
2. A has most reason to promote p and B has most reason to promote q.
3. q is more valuable than p.

If the preceding argument is correct, the distributor can have as much reason to provide A with the resources so that p is promoted as she can to provide resources to B to promote q. To use the preceding example, if A and B both have the same disorder, and A prefers to have the potency-preserving operation, and B prefers to have the paralysis-avoiding operation, distributors should not prefer B to A, even if the expected value of B's choice is greater.

C4 implies:

> A distributor of public resources can have the same strength reason to promote A in p (if that is what A most cares about) as to promote B in q (if that is what B most cares about), even if q is more valuable than p.

There are limits to this principle. When the expected value of one option greatly outweighs the expected value of another option, we are rationally required to choose the former. Thus while it may be up to A to choose between the potency-preserving operation and the paralysis-avoiding operation, distributors of public resources are not required to provide resources for A to have a herbal therapy of no benefit, even if A strongly desires this. Thus, the difference between the expected value of treating A and the expected value of treating B must be below some threshold or limiting amount, or else there is more reason to treat B.

If C4 and its implications are true, these together would justify Harris' intuitions about the medical example. Can we find other support for it? Consider a related but slightly different example. Two 70 year old men have cancer and will die without treatment. Each loves his family dearly. In each case, the man's family is grown up, but they are poor. Each old man is trying to decide whether to use all his remaining assets to pay for a new experimental treatment. If he spends his remaining assets, his family will be worse off, and this weighs heavily with him. With treatment:

> Man A has a 1/50 chance of survival.
> Man B has a 1/100 chance of survival.

Each asks, in a state of genuine uncertainty, 'What should I do?' Let's assume that each decides to take a chance on life.

Now my intuition about this case is that there is as much reason for each man to choose a chance on life rather than his family's welfare. Each has the same reason to bequeath money to his family. So, the reason giving force of a chance of continued life is the same for each, even though the chances are different.[55]

[55] An alternative explanation of these intuitions is that in each case the chance of life is so much greater in value than providing for one's family that setting aside

If you do not share the intuition that each is equally rational in choosing a chance of life, change the relevant probabilities to 1/50 and 1/51. At some point, the probabilities are so close that any difference is not relevant to the old man's reasons in this situation. There is a range of probabilities which are so close to 1/50, say, that the difference is not relevant to rational deliberation. For example, it may be down to 1/55, or 1/60, or 1/100, or 1/200. The point is that reasons are not so fine grained as to be sensitive to small changes in expected value.

Compare these men to Man C. He has the same disease, the same assets and the same concern for his family. However, he has a 1/1 000 000 chance of survival. Man C is relevantly different to Men A and B. If Man C cares greatly about his family's welfare, he should not spend his money on the experimental treatment. The expected value of the operation is so small that the strength of his reason to have the operation is weaker.

Importantly, whether the expected value of a course of action is good enough to justify performing that action depends on the expected value of the alternatives. If Man C did not have a family, he would have most reason to spend his money on the treatment. But not, I am suggesting, as much reason as A and B. In deciding whether to distribute resources to Man A, B or C, I am claiming that there is as much reason to provide resources to Man A as there is to Man B, but less reason to treat C.

To return to Harris' medical example, whether there is equal reason to treat those with appendicitis as those with brain tumours turns on whether individuals with appendicitis have as much reason to seek treatment as individuals with brain tumours. Here, Harris' intuitions may be justified. It seems plausible to suggest that if operation for brain tumour had a 50% chance of success then those with brain tumours have as much reason to seek treatment as those with appendicitis. They could plausibly say that it is their only chance at their own life, and 50% success is enough.

What are the limits to this principle? When is one option of lesser value good enough compared to other options? Kamm suggests that '[o]nly equal or approximately equal individual interests or rights should be matched against each other in deciding who or what may

one's family tells us nothing of the relative reason-giving force of the chosen alternative. But this does not seem to me to be a case of this kind. In each case, the alternative of bequeathing the money to the family is roughly of similar reason-giving force to taking a chance on prolonging one's own life. That is why it is a dilemma for each man.

be a contestant for a good'.[56] She notes that 'the fact that someone values his own cut-off toenail as much as someone's life is not thought . . . to make it morally acceptable for him not to give up the toenail to save a life.'[57] Kamm describes this as an objective constraint on the subjective weight we give to our own interests.

Taurek, for example, suggested that it would be permissible to prevent one person losing her arm rather than prevent one person dying.[58] Parfit, in what has been taken to be an endorsement of maximisation (289), denies this claim.[59] He and probably Harris (289) believe that there is more reason to save lives than limbs. Some people, however, choose to die rather than lose a limb. Is a person who chooses an operation for some serious spinal condition with a 10% chance of death in preference to another operation with a 10% chance of loss of limb irrational? It is hard to see that he is.

Kamm draws the line differently to Parfit and Harris. She claims that it would be rational for a person to give more weight to his own legs than another person's life, though (citing an observation of Parfit's) that same person would be rational if he chose to give up his legs to save another's life. My claim is that saving legs and preserving life both provide an equal reason for action, and are good enough relative to each other to ground an equal claim to health care resources. While Kamm believes that the loss a person's legs can be compared with the loss of a life, she believes that the loss of an arm cannot be compared with the loss of a life.[60] There is more work to be done in determining which options are good enough compared to available alternatives.

There are, of course, many situations arising in real life which are outside the scope of C2 and reasons-based consequentialism. What

[56] Kamm, *Morality, Mortality*, p. 148.
[57] *Ibid.*, p. 153.
[58] Taurek, 'Should the Numbers Count?'
[59] Parfit, 'Innumerate Ethics'.
[60] Kamm, *Morality, Mortality*, p. 154. She states that 'an arm and a life do differ too radically.'

While Kamm's approach is non-consequentialist, mine is consequentialist. Our views may differ in other ways. On Kamm's preferred analysis of interests, subjectivism, for the losses of A and B to be comparable, what A would lose must be as important to A as what B would lose would be to B (and these losses are roughly objectively comparable, *Ibid.*, p. 154). On Kamm's view, if A happens to care less about the loss of his legs than B (say, he is more stoic), though A still most wants to keep his legs, there is a reason for a distributor of public resources to give preference to B because B's loss is more important to him. On my view, the loss of one's legs in relevantly similar circumstances provides an equal reason for action (provided that a person cares more about that loss than other alternatives). Our views differ on the life-boat case, with Kamm's preferred analysis of interests, subjectivism, requiring that everyone be given equal chances or proportional chances (*Ibid.*, p. 156).

should be done when there are insufficient resources and we have a group of people who have equally rational claims which can be either partially or fully satisfied? For example, we could save A for 10 years, or B for 10 years, or A and B each for 6 years. When outcomes are comparable, like being saved for 6 or 10 years, we should defer to individual autonomy: would A (and B) prefer a 50% chance of surviving 10 years or a certainty of 6 years? The question here is whether a 50% chance of 10 years provides the same reason for action as a certainty of 6 years, and this is, I believe, a matter of individual judgement, which should be accommodated as far as possible.

Secondly, what should be done when not all individual claims on resources are equally strong? For example, we could save A for 10 years or B–Z for 6 months. Here there are at least two alternatives:

1. the priority view: we should satisfy those first whom individually there is most reason to save.[61]
2. the additive view: would be to give some but less weight to less rational claims and these claims can be summed together.

I am attracted to the priority view but I will not argue for this here.

OPPORTUNIST OR REASONS-BASED CONSEQUENTIALISM?

Reasons-based consequentialism avoids one absurd implication of Harris' argument: that any chance of good life, no matter how small, justifies a claim on resources. Thus, in the medical example, if those with brain tumours had only a 0.1% chance of survival, it would be appropriate to give preference to patients with appendicitis. That chance is too small compared to the certainty of saving patients with appendicitis.

It is worth noting that reasons-based consequentialism will not avoid one of Harris' criticisms of utilitarianism: the charge that it is economist. But nor, as I previously noted, will Harris' preferred solution: opportunist consequentialism. Economism cannot be avoided if we aim to maximise the satisfaction of the legitimate claims which people have on public resources. The cost of treatment may put a person in a group, in the same way as happening to find himself in a

[61] Is this still a version of consequentialism? I am not sure. It may be that the satisfaction of rational claims of differing strength are non-comparable goods and that consequentialism can accommodate a lexical priority to goods. That is, if A and B are non-comparable goods, and A has lexical priority to B, we should maximise A before we maximise B.

small life-boat puts a person in a group, a group in which, because of the circumstances, there is less reason for distributors of public resources to save than other groups.

FINAL REMARKS

Harris argues that utilitarianism unfairly neglects the claims of some people in its quest to maximise welfare. In his most recent contribution, Harris suggests that doctors should give up prioritising patients according to prognosis. Harris argues that the good of health care is opportunity. However, I have argued that this approach is problematic on several grounds. I have offered a different alternative: reasons-based maximising consequentialism. Justice, I have argued, involves giving weight to individual's reasons, and what is expected to be good them. On the conception of justice which I favour, distributors of public resources should seek to maximise the satisfaction of equally rational health-related claims which people make on those resources. However, reasons themselves are not tied to the expected value of actions in a scalar (maximising) way, but rather there is a plateau (submaximising) threshold relationship. There can be equal reason to save each of two people, even though the expected value of saving each is different. This approach requires that we give preference to the best prognosis patient in preference to someone with a *much worse* prognosis, but that when prognoses are comparable, we are not rationally required to provide treatment to the patient with the best prognosis.

Centre for Human Bioethics
Monash University