If you are a consequentialist, then what makes an action good or bad is solely determined by the consequences that that action produces. When your action produces the best overall results, then your action is *optimific*. Consequentialism gives us a template to follow so that we can determine which actions are in fact optimific:

1. Identify what is intrinsically good - valuable in and of itself, and worth having for its own sake. (e.g. happiness, autonomy, knowledge, and virtue.)
2. Identify what is intrinsically bad - bad all by itself. (e.g. physical pain, mental anguish, sadistic impulses, and the betrayal of innocents.)
3. Determine all of your options. Which actions are open to you at the moment?
4. For each option, determine the value of its results. How much of what is intrinsically good will each action bring about? How much of what is intrinsically bad?
5. Finally, pick the action that yields the best balance—the highest ratio of good to bad results. That is the optimific choice. That is your moral duty. Doing anything else—failing to strike the greatest balance of good over bad - is immoral.

Step 1 and 2 indicate that there are many ways to be consequentialist. Utilitarians, for example, claim that happiness is the sole intrinsic good, and pain is the sole intrinsic bad.[1] Hence, for Utilitarians, an action is optimific when that action produces the greatest net increase of happiness than any other option. This is called the *principle of utility*.

Now, the principle of utility is sometimes misunderstood. First, it may seem like we must always act in such a way that benefits the most people. This is incorrect. If action A produces a very minimal amount of happiness for a lot of people, but action B produces a tremendous amount of happiness (much more than A) for a small amount of people, the principle of utility tells us to do B. Second, it may seem like we must always act in such a way that creates the most amount of happiness. This is also incorrect. Notice that in the formulation of the principle of utility that the requirement is a *net* increase of happiness. If action A produces the most happiness but also the most pain, and action B produces some happiness but almost no pain, then if the net happiness increase in B is greater than A, then we are required to do B.

Now that we have clarified a bit, there is a further issue to consider. What is more morally important: the actual consequences of our actions, or what we expect the consequences to be? The most common response from Utilitarians is that it is the *actual* consequences of our actions. The actual happiness increases / decreases in the world are what determine the morality of an action. This seems like a straightforward way to understand the principle of utility, but it does lead to a potential issue.

We are quite limited little creatures. We almost never, for example, can say with one hundred percent certainty what the full consequences of our actions will be. How, then, do we satisfy

[1] There are, in fact, numerous kinds of Utilitarianism. Our chapter focuses on John Stuart Mill's version of Utilitarianism which endorses hedonism, and so I will use happiness as the sole intrinsic good, and pain as the sole intrinsic bad.

step 4? Some Utilitarians respond to this problem by focusing not on the actual consequences of our actions, but on what we *reasonably expected* to be the consequences of our actions. This seems like a, well, reasonable amendment. It is largely rejected, however, because of two problems. First, what is reasonable may produce a less than optimific result. If all that mattered was the expected results, then we would be *required* to do this less than optimific action, and that is inconsistent with our definition of consequentialism. Second, we may expect action A to produce bad consequences, but in actuality A will produce optimific consequences. If all that mattered was the expected results, then we would be required to not do A despite it being optimific. This, again, is inconsistent with our definition of consequentialism. Hence, if we focus on the expected results, then we may be required to not act optimifically. You may be fine with that, but a consequentialist is not.

That is not to say that what you expect to happen is not morally relevant. When you expect A to produce optimific consequences, but it doesn't, then your action was not morally correct. But, given that your intention was to act optimifically, you are praiseworthy. Hence, there is disconnect between actions and intentions. Your intentions earn you praise or blame, but your actions are either right or wrong.

Mill

John Stuart Mill was the nephew of Jeremey Bentham. Bentham's Utilitarianism is typically understood as not distinguishing between the qualities of happiness. Mill, however, claimed that there were *higher* and *lower* pleasures. The following famous quote captures this well:

> *"It is better to be a human being dissatisfied than a pig satisfied; better to be Socrates dissatisfied than a fool satisfied. And if the fool, or the pig, is of a different opinion, it is only because they only know their own side of the question"*
> Mill, 'Utilitarianism', p. 281.

Regardless of how you feel about this idea, there is a problem for both Mill and Bentham in regards to how we measure happiness. Pages 136 - 137 have a nice discussion of the measurement problem which will be worthwhile to study for the exam. Additionally, Utilitarianism is often described as a very demanding moral theory. Choosing the optimific action in any given circumstance seems likely to result in our needing to give tremendous amounts of our time and money away to others. Afterall, using our money to save lives and our time to serve those in need will usually be optimific (especially in contrast to what we usually spend our time and money pursuing). So, unless you are doing all that you can all of the time, then it would seem that you are acting immorally according to utilitarianism.

Not all consequentialists, however, accept the five step decision procedure stated above. Having to calculate every result of all relevant actions in every circumstance we find ourselves is, needless to say, quite a burden to bear. I, for example, am probably incapable of properly arriving at what can be reasonably expected by most of my possible actions simply because of time constraints. This is another way that some claim that utilitarianism is too demanding.

Some theorists who wish to retain the spirit of utilitarianism reject the 5 step decision procedure and opt instead for something less burdensome. Instead of focusing on optimific actions, they focus on establishing a set of optimific rules. Such rules, when followed, generally produce more optimific results than if the rules were not followed. A world which, for example, follows the rule ‘do not needlessly harm infants’ has less pain than a world which does not follow that rule. Hence, following that rule produces a world with less intrinsic wrongness.

The main issue for rule consequentialism and rule utilitarianism is what to do when the rules conflict with what is actually optimific. Sometimes breaking a rule will produce more intrinsically valuable results. Is it okay to break the rule? It’s hard to see how given that the basis for following the rules was to avoid having to constantly calculate. If you can break the rules when it is optimific, then we are just back to constantly calculating in order to see if the rule should be followed or not.

If we reject rule utilitarianism we do gain a bit of moral flexibility. No action will ever itself be always immoral or moral. Every action must be weighed against the relevant alternatives. So, in some weird and extreme cases, we may be required to do something that we usually think of as immoral because it will produce the right balance of good and bad consequences. Some see this as a positive of the view, and others see it as a negative.