

The Case Against Moral Luck

David Enoch and Andrei Marmor

USC Public Policy Research Paper No. 03-23



**PUBLIC POLICY
RESEARCH PAPER SERIES**

University of Southern California Law School
Los Angeles, CA 90089-0071

*This paper can be downloaded without charge from the Social Science Research Network
electronic library at <http://ssrn.com/abstract=475161>*

The Case against Moral Luck

DAVID ENOCH

Department of Philosophy and Faculty of Law, The Hebrew University, Jerusalem

ANDREI MARMOR

School of Law and School of Philosophy, The University of Southern California

There seems to be a morally significant difference between reckless driving and reckless driving that results in a fatal accident. There seems to be a morally significant difference between someone who has actually committed a wrong and his counterpart who did not, but would have committed the same wrong if only he had found himself in those same circumstances. And we often make very different moral assessments of persons according to their character traits and personal attributes. It has become extremely difficult to articulate, however, what the differences are, and how we can justify them from a moral point of view. On the one hand, there are good reasons to confine moral assessment of persons' conduct to those aspects of it which are within the persons' control. On the other hand, a great deal of what we do, and in particular, of what we manage to accomplish or fail to achieve, is due to factors beyond our control. Whether we succeed or fail in what we try to do is almost never entirely up to us. Upon closer reflection, we may come to realize that even who we are, what character traits we possess, are just a matter of sheer luck, that they are, to a significant extent at least, not up to us in any meaningful sense. Nevertheless, as Thomas Nagel and Bernard Williams have famously argued, we do not tend to excuse people for things they have done or failed to accomplish just because their success or failure is due to luck. Nor do we tend to abstain from a moral assessment of persons' constitution and character upon realizing that it is really just a matter of luck who they

are and what character traits they ended up having. Perhaps we are forced to this stance upon realizing, as Nagel claimed, that “if the condition of control is consistently applied, it threatens to erode most of the moral assessments we find it natural to make.”¹

This is the idea of moral luck: where we realize that a significant aspect of a person’s character or conduct is due to factors beyond that person’s control, and we nevertheless continue to treat the relevant character trait or conduct as an object of moral assessment, then we are forced to acknowledge that luck plays a constitutive role in our moral evaluations. The thought can be more conveniently put in comparative terms: If there can be a difference in the moral evaluation that is called for without a difference in a factor that is under the control of the relevant agents, then there is moral luck. Thus a denial of moral luck amounts to a supervenience claim: moral status of the relevant kind supervenes on what is under the relevant person’s control.

A dilemma emerges here. Either we adhere to the condition of control and consequently maintain that luck cannot play a role in our moral assessments, or else we acknowledge the inevitability of moral luck and consequently give up the condition of control. Since both the condition of control and the prevalence of moral luck seem philosophically plausible, perhaps even compelling, the dilemma points to a deep tension in any acceptable account of the nature of moral judgments.

Responses to this challenge take one of two general forms. Some philosophers are happy to admit that the moral phenomenology is, indeed, as Nagel and Williams describe. They admit that our moral evaluations are affected by luck, and that the phenomenology is reliable here, namely, that the moral facts themselves are affected by luck. Their main objective, however is to argue that moral luck is reconcilable with

plausible ways of thinking about the nature of morality. Luck is there, they say, but it is not so threatening as might be thought. Other philosophers, however, have taken a different tack, claiming that in fact there is no moral luck, and this is the line of argument that we will advance here as well². We argue that once some crucial distinctions are taken into account, our moral judgments are not as sensitive to luck as the proponents of moral luck suggest. The intuitions, or moral opinions, purportedly supporting moral luck, once carefully characterized, can be accommodated consistently with there being no moral luck. Intuitions of moral luck that cannot be thus accommodated do indeed have to be rejected, but rejecting them comes with an intuitive price that is not unreasonable given the importance of the relevant version of the condition of control. Indeed, one way of explaining away intuitions of moral luck is to attribute them to a failure to distinguish between them and the more benign intuitions that are consistent with there being no moral luck.

Two clarifications are in place before we proceed. The idea of luck is often associated with the occurrence of an event that somehow happens against the odds, as it were. More precisely, perhaps, when the occurrence of an event is deemed to have either objective or subjective probability that is very low, and the event nevertheless occurs, we would tend to say that it is a matter of sheer luck that the event did, in fact, occur.³ This is not the relevant notion of luck to our discussion, however. As Nagel uses the term here, and rightly so, we assume, the idea of luck is not relative to some initial assessment of probabilities. Rather, it points to factors affecting a person's character or behavior which are beyond the person's control. Thus anything that happens, whether likely or not, may be a matter of luck in the relevant sense, if it is not within the agent's control. Needless to say, this is not a satisfactory definition of

luck, but we hope that the kind of cases discussed below will be intuitively clear enough even without a very precise definition of what ‘luck’ means in our context.

The second requisite clarification is a methodological one. Throughout this essay we will assume, without much argument, that the condition of control that would be undermined by the existence of moral luck is philosophically sound and quite crucial to an adequate understanding of the nature of moral assessment and moral judgment. Nevertheless, we do not see it as part of this project to provide a comprehensive defense of the condition of control. Our only assumption is that the condition of control has enough intuitive support and philosophical credentials that one would need very strong reasons to discard it. The main purpose of our argument is to show that proponents of moral luck have not provided such reasons.

We will concentrate on three out of Nagel’s (1976, 28) four categories of (purportedly) moral luck: luck about consequences of actions (consequential luck), luck in the morally relevant circumstances one encounters (circumstantial luck), and luck about moral character (constitutive luck). As is now common in the literature on moral luck, we will try to abstract from the larger issues concerning the freedom of the will, which is why we will not discuss Nagel’s fourth category, concerning luck in how one’s will is caused. There is a worry, however, that abstracting from these (even) larger issues is not philosophically legitimate. We address this worry, though in a somewhat preliminary way, in the last section.

1. Consequential Luck

Arnold and Brian drive back home from a late night party. Both of them know that they have drunk too much and that therefore they should not be driving. Indeed, both of them pass a red light. Arnold is lucky, however; nobody is around that late at

night and he gets home without any incident. Brian's fate turns out to be different. A pedestrian crosses the junction and Brian hits her, killing her on the spot. Let us take this familiar story as our main example of consequential luck. Others would be just as familiar: Cases where instead of recklessness (or negligence) we have a morally bad intention, such as two people attempting murder, with only one of them hitting his target; cases of good intention that in one case but not the other misfires, such as Nagel's examples of political endeavors which turn out to be a disaster, or a glorious victory, thus affecting, allegedly, the *ex post* judgment of our leaders' wisdom and moral stature; cases of personal decisions whose justification may seem to depend on whether they lead to success or failure, or perhaps success or failure of a special kind, such as that of Williams' Gauguin. There may be important distinctions to be drawn between these examples, but for now they can all serve just as well. In all of them, and in numerous other cases, Nagel and Williams claim, the actual consequences of an agent's actions partly determine her moral status.

Let us return, then, to the drunken drivers example. Several aspects of the distinction between Arnold and Brian are uncontroversial. First, and quite obviously, Brian has killed a person, and Arnold has not. It is also true, at least in most jurisdictions, that Brian faces a much harsher punishment than Arnold. Furthermore, as Williams⁴ has pointed out, it is not only the hostile reaction of others that differs with respect to these two drivers. Brian is likely to feel 'agent-regret', that is, regret or perhaps remorse of a special kind, that is not available to spectators (who may ache the loss of a human life just like Brian does), or indeed to Arnold. It seems that Brian *should* feel such agent-regret; what, after all, would we think of someone in Brian's circumstances who did not? Arnold may regret the fact that he took such an unwarranted risk, reproach himself, and promise himself never to drive under the

influence again. But then he can just move on and get on with his life. This is not Brian's situation. Unless Brian is a moral wanton, he will feel remorse for years to come, and his life is not going to be the same ever again. But, *ex hypothesi*, the only difference between the fate of Arnold and Brian is due to sheer luck.

Nagel and Williams believe that the way things happen to turn out in such cases determines something of profound moral significance. But what is it, exactly, in the *ex post* moral evaluation of the situation that is really affected by the consequences? According to Nagel, Brian would be more culpable or more to blame than Arnold. The question is how to make best sense of this claim. Surely, no difference is warranted with respect to the judgment about the moral character of these two agents. They were both equally irresponsible in as much as both exemplified a similar flaw of character, and to a similar degree of severity. Generally speaking, it seems clear that there is nothing in the future consequences of an agent's actions that can affect his moral character at the time of the relevant action. If there is a morally significant difference between the cases of Arnold and Brian, it must reside elsewhere, presumably in their respective blame or culpability. There must be a sense in which Brian's conduct is somehow worse than Arnold's. As Nagel puts it "how things turn out determines what he has done." "Actual results influence culpability or esteem ...".⁵

The problem with this thesis is that in a clear sense it is trivially true, but shows nothing about moral luck, while in another sense, it is false. When we say that Brian has done something worse than Arnold, we could mean either one of the following two ideas. First, we could just mean to say that one state of affairs or complex event, Brian's drunken driving that ends up killing a person, is worse than another, Arnold's drunken driving that does not end up killing anyone. While this is

correct, it does not prove that there is moral luck. It only exemplifies that we can, as we often do, make ethical assessments of events, states of affairs, or consequences. We often incorporate our evaluation of consequences in the evaluation of the event or state of affairs whose consequences they are. For instance, we would say that the earthquake in Turkey was much worse than the earthquake in San Francisco because many more people have died in it. This is bad luck, but not moral luck: no person is being held morally responsible here for something that is beyond their control. Thus this is not the sense in which Brian's action is worse than Arnold's that Nagel could have in mind. Nagel must have meant that Brian is somehow more *blameworthy* or *culpable*⁶ than Arnold.

Perhaps Nagel's argument can be interpreted in the following way. People are morally responsible for their actions, but what action one performs is partly determined by its actual consequences. Brian, after all, performed a killing, whereas Arnold did not. Therefore, only Brian can be held responsible for a killing, which entails that his responsibility is sensitive to consequential luck. But as Nagel himself seems to admit⁷, this line of thought clearly fails. We are never morally responsible for actions *tout court*. Rather, we are responsible for actions *under descriptions*. Every action may have many descriptions under which the agent is not morally responsible for it, even if he is responsible for it in a factual, causal, sense. For example, even if it turns out that Brian has killed the person who was about to invent the cure for AIDS, we would be very reluctant to claim that Brian is morally responsible for millions of deaths that could have been avoided by his victim's survival, or that he is morally responsible for his action under the description 'preventing a life-saving cure from millions.' The fact that there is a sense in which Brian is responsible for causing this catastrophic result does not settle the moral issue

here. What we need is a reason to hold Brian morally responsible for his reckless drunken driving under the description of a killing, and not just under the description (of a reckless drunken driving) that equally applies to Arnold. But any such reason will just be a reason to acknowledge moral luck. At the present stage of the argument, just assuming from the start that Brian is morally responsible for his action under the description of a killing would be a clear case of begging the question against the denial of moral luck.⁸

Thus the relevant question is this: Is a person *morally* responsible for how things turned out as a consequence of his action? And more specifically: Can there be two cases alike in all respects except actual consequences, where nevertheless a differential judgment of moral responsibility is called for? According to the condition of control, the answer to this second question is "no". The scope of moral responsibility depends on how foreseeable, probable, or likely the consequence was, given the relevant information available to the agent at the time of action, that is, *ex ante*. We rightly hold the drunken driver morally responsible for the death of his victim because this consequence is precisely the kind of consequence that might occur if someone is driving under the influence of alcohol, and we assume that under normal circumstances, people know this⁹. In this respect, Arnold and Brian are on par. But now consider Cynthia, who unlike Brian drove her car as cautiously as possible, but in spite of no fault on her part, ran over a pedestrian who crossed the road unexpectedly, killing him on the spot. Once again, though there is a sense in which Cynthia is responsible for the victim's death, we would think that she is not morally responsible for killing him. According to the condition of control, the moral responsibility of an agent for the consequences of his or her action depends only on those factors which could have been assessed prior to the action's performance, and

the relevant responsibility should be judged only on the basis of the choice or decision which was made on the basis of such considerations.¹⁰

One further assumption is needed here: we will assume that an agent is blameworthy (or praiseworthy) for an aspect of an action that he or she performed only if, and to the extent that, he or she is *morally* responsible for it. If this assumption is correct, and we do not see it as begging any question here, it follows that blameworthiness cannot depend on actual consequences, on how things just happen to turn out.

Admittedly, this conclusion might leave many people uncomfortable. After all, they would say, we do differentiate suitable punishment, hostile reaction, and even appropriate first-person regret according to how things turn out. The drunken driver who killed a person is, appropriately, perhaps, liable to harsher punishment and greater condemnation than the drunken driver who killed no one. And, presumably, he ought to feel greater remorse (or perhaps agent-regret) than the lucky driver.

In order to address this kind of worry let us distinguish between moral *blameworthiness* (or praiseworthiness) which is, as we claim, solely a function of moral responsibility and therefore immune to luck, from the appropriateness or justification of what we will call *blame- (or praise-)related reactions*, such as punishment, social condemnation, and even such first-person attitudes as regret or remorse.

Very often, perhaps always, one of the relevant considerations to the justification of blame-related reactions is the blameworthiness of the agent. But it does not follow that this is the only consideration relevant for the justification of blame-related reactions.¹¹ In the case of punishment, according to all but the most extreme retributivist theories (to which we return below), many other considerations

are relevant as well. Consider, for example, considerations of deterrence, or those of prevention, questions about dangerousness, about what is the proper attitude society needs to express towards various crimes, considerations having to do with minimizing the costs of punishment-related practices, and so on. These and similar considerations are often mentioned as relevant to the question of what a justified punishment for a given offense is. Because of the relevance of such considerations to the justification of punishment, it is quite possible that two cases which are identical in degree of blameworthiness merit different punishment. Similar points hold with regard to other blame-related reactions. Seeing that there are many considerations relevant to the justification of, say, an expression of anger, many more than just the blameworthiness of the object of the anger, it is quite possible that two cases identical in blameworthiness are different with regard to the appropriateness of expressions of anger.

This point alone suffices to establish the distinction between blameworthiness and the justification of blame-related reactions, though we have yet to establish its relevance, of course. But even without concrete examples it is clear that the two are conceptually quite distinct. Questions of blameworthiness are about the truth or falsehood of attributions of blame. Questions of the justification of blame-related reactions are practical questions, about the justification of actions or attitudes. The former call for epistemic reasons to decide them, the latter for practical ones¹². And though they may be related, they are doubtless distinct. (You may still have some worries about this distinction, or about the use we are about to make of it. We return to this distinction and possible worries about it in section 4.1.)

Let us return now to the intuitive judgments according to which Arnold and Brian differ in how they can be justifiably punished, or in how they are otherwise

justifiably treated, or indeed in how they should feel. One thing that we can say about such intuitive judgments is that they are, ultimately, groundless; that at the end of the day, on due reflection, they should be rejected. But utilizing the distinction between blameworthiness and the justification of blame-related reactions, another line of reply now opens. For we can agree that there are distinctions in how Arnold and Brian should be treated, and consistently insist that these are not due to differences in blameworthiness, but rather to other considerations that govern the relevant blame-related reaction.

Let us make this point in more detail with regard to punishment. As noted above, most jurisdictions punish Brian more severely than Arnold. Seeing that the only difference between them is a matter of luck, we contend that there is no difference in their moral blameworthiness. At this point, two lines of reasoning are available to us. We can argue that the practice of differential punishment is unjustified, and that it calls for reform¹³. Or we can note that there may be other considerations, not deriving from moral blameworthiness, that call for differential punishment here.¹⁴ If this can be done, then the intuition according to which differential punishment is justified can be accommodated consistently with the denial of moral luck. These two suggestions can work together, so that a distinction between two equally blameworthy agents can be defended in some cases of blame-related reactions (consistently with the denial of moral luck), but other blame-related practices ought to be reformed.

Even the most extreme retributivists about punishment would agree that what determines the justification of punishment is not blameworthiness directly, but rather our *justified beliefs* about blameworthiness. To take an extreme example, if someone is blameworthy but we have no reason to believe that he is (because we have no

reason to believe that he has done whatever it is that he is blameworthy for), we are not justified in punishing him, or indeed in engaging in any other kind of blame-related reaction towards him. This is important, first, because such epistemological considerations are inescapable. Even someone who denies the relevance of any but blameworthiness-considerations for the justification of a certain blame-related reaction, such as punishment, cannot really deny the relevance of epistemological considerations, as the example above shows. Second, this point is significant because it can be argued that actual results often have considerable evidential value.¹⁵ In reality, it is often very hard to determine whether Arnold and Brian were equally reckless in their driving, and in their decision to drive knowing they are intoxicated. All other things being equal, the occurrence of an accident is plausibly considered as at least some *prima facie* evidence for recklessness, or indeed for a higher degree of recklessness. Such epistemic considerations, relevant to the justification of punishment and other blame-related reactions, can go a considerable way in accommodating the purported intuition that we are justified in treating Arnold and Brian differently, consistently with there being no moral luck.

Consequential luck typically arises in cases of risky activities, and this raises a wider question here about the need to internalize various risks of our activities. We have claimed that moral responsibility for the assumption of risks should only depend on those factors which are available to the agent *ex ante*, that is, prior to the performance of the action and based on the information available to the agent at that time. Moral responsibility, however, should be distinguished from the question of the appropriate payoff. Consider, for example, the idea of a morally justified gamble. Let us assume that there are cases in which an agent is warranted in making a certain gamble. The fact that making a gamble under the circumstances is justified entails that

the agent should not be held morally or otherwise blameworthy for the bad consequences, if they occur. But it does not mean that the agent is entitled to the reward of his action just as if it had turned out to be successful, regardless of what happens. Suppose, for example, that someone has made a sound investment decision but nevertheless, the investment turned sour and she lost her money. The fact that the investment decision was sound at the time only entails that the agent does not deserve to be condemned for it. She is not rationally, prudentially, or otherwise, blameworthy. It does not necessarily entail, however, that she ought to get her money back as if the investment had been successful.

There are many reasons for forcing people to internalize the risks of their activities and bear their costs. Presumably, such reasons concern the appropriate incentives for risk taking and risk aversion, depending, at least partly, on the social values of the relevant activity. Payoff allocation is at most an example of a blame- or praise-related reaction, rather than a matter of praise- or blameworthiness, and these are some of the other considerations governing it. Thus the appropriate payoff scheme for one's risky activities need not match the moral responsibility for the assumption of the relevant risk. A denial of moral luck does not entail that life in general should be immunized from luck, as it were. The question of who should bear the costs of various risky activities is, generally, a question of distributive justice. With the exception, perhaps, of those who maintain an uncompromising principle of equality of welfare, all conceptions of fairness allow for possible discrepancies between how well an agent exercises her control when acting and the eventual payoff she receives. A conception of fairness that requires agents to internalize the costs of their risky activities does not necessarily reflect a view of responsibility or blameworthiness. It

may simply reflect a judgment about the appropriate distribution of the costs of risky activities.

Thus the real challenge for proponents of moral luck is to come up with intuitively compelling examples of two cases where the protagonists differ only in what is for them a matter of luck, and yet they differ not merely in the goodness or badness of the state of affairs that includes the relevant actions and their consequences, and not merely in the kind of blame-related reactions or pay-off schemes that are *ex post* appropriate in the circumstances, but also in the degree of blameworthiness and responsibility of the agents for their actions. Without an example that satisfies these conditions, there is not even the beginning of a case for consequential moral luck.

Perhaps such examples can after all be found. But given the force of the intuition underlying the condition of control, we are entitled to reject consequential moral luck until forced to accept it by the kind of example described. Even if such examples are put forward, that is, clean examples, where relevant distinctions are not conflated, and it still seems intuitively plausible that different moral judgments are justified with regard to the moral responsibility or blameworthiness of agents identical in all features but lucky ones, it is not clear that the force behind these intuitions defeats the force behind the intuition underlying the relevant version of the control condition.

We want to conclude the discussion of consequential moral luck by reconsidering a special example of a blame-related reaction, one already mentioned, namely, that of agent-regret. We think that it is especially important not just because of its role in Williams's influential discussion, but also because it is an especially powerful example. : A call for change in current penal practices that will eliminate the

role of luck in distinguishing between otherwise similar cases seems at least somewhat plausible. Similarly, we think, a call for reform would make sense with respect to many other social pay-off schemes and blame-related reactions. But with agent-regret things look different. For it does seem like a robust intuition, one we too would be loathe to discard, that Brian should feel this special kind of regret, the kind that Arnold need not, and that if Brian does not have these feelings he is morally worse for that. This means that in the case of agent-regret, the only line we can plausibly pursue is the one relying on the distinction between the appropriateness of agent-regret and the question of blameworthiness. Fortunately, this option is not at all implausible here. To mention just one relevant consideration¹⁶: In real life situations you never have perfect epistemic access to the relevant facts regarding blameworthiness. In particular, you can always sensibly wonder whether the result was a matter of luck, or whether you were after all more reckless than you think you were, or whether there was more you could and should have done to avoid the horrible result, and so on. Such doubts may rationalize feelings such as agent-regret, but in a way perfectly consistent with the denial of consequential moral luck.

Consider a purer example. In front of you are two handles about which you know nothing but the following: If you do nothing, a disaster will follow, say, a hundred people will die. Pulling one of the handles prevents the disaster, while pulling the other does nothing at all. Assume that you only get to pull, at most, one of the handles. In such a case you are morally required to make a guess, and just pull one of the handles, hoping for the best. You proceed to do just that, only to find out that you have pulled the wrong one. One hundred people die.

Those who think of agent-regret as a guide to moral luck seem committed to saying that in such a case too you are likely to feel, and indeed you should feel, agent-

regret. After all, if you had not pulled the wrong handle many lives would have been saved. In other words, your agency was involved in a causal chain that brought about a horrible result in exactly the way which is characteristic of paradigmatic cases where agent-regret is purportedly called for. But our intuitions about this and similar cases are certainly not as clear as they are in the more common cases of moral luck. We think that the difference is epistemic in nature. In the example described, there is no doubt that you were not at fault in your behavior. You had to take a shot, and you did. There is *ex hypothesi* nothing further you knew or, indeed, could have known, about the two handles in front of you. But if the difference between our hypothetical example and the more common cases that makes a difference with regard to the appropriateness of agent-regret has to do with your warranted confidence in your actual degree of fault, then this lends support to the hypothesis that even in the more common cases agent-regret is fueled at least partly by (reasonable) doubts about how much at fault you really were. And if this is so, the appropriateness of agent-regret is entirely compatible with the denial of moral luck.

The point here generalizes. Even granting Williams that agent-regret is often morally appropriate, no conclusion about moral luck immediately follows. This is so simply because no argument has been given that licenses taking agent-regret as a global indication of moral blameworthiness. Indeed, Williams's argument from the (moral) appropriateness of agent-regret to moral luck may be trading on a crucial ambiguity that is relevant here. Agent-regret is supposedly distinct from remorse, because, unlike remorse, it is compatible with the realization that the regrettable event was not under one's control, that as far as things were up to the agent, she did nothing wrong. And yet for Williams's argument to go through, agent-regret must be taken as an indication of moral blameworthiness or responsibility. But Williams provides no

reason to believe, and it is hard to see how such a reason could be given, that a feeling compatible with the realization that one has done nothing wrong should nevertheless be taken as a guide to moral blameworthiness or responsibility.

Phenomenologically speaking, it may still be true that we feel agent-regret even when reasonably confident that we did nothing wrong. But in such cases we are willing to bite the bullet. For imagine the protagonist in our hypothetical example above, one who through no fault of hers guesses wrong and pulls the wrong handle, who nevertheless reproaches herself for failing to have saved all those innocent lives. It seems to us that in such a case accusing her of irrationality is in no way counterintuitive.¹⁷ Or, at the very least, we would strongly suspect that she confuses regret with remorse.

One final point about agent-regret. At times Williams seems to argue that the appropriateness of agent-regret shows not so much that there is moral luck, but rather that whatever form of evaluation we genuinely care about is subject to luck. When in this mood, Williams would let us have our no-moral-luck thesis if we wanted it, but he would demand that we pay the price of admitting that moral responsibility is of little significance.¹⁸ We cannot discuss here the most general questions about the significance of morality. Nor do we claim that only evaluations that are immune to luck are important in our lives. But we do want to emphasize that there is nothing artificial or ad-hoc about the sense of moral responsibility or blameworthiness that is immune to luck. Moral responsibility thus understood captures the mode of evaluation relevant, somewhat roughly, to how (morally) well one exercises one's control. While Williams may very well be right that agent-regret shows that this is not the only mode of evaluation we do, or should, care about, agent-regret in no way shows that this is not a mode of evaluation worthy of serious consideration and respect.

2. Circumstantial Luck

It is hardly deniable that various kinds of circumstances determine, to a large extent, who we are, what we can do, and what we end up doing. Character traits and human behavior are not exempt from the causal chains of nature. In obvious ways, it must be true that numerous contingent circumstances affect our personal character, inclinations, actions, and what not. However, it is not this causal impact of circumstances on people's character and behavior that Nagel labels 'circumstantial luck'. 'Circumstantial luck' is the kind of luck that creates the morally relevant challenges and opportunities people encounter in their lives. We will follow Nagel in this respect and discuss here only the role of circumstances in creating the moral opportunities we face, and not their causal bearing on who we are and what we do. The latter will be discussed in the next sections.

Thus, consider the following case. Green and White are both equally, but separately, determined to kill Orange. Green carries out his plan and kills Orange. By the time White is about to carry out his own plan, he realizes that Green had already done the job, so he does not need to do the killing himself. Let us assume, however, that White would have killed Orange himself in case Green had not done it ahead of him. Of such cases Nagel claims that "here again, morality is at the mercy of fate", as we cannot resist the inclination to "judge people for what they actually do or fail to do, not just for what they would have done if circumstances had been different."¹⁹ Surely the only difference between Green and White here is due to luck. White was lucky that Green succeeded in his assassination, so he never got the chance to kill Orange himself. But, again, he would have done so if Green failed.

The problem of circumstantial moral luck is very similar to the problem of consequential luck. It is difficult to deny that we do have rather strong intuitions according to which Green is more to blame than White, even when we realize, upon reflection, that the difference between them is only due to luck in the circumstances they confronted. The question is, again, whether we can make sense of these intuitions. Here is one way to look at it: suppose Green confronts White telling him, ‘Don’t you condemn me! You are just as bad!’. Surely, in one clear sense, this is, ex hypothesis, true: we have assumed that they were equally determined to assassinate Orange. So if the assertion of Green, ‘you are just as bad’ is meant to refer to White’s intentions, or moral character, or such, then it is trivially true. On the other hand, if the assertion is meant to refer to White’s actions, the statement would be trivially false: it is a fact that Green killed Orange and White did not. But of course, this only shows that circumstances may determine, to some extent, what people do. It shows that luck plays a role in our lives, but the question remains whether it is in any significant sense an issue of moral luck. So if there is a difference in moral appraisal between Green and White, it must reside elsewhere. Yet again, it must relate to a difference in blameworthiness, culpability or responsibility. We would have to maintain that Green is just more blameworthy than White. After all, he killed Orange, and White did not. In this case, however, we suspect that epistemic considerations account for most of the difference in our intuitive reactions. Suppose White responds to Green’s accusation that he is just as bad by saying: ‘Yes, perhaps I’m just as bad as you are, but we will never know for sure, will we?’ As annoying as such a retort may be, it does seem to be very much to the point. Unless people actually perform an (intentional) action or, at least, take unequivocal steps in an attempt to perform it, the inference from intention (or inclination, propensity, etc.,) to moral blameworthiness

will be shadowed by some uncertainty, no matter how strong the intention is, or how certain we feel about it. Consider, as another example, the case Nagel mentions of the Nazi collaborators²⁰: we tend to condemn many Germans who have succumbed to the terror of the Nazi regime and collaborated with it even though we may well suspect that most of us would have succumbed to the same forces as well, had we been in those terrible circumstances. But again, the truth is that we will never know for sure. So when we are much more inclined to blame a person for what he has actually done or refrained from doing, rather than for what he would have done under different circumstances, this is, we submit, mostly because we have some doubts about what it is, exactly, that the person would have done if circumstances had indeed been different. Even if we are quite certain that White would have shot Orange if Green failed to accomplish his mission, we cannot be entirely sure. And even if we are, we cannot know the specific features of the relevant counterfactual action.

Notice further that our employment of the epistemic argument here enjoys the following advantage. The epistemic argument can be employed, as it has been employed in the previous section, to accommodate intuitions about consequential luck as well. But there it can be countered by considerations of other relevant evidence.²¹ Perhaps actually causing harm is some evidence for reckless driving, but where other strong evidence is also available there may be an all-things-considered reason to believe that lucky Arnold was more reckless than unlucky Brian. The possibility of counter evidence restricts the force of the epistemic line in accommodating the intuitions about consequential moral luck. But the case of circumstantial moral luck is different, in that countervailing evidence, though possible, is unlikely to be very strong, and is guaranteed never to be as strong as the evidence it is supposed to outweigh. For what evidence about a purportedly would be Nazi-collaborator could

possibly warrant a belief that he would have really collaborated with the Nazis to the same degree of confidence we have about the actual a collaborator? Thus, the epistemic line is much more promising with regard to purported cases of circumstantial moral luck than it is with regard to cases of consequential moral luck, where, as we have seen, other replies do most of the work.

Further evidence for the key role of epistemic considerations in underlying the intuitions about circumstantial luck can be gained from the following observation. The closer the agent gets to the relevant wrong action, the more inclined we are to condemn her. In the examples already mentioned, we are more inclined to think badly of White for the murder he would have committed than we are to think badly of the would be Nazi collaborator for his counterfactual collaboration. But if we really do judge people for what they do and not for what they would have done, this should seem puzzling. So long as the agent did not perform the action, why should it matter how close to performing it he got? At least one plausible answer is that we do judge people for what they would have done, and the closer they get to action the better our evidence that they would have indeed performed it. This is not, of course, the only possible explanation, and some alternative explanations are not as supportive of our case against moral luck. It may be argued, for instance, that the closer one gets to action the more other, preparatory actions one performs, or the more bad intentions one has, and the additional blame can be attributed to those. But examples can easily be constructed where such alternative explanations seem less adequate than the one suggested above, as when there are no significant preparatory actions or intentions. Thus, the hypothesis that what explains the intuitions regarding circumstantial luck are mostly epistemic considerations gains some support from the intuitive relevance of how close to action the relevant agent got.

Furthermore, if we focus our attention on blame-related reactions, as opposed to mere blameworthiness, it is much easier to explain the reluctance to condemn, denounce or punish people for things that they would have done under different circumstances. In addition to the epistemic considerations we have been emphasizing, there are other familiar considerations against engaging in blame-related reactions that are not a reaction to things people have actually done, or refrained from doing, under the circumstances which did, in fact, obtain. Consider, for example, a case of praise-related reaction that is affected by circumstances. Even if you tend to deny that there is, generally, a complete symmetry between praise and blame²², surely the problem of moral luck arises with regard to both. Thus suppose that Debbie and Edward are particularly good natured, benevolent, and willing to sacrifice their own well being for the sake of others. Both of them are potential heroes, so to speak, and to the same extent. As it happens, only Debbie gets the opportunity to exercise her heroism: While taking her morning walk on the beach she observes a child drowning in the ocean, so she saves him while risking her own life. Edward never gets a similar chance. His potential heroism is never manifest in action. Naturally, Debbie will be praised and her heroism celebrated but, of course, it is very unlikely that Edward will be equally praised because he would have done exactly the same if he had been walking on the beach that morning. Assume now, for the sake of the argument, that we are convinced that Edward would have done exactly the same act of heroism under similar circumstances. Nevertheless, it is not difficult to see that we have reasons to praise Debbie and not Edward. For example, one of the values we attach to praise and public celebration of good deeds is educational. We would want Debbie to become a role model, to be emulated by others. It would be difficult to achieve such educational goals without being able to point to the concrete action that Debbie had performed

and tell the story that provides it with the educational value it has. In this case, as in countless others, praise-related reactions have their own rationales which we can distinguish, quite clearly, from the questions of praiseworthiness themselves.

And what if Edward remains somewhat resentful that, after all, he is just as good? To some extent we can come to understand and sympathize, but we might also feel that it is a childish reaction. In fact, it often happens with young siblings that one of them gets an odd chance to exhibit good behavior while the other does not. Sensitive parents would be inclined to restrain their praise in such cases, reassuring the other sibling that they know that he or she would have been just as good if circumstances allowed. Grown ups are expected to know this without reassurances.

There is a general point worth emphasizing here. It is certainly not our aim to argue that people's lives, or how well they go, can somehow be immune to luck. The question is not whether our life is at the mercy of fate. Unfortunately, often it is. The relevant question is whether our moral standing and judgments are at the mercy of fate. The discussion of circumstantial luck does not seem to support the contention that it is.

3. Constitutive Luck

Some people have a bad moral character, and although some aspects of one's character may be amenable to change, Nagel says, "it is largely a matter of constitutive bad luck. Yet people are morally condemned for such qualities, and esteemed for others equally beyond control of the will: they are assessed for what they are *like*."²³ In this section we discuss this element of luck in the constitution of people's moral character.

We assume in what follows that there are considerable differences in people's character traits. We will also assume that some aspects of a person's character may be amenable to change by her will and effort, but others are not. We will also take it for granted that there is constitutive luck. The question we want to examine is whether there is constitutive moral luck. In other words, we do not wish to deny that a great deal of what makes a person the kind of person she is, is determined by luck: Luck in the genes she inherited, in the environment in which she grew up, or in anything else which might have a causal role to play in the constitution of her character. Furthermore, even conceding that certain aspects of people's character are within their control, it may still be a matter of luck how much change, and of which kind, is achievable by one's will and effort.

There are two ways in which constitutive luck becomes a problem of moral luck. First, as Nagel observes, judgments and appraisals of moral character, *per se*, seem unavoidable. We condemn and praise people simply for what they are. But in addition to the appraisal of moral character as such, there is also a concern about the causal role character plays in the performance of actions. It is difficult to deny that character traits may causally affect actions or, at the least, the propensity to perform certain actions which may be morally significant. A miser is less likely to give to charity than others, an egocentric is much more likely to decline your call for help, and so on. Let us call these *direct* and *indirect* constitutive moral luck, respectively. Notice, by the way, that if any of these assumptions are empirically false and it is actually the case that people do not have anything worth calling character traits, this *strengthens* the case against constitutive moral luck.

The importance of indirect constitutive moral luck is the following: Some philosophers (allegedly Kant) sought to avoid the problem of constitutive luck simply

by denying its moral significance. They claim that judgments of character are just not within the bounds of morality: We should never morally condemn or praise people for what they are, only for what they do. But if indirect constitutive luck is taken into account, then it must be acknowledged that what people do is partly due to who they are, and this brings constitutive luck back into the picture. Even if we can avoid moral judgment of character *per se* (though it is far from clear that we either can or should, and we shall discuss this below), we cannot ignore the causal role that character traits may play in the realm of actions. To the extent that there is any significant impact of character traits on persons' actions, constitutive luck becomes a problem of moral luck even for those who may wish to deny that moral judgment of character, *per se*, is ever a sensible type of moral evaluation. The good news is, however, that if direct constitutive moral luck can be plausibly denied, the indirect version ceases to be a matter of concern. Therefore, we will focus our discussion on the question of moral evaluation of character *per se*.

The problem, of course, is this: On the one hand, it seems that we just cannot avoid making moral judgments about people's character and, more importantly, such judgments seem mostly quite appropriate. It is not just that we have an urge to condemn the miser or the egocentric and praise the generous or the kind, but it also seems quite appropriate to do so. At the very least, we would need very strong reasons to retract or discard these intuitive judgments. On the other hand, when we condemn someone for what he *is*, we seem to be making a moral judgment about things that are out of his control, things that just happen to be as they are, purely a matter of luck. Are we committed, then, to the existence of constitutive moral luck?

The problem of moral luck arises, remember, with regard to those moral judgments that are closely tied to agency, judgments of responsibility, culpability,

blame or praiseworthiness. But there is a wide range of moral and ethical evaluations that is not related to moral agency in such a way. Consider, first, this analogy with aesthetic or artistic evaluation. You look at a work of art, thinking how great it is. You may have two kinds of aesthetic appraisal here, and they are not mutually exclusive. You may think about the work of art as the product of deliberate design, as an accomplishment of an artistic scheme. In this case, you focus on the artist as an agent and admire her accomplishment. But you may also ignore the artist and just think about the aesthetic and artistic qualities that are in the work, admiring them as such. The qualities of the work, you may think, manifest certain artistic values. They highlight certain subtleties of human emotion, bring out an aspect of the human condition hitherto unnoticed, or such. And you need not care about the question of whether these qualities were intended by the artist or deliberately planned. The work of art is a good one, you think, just because it has these qualities.

A similar duality is present in the realm of moral and ethical evaluations. We make moral judgments about aspects of the world, including human character, that are quite independent of agency. Such evaluations do not involve attributions of blame and responsibility²⁴. For example, you may think that it is a morally bad state of affairs that many people are so greedy, or indeed that Frank is. It is a morally regrettable fact, a bad aspect of the world we live in. And this is so even if there is no one to blame for it. Notice that this is still very much a moral, or at least an ethical evaluation, not merely an expression of a subjective preference for things to be, or to have been, otherwise. Just as a work of art may have artistic qualities even if they were not intentionally produced as such, aspects of the world have morally or ethically significant qualities even if no one is responsible for bringing them about. When we point out the moral qualities of a certain state of affairs, we purport to

describe a moral fact, a moral aspect of the world, and therefore we engage in a certain form of judgment.²⁵

Now consider moral judgments about character, *per se*. Would it be correct to assume that such judgments of character are of this non-agency-related kind? Often they are, but not always. Nagel is right to claim that a criticism of a certain character trait often implies more than an evaluation of something which is just given.

“Condemnation implies that you should not be like that, not that it is unfortunate that you are”.²⁶ Condemnation of a character trait often implies that you need to change, and this involves an assumption about agency and control. But this is often quite sensible. People can change, to some extent, or at least they could try. Even when we understand that Frank is not to blame for his greediness we may think that he ought to try to change or to overcome his natural inclinations, at least to some extent, and then we may criticize him for not even trying, or not trying hard enough.

The distinction between those cases in which there is room for an agency - related judgment, and those cases where there is only room for the non-agency kind of judgments is not at all sharp, and one can easily think of borderline cases between the two. For example, think of those friends of ours who are always late. Chronically, they just never make it on time for anything! It might be very difficult to say which aspects of this kind of behavior are a matter of moral agency and action, and which are a matter of pure character trait. We expect such people to realize that their behavior is unacceptable, and that they should really try to make an effort not to be late. On the other hand, we may also come to realize that some people just find it extremely difficult, almost impossible, to get their schedule under control. To some extent they just cannot really do much about it. In some such cases it may be very hard to know whether a responsibility judgment is in place. In others, it may be

genuinely indeterminate. But here, as in other cases of vague distinctions, vagueness implies neither the absence of a conceptual distinction nor its normative insignificance.

Indeed, the nature of the distinction between cases where responsibility judgments are in place and those in which they are not helps to explain why our reaction to bad moral character is often one of condemnation and not just disappointment (and to good character it is typically one of praise and not just admiration). We expect people to change, or at least to make an effort in that direction. But it is noteworthy that on closer reflection we often realize that the temptation to condemn should be restrained. This is one aspect of intimacy; the more you know about a person's biography and come to a better understanding of her situation, the more you may come to realize that even if she has bad character traits, they are of the kind which should be regretted, not necessarily condemned. If you find out that a friend who is chronically late has really made every reasonable effort to get her schedule under control, but all of it is to no avail, would you not be inclined to curb your condemnation? You can still, of course, regret the fact that you can never rely on her to be on time, and perhaps you can even treat this as a reason to avoid her company on some occasions, but should you condemn her for this? We think not, and we think that what explains this fact is your tacit acknowledgment that when something really is out of your friend's control, she cannot be morally responsible for it, and therefore moral condemnation cannot be justified. A similar point can be made with regard to morally good character traits. The typical reaction to a person's good moral character is admiration, a recognition that such character traits should be emulated, and we normally reserve reactions of moral praise or acclaim for people's achievement or effort.

Let us conclude the discussion of constitutive moral luck: We can make moral judgments about character traits regardless of responsibility and agency. Often, however, a condemnation of an aspect of a person's character trait makes perfect sense, namely, when it is based on an expectation that the person take certain actions or avoid them as an attempt to improve his or her character. The distinction between the cases where this holds and the cases where it does not is vague and epistemically hard to track, and this already explains some of our inclination to blame and praise people for their moral character even in cases where it is not entirely clear that such attitudes are in place. Nevertheless, our evaluative practices are sensitive to such distinctions, as can be seen from the fact that the more convinced we are that the relevant agent has done all that he could to become of better moral character, the less we are inclined to blame him. This does not mean that in such cases there is no room for any relevant moral or ethical evaluation, for there is still room for the kind of evaluation that does not relate directly to agency.

One final point: Some philosophers espouse a character-based theory of blame or of responsibility.²⁷ They believe, roughly, that you are responsible, or blameworthy, for a wrong action of yours if and only if, and to the extent that, it reflects badly on your character. How does accepting a character-based theory of blame and responsibility affect the discussion of moral, and in particular constitutive moral, luck? A character-based theory of blame and responsibility straightforwardly entails that there is neither consequential nor circumstantial moral luck (for in the relevant examples it is conceded from the start that, say, Arnold and Brian, or the Nazi-collaborator and the would-be-Nazi-collaborator are alike as far as their morally relevant character-traits are concerned). But character-based theorists have to pay a price for this elegant result when it comes to constitutive luck. For it follows from

their view, unless it is suitably qualified, that you are responsible for your morally relevant character traits just in case they reflect badly on your morally relevant character traits, which they trivially do in all cases. On this view, in other words, you are responsible for *all* your morally relevant character traits, regardless of whether they are or ever have been in any interesting way under your control. So proponents of such a view must accept constitutive moral luck (unless they are willing to claim implausibly that our character-traits are entirely under our control²⁸). We find this result problematic, of course, but we do not want to pretend that we have here a clear-cut argument against character-based theories of responsibility and blame. A comprehensive critique of such theories exceeds the scope of this essay.

4. Two Remaining Worries

In this concluding section we want to address two remaining worries. The first is about the distinction between blameworthiness and blame-related reactions, and it threatens much of the argument in previous sections. The second is a more general worry about the philosophical motivation for the denial of moral luck.

4.1. Blameworthiness and Blame-related reactions Again

The distinction between blameworthiness and blame-related reactions played a central role in the discussion so far, both in accommodating those moral-luck-intuitions that are consistent with the denial of moral luck and in explaining away those that are not. But this distinction may raise the following worry.

At least one philosophical account of blameworthiness, inspired by Strawson's discussion of reactive attitudes in 'Freedom and Resentment' is *in terms of* blame-related reactions or responses.²⁹ Blameworthiness, according to this argument, just

consists in the appropriateness of certain blame-related reactions, or attitudes, emotions, beliefs, desires etc., (we will just use ‘reactions’ from now on to denote all of these). Now, if this is really how blameworthiness is to be understood, there is room for the worry that the distinction we have been making so much of collapses, and with it much of the case against moral luck. For we have been arguing, in effect, that blameworthiness is one thing, and the appropriateness of blame-related reactions another, so that intuitions about the latter need not carry over to the former. But if blameworthiness just is the appropriateness of certain blame-related reactions, this line of thought clearly fails.

In reply, let us make two points, the cumulative weight of which suffices, we think, to allay this worry. First, though we cannot, of course, offer here a full philosophical account of blameworthiness, we want to voice our doubts about the prospects of understanding blameworthiness in terms of the appropriateness of blame-related reactions. Judgments about blameworthiness, it seems to us, and related judgments about responsibility, culpability and the like, are not logically equivalent to judgments about the appropriateness of blame-related reactions. Rather, the former serve *to ground* or *justify* the latter. It is because a person is blameworthy that certain blame-related reactions towards him are appropriate. It is not the case that because certain blame-related reactions towards a person are appropriate that he is blameworthy³⁰. As is often the case with response-dependence theories of normative concepts, an understanding of blameworthiness in terms of the appropriateness of certain blame-related reactions fails to respect the nature and direction of this because-relation and is thus unacceptable.³¹

Second, even if at the end of the day the best theory of blameworthiness is going to be in terms of the appropriateness of certain blame-related reactions, still the

main line of our argument against moral luck can stand almost without change. For such a theory to be plausible, it must designate certain reactions – crucially not all of them – as the ones the appropriateness of which determines, or indeed, just is, blameworthiness. Otherwise, if all blame-related reactions are deemed constitutive of blame, no room will be left for any critical appraisal of blame-related reactions. Surely, that does not make any sense. Any reasonable account of blameworthiness must make some room for ethical criticism of certain blame-related reactions. So let us assume that there are some *core reactions* that constitute blame. For example, the theory may designate resentment and social condemnation as the core reactions, arguing that blameworthiness consists in the appropriateness of resentment and social condemnation, but conceding that blameworthiness and, say, the appropriateness of punishment are conceptually distinct. Now, when a certain reaction is such that its appropriateness is determined by a number of different considerations, some of which are not directly related to the agent and her relevant action or character trait (like, in the case of punishment, considerations of general deterrence), this will be a strong reason not to designate this reaction as a core reaction.

Therefore, whenever we utilized the distinction between blameworthiness and blame-related reactions, the distinction between core and non-core reactions may just as comfortably be used. For now intuitions that seem to suggest that blameworthiness, that is, the appropriateness of core reactions such as resentment and social condemnation, depends on luck, can be accommodated by characterizing them as really applying to non-core reactions, such as appropriateness of punishment. Putting the distinction in these terms does not, it seems to us, take away any of the plausibility of utilizing it in order to accommodate some of our moral-luck-intuitions and explain away others.

To sum up: Views that reduce blameworthiness to the appropriateness of blame-related reactions do not pose a serious threat to our case against moral luck. This is so, first, because they are, as philosophical accounts of blameworthiness, implausible, unless considerably modified, and second, because more plausible versions of such views will themselves be committed to a related distinction, that between core and non-core reactions (that is, between the reactions purportedly constitutive of blameworthiness and those that are not). In this case, one can accommodate our argument and employ the distinction between core and non-core reactions rather than the one between blameworthiness and the appropriateness of blame-related reactions.

4.2 Can We Ignore the Freedom-of-the-Will Debate?

The literature on moral luck by and large avoids entering the debate about freedom of the will.³² Indeed, the questions regarding moral luck can be formulated independently of that (even) larger topic. At the outset, we characterized the moral luck issue as a debate over a supervenience claim: The enemies of moral luck affirm, and its friends deny, that moral standing of agents supervenes on what is under their control. According to the foes of moral luck, in other words, there can be no difference in moral standing, such as responsibility, blame or praiseworthiness, or culpability, without a difference in something that is under the relevant agent's control. Although such a characterization of the issue is not as clear and precise as one could hope, as it remains to be seen, for example, how 'moral standing' and 'control' are best understood, it is clear enough to capture the paradigmatic cases that are in dispute between the friends and foes of moral luck. This understanding of what is at issue does not directly depend on one's views about the freedom of the will. The

question of whether Arnold and Brian from section 1 above are morally on a par does not, it seems, depend for its significance or relevance on one's favorite view regarding the freedom of the will. Unless, that is, your favorite view regarding the freedom of the will is that there is none, and you maintain that it follows that none of us is morally responsible for anything, and so that everyone, Arnold and Brian included, are morally on a par, at least with regard to responsibility and blameworthiness. We return to this shortly.

Let us briefly address the following nagging doubt. Though the two discussions, the one about moral luck and the one about the freedom of the will, are distinct, they are certainly related. Indeed, they may be very closely related, and in a way that supports moral luck, for it may be thought that an examination of the difficulties concerning the freedom of the will and the most general conditions of moral responsibility undermines any philosophical motivation for the condition of control underlying the denial of moral luck.

The thought we are gesturing at here is that the view according to which we are only morally responsible for what is under our control is one we are going to have to discard even regardless of the discussion of moral luck, because it is ultimately at odds with any plausible metaphysics. It requires agent-causation, or perhaps the denial both of human actions being pre-determined and of them being random, or something of this sort. If so, we have independent reason to doubt any version of the control condition, and thus will be left with no intuitive support for the denial of moral luck. True, the thought goes, one can intuitively distinguish between things we have more and less control over, such as our pulling the trigger and the bullet hitting our target respectively, and one can, if one wants, base a discussion of moral luck on such distinctions. But given the allegedly deeper doubts about freedom and

responsibility, such distinctions are philosophically unmotivated. The difference in degree of control can seem irrelevant, more or less of the same, not a difference that makes a philosophical and moral difference.

Now, this worry may be raised with regard to all kinds of moral luck, but it may be thought of as somewhat more limited in scope. We can think of a continuum, starting with cases of (purported) consequential luck like that of Arnold and Brian, then moving to cases of close-call circumstantial luck (where the relevant agent gets very close to the action he would have performed), then gradually to much more remote cases of circumstantial luck (like that of the would-be Nazi collaborator), then to cases of constitutive luck with regard to rather easily changeable and controllable character traits, then finally to cases of constitutive moral luck with regard to character traits we have no control over. And, it may be argued, the further you go along this continuum the closer you get to the problems of free will, and indeed the more acutely the problem in the previous paragraph arises. As we proceed along the continuum, then, affirming a control condition in the face of the problem of free will seems less and less well motivated.

This is a serious worry, and certainly not one we want to underestimate. But what would really follow if things were as they are suggested to be by such a worry? If we are no more in control in the philosophically and morally relevant sense of our pulling the trigger, or even of our intending to pull the trigger, than we are in control of the bullet hitting the target, or even of the competence of the doctor who treats the person we shot, then our case against moral luck does indeed collapse, but with it, we think, collapses all of morality. That there is *some* distinction along these lines that can be defended – a rough and vague one, certainly, and perhaps also context-dependent, but a distinction nonetheless – is, we think, so deeply entrenched in what

we do when we moralize that there can be no morality without it. It is here that Nagel's warning seems to be accurate: "The area of genuine agency, and therefore of legitimate moral judgment, seems to shrink under this scrutiny to an extension-less point."³³

It is not, of course, philosophical news that some views about the metaphysics of freedom threaten to destroy morality (or at least its applicability to creatures such as ourselves), and it is no surprise that in such a situation nothing remains from our case against moral luck. In such a case, by the way, nothing remains from any case *for* moral luck either. If, on the other hand, such a catastrophe can be avoided, then the case against moral luck stands as before. Either way, then, the freedom of the will debate does not threaten the particular case against moral luck: either it does not threaten it at all, or it threatens it only as a part of the threat it poses to morality as a whole.³⁴

¹ Nagel, T. (1976), 'Moral Luck', *Proceedings of the Aristotelian Society*, suppl. 50, 137-152; reprinted (in a revised version) in Nagel, T. (1979), *Mortal Questions* (Cambridge: Cambridge University Press), 24-38, at 26.

² Attempts at the first line of response include Adler, J. E. (1987), 'Luckless Desert is Different Desert', *Mind* 96, 247-249; Walker, M. U. (1991), 'Moral Luck and the Virtues of Impure Agency', *Metaphilosophy* 22, 14-27; reprinted in Statman, D. (ed.) (1993) *Moral Luck* (Albany: State University of New York Press), 235-250; and perhaps also Conception, D. W. "Moral Luck, Control, and the Bases of Desert", *The Journal of Value Inquiry* 36 (2002), 455-461. For the second line of response, see Jensen, H. (1984) 'Morality and Luck', *Philosophy* 59, 323-330; reprinted in Statman, 131-140; Richards, N. (1986), 'Luck and Desert', *Mind* 65, 198-209; reprinted in Statman, 167-180; Thomson, J. J. (1989), 'Morality and Bad Luck', *Metaphilosophy* 20, 203-221; reprinted in Statman, 195-215; Rescher, N. (1990), 'Moral Luck', *Proceedings of the American Philosophical Association* 64, 5-20; a revised version printed in Statman, 141-166. Rosebury, B. (1995), 'Moral Responsibility and 'Moral Luck'', *The Philosophical Review* 104, 499-524; Zimmerman, M. (2002), 'Taking Luck Seriously', *The Journal of Philosophy* 99, 553-576; Domsly, D. (2004) "There Is No Door: Finally Solving the Problem of Moral Luck", *The Journal of Philosophy* 99, 445-464; and Royzman, E. and Kumar, R. (2004), "Is Consequential Luck Morally Inconsequential? Empirical Psychology and the Reassessment of Moral Luck", *Ratio* 17, 329-344.

For attempts at intermediate positions, see Andre, J. (1983), 'Nagel, Williams, and Moral Luck', *Analysis* 43, 202-207; reprinted in Statman, 123-129; and Wolf, S. (2000) "The Moral of Moral Luck", *Philosophical Exchange* 31, 5-19.

³ See, for instance, Latus, A. (2003), 'Constitutive Luck', *Metaphilosophy* 34, 460-475.

⁴ Williams, B. (1976), "Moral Luck", *Proceedings of the Aristotelian Society*, suppl. 50, 115-136; reprinted (in a revised version) in Williams, B. (1981), *Moral Luck* (Cambridge: Cambridge University Press), 20-39, at 27-30.

⁵ Nagel, op. cit., at 29-30.

⁶ Cf., Zimmerman, op. cit., at 554.

⁷ Nagel, op. cit., at 31.

⁸ Cf. Wolf, op. cit.

⁹ Cf. Rosebury, op. cit., at 518-9.

¹⁰ See Dan-Cohen, M. (1992) 'Responsibility and the Boundaries of the Self', 105 *Harvard L.R.*, 959-1003.

¹¹ Cf. Jensen, op. cit. and Zimmerman, op. cit., at 556.

¹² See Zimmerman, M. J. (1987), 'Luck and Moral Responsibility', *Ethics* 97, 374-386; reprinted in Statman, 217, at 218-9 and 229-230 for a similar point.

¹³ See, for instance, Schulhofer, S. J. (1974), 'Harm and Punishment: A Critique of Emphasis on the Results of Conduct in the Criminal Law', 122 *University of Pennsylvania Law Review* 1497-1607; Shachar, Y. (1987) 'The Fortuitous Gap in Law and Morality', 6 *Criminal Justice Ethics*, 12-36; Kadish, S. H. (1994), 'Foreword: The Criminal Law and the Luck of the Draw', 84 *The Journal of Criminal Law and Criminology*, 679-702; Kessler, K. D. (1994), 'The Role of Luck in the Criminal Law', 142 *University of Pennsylvania Law Review*, 2183-2237; and Feinberg, J. (2003) 'Criminal Attempts: Equal Punishments for Failed Attempts', in his *Problems at the Roots of Law* (Oxford: Oxford University Press), 77-102.

¹⁴ See, for example, Shavell, S. (1990), 'Deterrence and the Punishment of Attempts', 19 *Journal of Legal Studies*, 435-466, Ben-Shahar, O. and Harel, A. (1996), 'The Economics of the Law of Criminal Attempts: A Victim-Centered Perspective', 145 *University of Pennsylvania Law Review*, 299-351, and Lewis, D. (1989), 'The Punishment that Leaves Something to Chance', *Philosophy and Public Affairs* 18, 53-67; reprinted in his *Papers in Ethics and Social Philosophy* (Cambridge: Cambridge University Press, 2000), 227-243.

¹⁵ See Richards, op. cit.; Thomson, op. cit., at 205; and Rescher, op. cit., at 154.

¹⁶ One also emphasized by Rosebury, op. cit., at 515, and mentioned by Latus, A. (2001), 'Moral Luck', *The Internet Encyclopedia of Philosophy*, available at <http://www.utm.edu/research/iep/m/moralluc.htm>

¹⁷ Cf., Levi, D. S. (1989) 'What's Luck Got to Do with It?', *Philosophical Investigations* 12, 1-13; reprinted in Statman, 109-121, at 120.

¹⁸ See, for instance, Williams, B. (1993), 'Postscript', in Statman, 251-258, at 254.

¹⁹ Nagel, op. cit., at 34

²⁰ Ibid., at 26.

²¹ Cf., Adler, op. cit.

²² Cf. Wolf, S. (1980) "Asymmetrical Freedom", *The Journal of Philosophy* 77, 151-166.

²³ Nagel, op. cit., at 33.

²⁴ For a similar point see Zimmerman, 'Taking Luck Seriously', op. cit., at 557.

²⁵ Cf., Andre, op. cit., at 126-7.

²⁶ Nagel, op. cit., at 33.

²⁷ See, for example, Richards, op. cit., at 169; and Thomson, op. cit.

²⁸ See Richards, op. cit., 172-3.

²⁹ Strawson, P. F. (1962), 'Freedom and Resentment', *Proceedings of the British Academy* 48, 1.

³⁰ See Zimmerman, 'Taking Luck Seriously', op. cit., at 555.

³¹ For our own stab at this, see Enoch, D. 'Why Idealize?', forthcoming in *Ethics*; and Marmor, A. (2001) *Positive Law and Objective Values*, (Oxford: Oxford University Press), ch 8. .

³² See, for instance, Andre, op. cit., at 124-5.

³³ Nagel, op. cit., at 35.

³⁴ For helpful conversations and comments on earlier drafts, we thank Meir Dan Cohen, Chaim Gans, Alon Harel, David Heyd, Sharon Lloyd, Joseph Raz, Josh Schechter, Danny Statman and Gideon Yaffe. Previous versions were presented at the Southern California Law and Philosophy Reading Group and at a seminar on moral and legal luck at the Hebrew University in Jerusalem, in the fall of 2003. We thank the audiences for their helpful comments.